



# The VoltDB Codeline

John Hugg

I was hired by Mike Stonebraker to commercialize the H-Store<sup>1</sup> research [Stonebraker et al. 2007b] in early 2008. For the first year, I collaborated with academic researchers building the prototype, with close oversight from Mike Stonebraker.<sup>2</sup> Andy Pavlo and I presented our early results at VLDB 2008 [Kallman et al. 2008] in August of that year. I then helped lead the efforts to commercialize VoltDB, ultimately spending the next ten years developing VoltDB with a team I was privileged to work with. In my time at VoltDB, Inc., Mike Stonebraker served as our CTO and then advisor, offering wisdom and direction for the team.

VoltDB was conceived after the success of Vertica<sup>3</sup>; if, Vertica, a system dedicated to analytical data, could beat a general-purpose system by an order of magnitude at analytical workloads, could a system dedicated to operational data do the same for operational workloads? This was the next step in Mike Stonebraker's crusade against the one-size-fits-all database.

VoltDB was to be a shared-nothing, distributed OLTP database. Rethinking assumptions about traditional systems, VoltDB threw out shared-memory concurrency, buffer pools and traditional disk persistence, and client-side transaction control. It assumed that high-volume OLTP workloads were mostly horizontally partitionable, and that analytics would migrate to special-purpose systems, keeping queries short.

The proposed system would dramatically reduce scaling issues, support native replication and high availability, and reduce costs for operational workloads without sacrificing transactions and strong consistency.

---

1. For more on H-Store see Chapter 19: H-Store/VoltDB.

2. See <https://dl.acm.org/citation.cfm?id=1454211> for the list of collaborators.

3. For more on Vertica see Chapters 18 and 27.

VoltDB 1.0 was originally released in April 2010, after nearly two years of internal development. Work on the H-Store academic project continued in parallel. Over the years, many ideas and experimental results were shared between the researchers and the VoltDB engineering team, but code diverged as the two systems had different purposes. VoltDB also hired a number of graduate students who worked on the H-Store project.

### **Compaction<sup>4</sup>**

In the Fall of 2010, the very first customer, who was equal parts brave and foolish, was using VoltDB 1.x in production and was running into challenges with memory usage.

This customer was using the resident set size (RSS) for the VoltDB process as reported by the OS as the key metric. While memory usage monitoring is more complex than disk usage monitoring, this is a good metric to use in most cases.

The problem was that the RSS was increasing with use, even though the data was not growing. Yes, records were being updated, deleted, and added, but the total number of records and the size of the logical data they represented was not growing. However, eventually, VoltDB would use all of the memory on the machine. This early customer was forced to restart VoltDB on a periodic basis—not great for a system designed for uptime. Needless to say, this was unacceptable for an in-memory database focused on operational workloads.

The problem was quickly identified as allocator fragmentation. Under it all, VoltDB was using GNU LibC malloc, which allocated big slabs of virtual address space and doled out smaller chunks on request. Allocator fragmentation happens when a slab is logically only half used, but the “holes” that can be used to service new allocations are too small to be useful.

There are two main ways to deal with this problem. The most common approach is to use a custom allocator. The two most common alternatives are JEMalloc and TCMalloc. Both are substantially more sophisticated at avoiding fragmentation waste than the default Glibc malloc.

The VoltDB team tried these options first but ran into challenges because VoltDB mixed C++ and Java in the same process. Using these allocators with the in-process JVM was challenging at the time.

---

4. Compaction, which is critical to running VoltDB for more than a few hours, didn’t come up in the initial design or research because academics don’t always run things the way one might in production. It ended up being critical to success.

The second approach, which is both more challenging and more effective, is do all the allocation yourself. You don't actually have to manage 100% of allocations. Short-lived allocations and permanent allocations tend not to contribute to allocator fragmentation. You primarily have to worry about data with unknown and variable life cycles, which is really critical for any in-memory database.

The team focused on three main types of memory usage that fit this profile.

- Tuple storage—a logical array of fixed size tuples per table.
- Blob storage—a set of variable-sized binary objects linked from tuples.
- Index storage—trees and hash tables that provide fast access to tuples by key.

Two teams set about implementing two different approaches to see which might work best.

The first team took on indexes and blob storage. The plan was to remake these data structures in such a way that they never had any “holes” at all. For indexes, all allocations for a specific index with a specific key width would be done sequentially into a linked list of memory-mapped slabs. Whenever a tree node or hash entry was deleted, the record at the very end of the set of allocations would be moved into the hole, and the pointers in the data structure would be reconfigured for the new address. Blob storage was managed similarly, but with pools for various size blobs.

There was a concern that the extra pointer fixups would impact performance, but measurements showed this was not significant. Now indexes and blobs *could not fragment*. This came at an engineering cost of several engineer-months, but without much performance impact to the product.

Tuple storage took a different approach. Tuples would be allocated into a linked list of memory-mapped slabs, much like index data, but holes from deletion would be tracked, rather than filled. Whenever the number of holes exceeded a threshold (e.g., 5%), a compaction process would be initiated that would rearrange tuples and merge blocks. This would bind fragmentation to a fixed amount, which met the requirements of VoltDB and the customer.

In the end, we didn't pick a winner; we used both schemes in different places. Both prototypes were sufficient and with an early product, there were many other things to improve. The anti-fragmentation work was a huge success and is considered a competitive advantage of VoltDB compared to other in-memory stores that often use memory less efficiently.<sup>5</sup> Without it, it would be hard to use VoltDB in any production workloads.

---

5. The competition catch-up is a long story. Most systems can't do what VoltDB does because they use shared-memory multi-threading and even lock-free or wait-free data structures. These are

These kinds of problems can really illustrate the gulf between research and production.

It turns out compaction is critical to running VoltDB for more than a few hours, but this didn't come up because of the research results. We previously assumed that if a steady state workload worked for an hour, it would work forever, but this is absolutely not the case.

**Lesson.** Memory usage should closely track the actual data stored, and systems should be tested for much longer periods of time.

## Latency

Version 1.0 of the VoltDB database, like the H-Store prototype it was based on, used a transaction ordering and consensus scheme that was based on the ideas described in the original H-Store paper [Stonebraker et al. 2007b], but with additional safety. Oversimplifying a bit, nodes would collect all candidate work in a 5 ms epoch and then exchange between all nodes the work inside the cluster for that 5 ms. This work would then be ordered based on a scheme similar to Twitter Snowflake.<sup>6</sup>

This scheme guaranteed a total, global pre-order for all submitted transactions. That is, before a transaction was run, its serializable order with respect to all other transactions was known.

Compared to contemporary transaction ordering schemes, VoltDB offered more fault tolerance than two-phase-commit and was dramatically simpler than using a schema like Paxos for ordering. It also supported significantly higher throughput than either.

Having a global pre-ordering of all transactions required less coordination between cluster nodes when the work itself was being done [Stonebraker et al. 2007b]. In theory, participants have broad leeway to re-order work, so it can be executed more efficiently, provided it produces results effectively equivalent to the specified order. This was all part of the original H-Store research [Stonebraker et al. 2007b].

So, what's the catch? This scheme used wall clocks to order transactions. That meant transactions must wait up to 5 ms for the epoch to close, plus network round trip time, plus any clock skew. In a single data center, Network Time Protocol (NTP)

---

*much* harder to compact. Other systems *can* use TCMalloc or JEMalloc because they don't embed the JVM.

6. "Announcing Snowflake," the Twitter blog, June 1, 2010. [https://blog.twitter.com/engineering/en\\_us/a/2010/announcing-snowflake.html](https://blog.twitter.com/engineering/en_us/a/2010/announcing-snowflake.html). Last accessed March 29, 2018.

is capable of synchronizing clocks to about 1 ms, but that configuration isn't trivial to get right. Network skew is also typically low but can be affected by common things like background network copies or garbage collections.

To put it more clearly, on a single-node VoltDB instance, client operations would take at least 5 ms even if it did no actual work. That means a synchronous benchmark client could do 200 trivial transactions per second, substantially slower than MySQL for most workloads.

In a cluster, it was worse. Getting NTP set up well in order to evaluate VoltDB was a stumbling block, especially in the new world of the cloud. This meant the delay might be 10–20 ms. The original VoltDB paper assumes achieving clock synchronization is trivial, but we found that to be just false *enough* to cause problems. We didn't just need synced-clocks, we needed them to stay synced for days, months, or even years without issue.

None of this affected throughput. The VoltDB client was fully asynchronous by design and could process responses in the order they arrived. A proper parallel workload could achieve millions of transactions per second on the right cluster, but asking prospective users to build fully asynchronous apps proved too much of a challenge. Users were not used to developing that way and changing user habits is difficult.

VoltDB needed to be faster than MySQL without application wizardry.

Many months of disagreement and thought from the engineering team culminated in a small meeting where a decision had to be made.

A rough plan was hashed out to replace VoltDB consensus with a post-order system that would slash latency to near zero while keeping throughput. The new system would limit some performance improvements to cross-partition transactions (which are typically rare for VoltDB use cases) and it would require several engineers working for almost a year, time that could be spent on more visible features.

Engineering came out of that meeting resolved to fix the latency issues. As part of the plan, the VoltDB 1.0 consensus scheme would be kept, but only to bootstrap a new system of elected partition leaders that serialized all per-partition work and a single, global cross-partition serializer that determined the order of cross-partition work.

This scheme was launched with version 3.0, and average cluster latency was reduced to nearly nothing now that we did not have to hold transactions for clock skew and the all-to-all exchange. Typical response latencies were less than a millisecond with a good network.

This directly led to VoltDB use in low-latency industries like ad-tech and personalization.

**Lesson.** Response time is as important as throughput.

## Disk Persistence

When VoltDB launched, the high-availability story was 100% redundancy through clustering. There were periodic disk snapshots, so you would see data loss only if you lost multiple nodes, and then you might only lose minutes of recent data. The argument was that servers were more reliable, and per-machine UPSs (uninterruptive power supplies) were increasingly common, so multiple failures weren't a likely occurrence.

The argument didn't land.

VoltDB technical marketing and sales spent too much time countering the idea that VoltDB wouldn't keep your data safe. Competitors reinforced this narrative. In early 2011, it got to the point where lack of disk persistence was severely limiting customer growth.

VoltDB needed per-transaction disk persistence without compromising the performance it was known for. Part of the original H-Store/VoltDB thesis was that logging was one of the things holding traditional RDBMSs back when they moved to memory [[Harizopoulos et al. 2008](#)], so this posed quite a challenge.

To address this problem, Engineering added an inter-snapshot log to VoltDB but broke with the ARIES (Algorithms for Recovery and isolation Exploiting Semantics) style logs used by traditional RDBMSs. VoltDB already heavily relied on determinism and logical descriptions of operations to replicate between nodes. Engineering chose to leverage that work to write a logical log to disk that described procedure calls and SQL statements, rather than mutated data.

This approach had a huge technical advantage for VoltDB. As soon as transactions were ordered for a given partition (but before they were executed), they could be written to disk. This meant disk writes *and* the actual computation could be done *simultaneously*. As soon as both were completed, the transaction could be confirmed to the caller. Other systems performed operations and *then* wrote binary change-logs to disk. The logical approach and VoltDB implementation meant disk persistence didn't have substantial impact on throughput, and only minimal impact on latency.

Per-transaction disk-persistence was added in VoltDB 2.5 in Fall 2011 and almost immediately silenced persistence-based criticism of VoltDB. It's clear that without this feature, VoltDB would have seen much more limited use.

As an addendum, we have a lot more data today about how common complete cluster failure is with VoltDB. Cluster failures for well-run VoltDB instances are rare, but not always 100% unavoidable, and not all VoltDB clusters are well run. Disk persistence is a feature that not only cut off a line of criticism, but also gets exercised by users from time to time.

**Lesson.** People don't trust in-memory systems as system of record.

## Latency Redux

In 2013, within a year of reducing average latency in VoltDB to nil, VoltDB was courted by a major telecommunications OEM (original equipment manufacturer) looking to replace Oracle across their stack. Oracle's pricing made it hard for them to compete with upstart Asian vendors who had built their stacks without Oracle, and Oracle's deployment model was poorly suited to virtualization and data-center orchestration.

Replacing Oracle would be a substantial boost to competitiveness.

During the OEM's VoltDB evaluation, latency quickly became an issue. While average latency met requirements, long tail latency did not. For a typical call authorization application, the service level agreement might dictate that any decision not made in 50 ms can't be billed to the customer, forcing the authorization provider to pay the call cost.

VoltDB created a new automated test to measure long tail latency. Rather than measure average latency or measure at the common 99th percentile or even the 99.999th percentile, Engineering set out to specifically count the number of transactions that took longer than 50 ms in a given window. The goal was to reduce that number to zero for a long-term run in our lab so the customer could support P99.999 latency under 50 ms in their deployments.

Once you start measuring the right things, the problem is mostly solved, but there was still code to write. We moved more of the statistics collection and health monitoring code out of blocking paths. We changed how objects were allocated and used to nearly eliminate the need for stop-the-world garbage collection events. We also tuned buffer sizes and Java virtual machine parameters to get everything running nice and "boring."

If there's one thing VoltDB Engineering learned over the course of ten years of development, it's that customers want their operational databases to be as boring and unsurprising as possible. This was the final piece of the puzzle that closed the first major telecommunications customer, with more coming right on their heels.

Today, a significant portion of the world's mobile calls and texts are authorized through a VoltDB-based system.

**Lesson.** P50 is a bad measure—P99 is better—P99.999 is best.

## **Conclusion**

Of course, the incidents described here are just a tiny sliver of the challenges and adventures we encountered building VoltDB into the mature and trusted system it is today. Building a system from a research paper, to a prototype, to a 1.0, and to a robust platform deployed around the world is an unparalleled learning experience.



# The Collected Works of Michael Stonebraker

- D. J. Abadi, D. Carney, U. Çetintemel, M. Cherniack, C. Convey, C. Erwin, E. F. Galvez, M. Hatoun, A. Maskey, A. Rasin, A. Singer, M. Stonebraker, N. Tatbul, Y. Xing, R. Yan, and S. B. Zdonik. 2003a. Aurora: A data stream management system. In *Proc. ACM SIGMOD International Conference on Management of Data*, p. 666. DOI: [10.1145/872757.872855](https://doi.org/10.1145/872757.872855). 225, 228, 229, 230, 232
- D. J. Abadi, D. Carney, U. Çetintemel, M. Cherniack, C. Convey, S. Lee, M. Stonebraker, N. Tatbul, and S. B. Zdonik. 2003b. Aurora: a new model and architecture for data stream management. *VLDB Journal*, 12(2): 120–139. DOI: [10.1007/s00778-003-0095-z](https://doi.org/10.1007/s00778-003-0095-z). 228, 229, 324
- D. J. Abadi, R. Agrawal, A. Ailamaki, M. Balazinska, P. A. Bernstein, M. J. Carey, S. Chaudhuri, J. Dean, A. Doan, M. J. Franklin, J. Gehrke, L. M. Haas, A. Y. Halevy, J. M. Hellerstein, Y. E. Ioannidis, H. V. Jagadish, D. Kossmann, S. Madden, S. Mehrotra, T. Milo, J. F. Naughton, R. Ramakrishnan, V. Markl, C. Olston, B. C. Ooi, C. Ré, D. Suciu, M. Stonebraker, T. Walter, and J. Widom. 2014. The Beckman report on database research. *ACM SIGMOD Record*, 43(3): 61–70. DOI: [10.1145/2694428.2694441](https://doi.org/10.1145/2694428.2694441). 92
- D. Abadi, R. Agrawal, A. Ailamaki, M. Balazinska, P. A. Bernstein, M. J. Carey, S. Chaudhuri, J. Dean, A. Doan, M. J. Franklin, J. Gehrke, L. M. Haas, A. Y. Halevy, J. M. Hellerstein, Y. E. Ioannidis, H. V. Jagadish, D. Kossmann, S. Madden, S. Mehrotra, T. Milo, J. F. Naughton, R. Ramakrishnan, V. Markl, C. Olston, B. C. Ooi, C. Ré, D. Suciu, M. Stonebraker, T. Walter, and J. Widom. 2016. The Beckman report on database research. *Communications of the ACM*, 59(2): 92–99. DOI: [10.1145/2845915](https://doi.org/10.1145/2845915). 92
- Z. Abedjan, C. G. Akcora, M. Ouzzani, P. Papotti, and M. Stonebraker. 2015a. Temporal rules discovery for web data cleaning. *Proc. VLDB Endowment*, 9(4): 336–347. <http://www.vldb.org/pvldb/vol9/p336-abedjan.pdf>. 297
- Z. Abedjan, J. Morcos, M. N. Gubanov, I. F. Ilyas, M. Stonebraker, P. Papotti, and M. Ouzzani. 2015b. Dataxformer: Leveraging the web for semantic transformations. In *Proc. 7th Biennial Conference on Innovative Data Systems Research*. [http://www.cidrdb.org/cidr2015/Papers/CIDR15\\_Paper31.pdf](http://www.cidrdb.org/cidr2015/Papers/CIDR15_Paper31.pdf). 296, 297

- Z. Abedjan, X. Chu, D. Deng, R. C. Fernandez, I. F. Ilyas, M. Ouzzani, P. Papotti, M. Stonebraker, and N. Tang. 2016a. Detecting data errors: Where are we and what needs to be done? *Proc. VLDB Endowment*, 9(12): 993–1004. <http://www.vldb.org/pvldb/vol9/p993-abedjan.pdf>. 298
- Z. Abedjan, J. Morcos, I. F. Ilyas, M. Ouzzani, P. Papotti, and M. Stonebraker. 2016b. Dataxformer: A robust transformation discovery system. In *Proc. 32nd International Conference on Data Engineering*, pp. 1134–1145. DOI: [10.1109/ICDE.2016.7498319](https://doi.org/10.1109/ICDE.2016.7498319). 296
- S. Abiteboul, R. Agrawal, P. A. Bernstein, M. J. Carey, S. Ceri, W. B. Croft, D. J. DeWitt, M. J. Franklin, H. Garcia-Molina, D. Gawlick, J. Gray, L. M. Haas, A. Y. Halevy, J. M. Hellerstein, Y. E. Ioannidis, M. L. Kersten, M. J. Pazzani, M. Lesk, D. Maier, J. F. Naughton, H. Schek, T. K. Sellis, A. Silberschatz, M. Stonebraker, R. T. Snodgrass, J. D. Ullman, G. Weikum, J. Widom, and S. B. Zdonik. 2003. The Lowell database research self assessment. *CoRR*, cs.DB/0310006. <http://arxiv.org/abs/cs.DB/0310006>. 92
- S. Abiteboul, R. Agrawal, P. A. Bernstein, M. J. Carey, S. Ceri, W. B. Croft, D. J. DeWitt, M. J. Franklin, H. Garcia-Molina, D. Gawlick, J. Gray, L. M. Haas, A. Y. Halevy, J. M. Hellerstein, Y. E. Ioannidis, M. L. Kersten, M. J. Pazzani, M. Lesk, D. Maier, J. F. Naughton, H. Schek, T. K. Sellis, A. Silberschatz, M. Stonebraker, R. T. Snodgrass, J. D. Ullman, G. Weikum, J. Widom, and S. B. Zdonik. 2005. The Lowell database research self-assessment. *Communications of the ACM*, 48(5): 111–118. DOI: [10.1145/1060710.1060718](https://doi.org/10.1145/1060710.1060718). 92
- R. Agrawal, A. Ailamaki, P. A. Bernstein, E. A. Brewer, M. J. Carey, S. Chaudhuri, A. Doan, D. Florescu, M. J. Franklin, H. Garcia-Molina, J. Gehrke, L. Gruenwald, L. M. Haas, A. Y. Halevy, J. M. Hellerstein, Y. E. Ioannidis, H. F. Korth, D. Kossmann, S. Madden, R. Magoulas, B. C. Ooi, T. O'Reilly, R. Ramakrishnan, S. Sarawagi, M. Stonebraker, A. S. Szalay, and G. Weikum. 2008. The Claremont report on database research. *ACM SIGMOD Record*, 37(3): 9–19. DOI: [10.1145/1462571.1462573](https://doi.org/10.1145/1462571.1462573). 92
- R. Agrawal, A. Ailamaki, P. A. Bernstein, E. A. Brewer, M. J. Carey, S. Chaudhuri, A. Doan, D. Florescu, M. J. Franklin, H. Garcia-Molina, J. Gehrke, L. Gruenwald, L. M. Haas, A. Y. Halevy, J. M. Hellerstein, Y. E. Ioannidis, H. F. Korth, D. Kossmann, S. Madden, R. Magoulas, B. C. Ooi, T. O'Reilly, R. Ramakrishnan, S. Sarawagi, M. Stonebraker, A. S. Szalay, and G. Weikum. 2009. The Claremont report on database research. *Communications of the ACM*, 52(6): 56–65. DOI: [10.1145/1516046.1516062](https://doi.org/10.1145/1516046.1516062). 92
- A. Aiken, J. Chen, M. Lin, M. Spalding, M. Stonebraker, and A. Woodruff. 1995. The Tioga-2 database visualization environment. In *Proc. Workshop on Database Issues for Data Visualization*, pp. 181–207. DOI: [10.1007/3-540-62221-7\\_15](https://doi.org/10.1007/3-540-62221-7_15).
- A. Aiken, J. Chen, M. Stonebraker, and A. Woodruff. 1996. Tioga-2: A direct manipulation database visualization environment. In *Proc. 12th International Conference on Data Engineering*, pp. 208–217. DOI: [10.1109/ICDE.1996.492109](https://doi.org/10.1109/ICDE.1996.492109).
- E. Allman and M. Stonebraker. 1982. Observations on the evolution of a software system. *IEEE Computer*, 15(6): 27–32. DOI: [10.1109/MC.1982.1654047](https://doi.org/10.1109/MC.1982.1654047).

- E. Allman, M. Stonebraker, and G. Held. 1976. Embedding a relational data sublanguage in a general purpose programming language. In *Proc. SIGPLAN Conference on Data: Abstraction, Definition and Structure*, pp. 25–35. DOI: [10.1145/800237.807115](https://doi.org/10.1145/800237.807115). 195
- J. T. Anderson and M. Stonebraker. 1994. SEQUOIA 2000 metadata schema for satellite images. *ACM SIGMOD Record*, 23(4): 42–48. DOI: [10.1145/190627.190642](https://doi.org/10.1145/190627.190642).
- A. Arasu, M. Cherniack, E. F. Galvez, D. Maier, A. Maskey, E. Ryvkina, M. Stonebraker, and R. Tibbetts. 2004. Linear road: A stream data management benchmark. In *Proc. 30th International Conference on Very Large Data Bases*, pp. 480–491. <http://www.vldb.org/conf/2004/RS12P1.pdf>. 326
- T. Atwoode, J. Dash, J. Stein, M. Stonebraker, and M. E. S. Loomis. 1994. Objects and databases (panel). In *Proc. 9th Annual Conference on Object-Oriented Programming Systems, Languages, and Applications*, pp. 371–372. DOI: [10.1145/191080.191138](https://doi.org/10.1145/191080.191138).
- H. Balakrishnan, M. Balazinska, D. Carney, U. Çetintemel, M. Cherniack, C. Convey, E. F. Galvez, J. Salz, M. Stonebraker, N. Tatbul, R. Tibbetts, and S. B. Zdonik. 2004. Retrospective on Aurora. *VLDB Journal*, 13(4): 370–383. DOI: [10.1007/s00778-004-0133-5](https://doi.org/10.1007/s00778-004-0133-5). 228, 229
- M. Balazinska, H. Balakrishnan, and M. Stonebraker. 2004b. Load management and high availability in the Medusa distributed stream processing system. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 929–930. DOI: [10.1145/1007568.1007701](https://doi.org/10.1145/1007568.1007701). 325
- M. Balazinska, H. Balakrishnan, and M. Stonebraker. 2004a. Contract-based load management in federated distributed systems. In *Proc. 1st USENIX Symposium on Networked Systems Design and Implementation*. <http://www.usenix.org/events/nsdi04/tech/balazinska.html>. 228, 230
- M. Balazinska, H. Balakrishnan, S. Madden, and M. Stonebraker. 2005. Fault-tolerance in the Borealis distributed stream processing system. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 13–24. DOI: [10.1145/1066157.1066160](https://doi.org/10.1145/1066157.1066160). 228, 230, 234, 325
- M. Balazinska, H. Balakrishnan, S. Madden, and M. Stonebraker. 2008. Fault-tolerance in the borealis distributed stream processing system. *ACM Transactions on Database Systems*, 33(1): 3:1–3:44. DOI: [10.1145/1331904.1331907](https://doi.org/10.1145/1331904.1331907).
- D. Barbará, J. A. Blakeley, D. H. Fishman, D. B. Lomet, and M. Stonebraker. 1994. The impact of database research on industrial products (panel summary). *ACM SIGMOD Record*, 23(3): 35–40. DOI: [10.1145/187436.187455](https://doi.org/10.1145/187436.187455).
- V. Barr and M. Stonebraker. 2015a. A valuable lesson, and whither hadoop? *Communications of the ACM*, 58(1): 18–19. DOI: [10.1145/2686591](https://doi.org/10.1145/2686591). 50
- V. Barr and M. Stonebraker. 2015b. How men can help women in cs; winning 'computing's nobel prize'. *Communications of the ACM*, 58(11): 10–11. DOI: [10.1145/2820419](https://doi.org/10.1145/2820419).
- V. Barr, M. Stonebraker, R. C. Fernandez, D. Deng, and M. L. Brodie. 2017. How we teach cs2all, and what to do about database decay. *Communications of the ACM*, 60(1): 10–11. <http://dl.acm.org/citation.cfm?id=3014349>.

- L. Battle, M. Stonebraker, and R. Chang. 2013. Dynamic reduction of query result sets for interactive visualization. In *Proc. 2013 IEEE International Conference on Big Data*, pp. 1–8. DOI: [10.1109/BigData.2013.6691708](https://doi.org/10.1109/BigData.2013.6691708).
- L. Battle, R. Chang, and M. Stonebraker. 2016. Dynamic prefetching of data tiles for interactive visualization. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 1363–1375. DOI: [10.1145/2882903.2882919](https://doi.org/10.1145/2882903.2882919).
- R. Berman and M. Stonebraker. 1977. GEO-OUEL: a system for the manipulation and display of geographic data. In *Proc. 4th Annual Conference Computer Graphics and Interactive Techniques*, pp. 186–191. DOI: [10.1145/563858.563892](https://doi.org/10.1145/563858.563892).
- P. A. Bernstein, U. Dayal, D. J. DeWitt, D. Gawlick, J. Gray, M. Jarke, B. G. Lindsay, P. C. Lockemann, D. Maier, E. J. Neuhold, A. Reuter, L. A. Rowe, H. Schek, J. W. Schmidt, M. Schrefl, and M. Stonebraker. 1989. Future directions in DBMS research—the Laguna Beach participants. *ACM SIGMOD Record*, 18(1): 17–26. [92](#)
- P. A. Bernstein, M. L. Brodie, S. Ceri, D. J. DeWitt, M. J. Franklin, H. Garcia-Molina, J. Gray, G. Held, J. M. Hellerstein, H. V. Jagadish, M. Lesk, D. Maier, J. F. Naughton, H. Pirahesh, M. Stonebraker, and J. D. Ullman. 1998a. The Asilomar report on database research. *ACM SIGMOD Record*, 27(4): 74–80. DOI: [10.1145/306101.306137](https://doi.org/10.1145/306101.306137). [92](#)
- P. A. Bernstein, M. L. Brodie, S. Ceri, D. J. DeWitt, M. J. Franklin, H. Garcia-Molina, J. Gray, G. Held, J. M. Hellerstein, H. V. Jagadish, M. Lesk, D. Maier, J. F. Naughton, H. Pirahesh, M. Stonebraker, and J. D. Ullman. 1998b. The Asilomar report on database research. *CoRR*, cs.DB/9811013. <http://arxiv.org/abs/cs.DB/9811013>. [92](#)
- A. Bhide and M. Stonebraker. 1987. Performance issues in high performance transaction processing architectures. In *Proc. 2nd International Workshop High Performance Transaction Systems*, pp. 277–300. DOI: [10.1007/3-540-51085-0\\_51](https://doi.org/10.1007/3-540-51085-0_51). [91](#)
- A. Bhide and M. Stonebraker. 1988. A performance comparison of two architectures for fast transaction processing. In *Proc. 4th International Conference on Data Engineering*, pp. 536–545. DOI: [10.1109/ICDE.1988.105501](https://doi.org/10.1109/ICDE.1988.105501). [91](#)
- M. L. Brodie and M. Stonebraker. 1993. Darwin: On the incremental migration of legacy information systems. Technical Report TR-0222-10-92-165, GTE Laboratories Incorporated.
- M. L. Brodie and M. Stonebraker. 1995a. *Migrating Legacy Systems: Gateways, Interfaces, and the Incremental Approach*. Morgan Kaufmann. [91](#)
- M. L. Brodie and M. Stonebraker. 1995b. *Legacy Information Systems Migration: Gateways, Interfaces, and the Incremental Approach*. Morgan Kaufmann.
- M. L. Brodie, R. M. Michael Stonebraker, and J. Pei. 2018. The case for the co-evolution of applications and data. In *New England Database Days*.
- P. Brown and M. Stonebraker. 1995. Bigsur: A system for the management of earth science data. In *Proc. 21th International Conference on Very Large Data Bases*, pp. 720–728. <http://www.vldb.org/conf/1995/P720.pdf>.

- M. J. Carey and M. Stonebraker. 1984. The performance of concurrency control algorithms for database management systems. In *Proc. 10th International Conference on Very Large Data Bases*, pp. 107–118. <http://www.vldb.org/conf/1984/P107.pdf>. 91, 200
- D. Carney, U. Çetintemel, M. Cherniack, C. Convey, S. Lee, G. Seidman, M. Stonebraker, N. Tatbul, and S. B. Zdonik. 2002. Monitoring streams—A new class of data management applications. In *Proc. 28th International Conference on Very Large Data Bases*, pp. 215–226. DOI: [10.1016/B978-155860869-6/50027-5](https://doi.org/10.1016/B978-155860869-6/50027-5). 228, 229, 324
- D. Carney, U. Çetintemel, A. Rasin, S. B. Zdonik, M. Cherniack, and M. Stonebraker. 2003. Operator scheduling in a data stream manager. In *Proc. 29th International Conference on Very Large Data Bases*, pp. 838–849. <http://www.vldb.org/conf/2003/papers/S25P02.pdf>. 228, 229
- U. Çetintemel, J. Du, T. Kraska, S. Madden, D. Maier, J. Meehan, A. Pavlo, M. Stonebraker, E. Sutherland, N. Tatbul, K. Tufte, H. Wang, and S. B. Zdonik. 2014. S-store: A streaming NewSQL system for big velocity applications. *Proc. VLDB Endowment*, 7(13): 1633–1636. <http://www.vldb.org/pvldb/vol7/p1633-cetintemel.pdf>. 234, 251
- U. Çetintemel, D. J. Abadi, Y. Ahmad, H. Balakrishnan, M. Balazinska, M. Cherniack, J. Hwang, S. Madden, A. Maskey, A. Rasin, E. Ryvkina, M. Stonebraker, N. Tatbul, Y. Xing, and S. Zdonik. 2016. The Aurora and Borealis stream processing engines. In M. N. Garofalakis, J. Gehrke, and R. Rastogi, editors, *Data Stream Management—Processing High-Speed Data Streams*, pp. 337–359. Springer. ISBN 978-3-540-28607-3. DOI: [10.1007/978-3-540-28608-0\\_17](https://doi.org/10.1007/978-3-540-28608-0_17).
- R. Chandra, A. Segev, and M. Stonebraker. 1994. Implementing calendars and temporal rules in next generation databases. In *Proc. 10th International Conference on Data Engineering*, pp. 264–273. DOI: [10.1109/ICDE.1994.283040](https://doi.org/10.1109/ICDE.1994.283040). 91
- S. Chaudhuri, A. K. Chandra, U. Dayal, J. Gray, M. Stonebraker, G. Wiederhold, and M. Y. Vardi. 1996. Database research: Lead, follow, or get out of the way?—panel abstract. In *Proc. 12th International Conference on Data Engineering*, p. 190.
- P. Chen, V. Gadepally, and M. Stonebraker. 2016. The BigDAWG monitoring framework. In *Proc. 2016 IEEE High Performance Extreme Computing Conference*, pp. 1–6. DOI: [10.1109/HPEC.2016.7761642](https://doi.org/10.1109/HPEC.2016.7761642). 373
- Y. Chi, C. R. Mechoso, M. Stonebraker, K. Sklower, R. Troy, R. R. Muntz, and E. Mesrobian. 1997. ESMDIS: earth system model data information system. In *Proc. 9th International Conference on Scientific and Statistical Database Management*, pp. 116–118. DOI: [10.1109/SSDM.1997.621169](https://doi.org/10.1109/SSDM.1997.621169).
- P. Cudré-Mauroux, H. Kimura, K. Lim, J. Rogers, R. Simakov, E. Soroush, P. Velikhov, D. L. Wang, M. Balazinska, J. Becla, D. J. DeWitt, B. Heath, D. Maier, S. Madden, J. M. Patel, M. Stonebraker, and S. B. Zdonik. 2009. A demonstration of SciDB: A science-oriented DBMS. *Proc. VLDB Endowment*, 2(2): 1534–1537. DOI: [10.14778/1687553.1687584](https://doi.org/10.14778/1687553.1687584).
- J. DeBrabant, A. Pavlo, S. Tu, M. Stonebraker, and S. B. Zdonik. 2013. Anti-caching: A new approach to database management system architecture. *Proc. VLDB Endowment*, 6(14): 1942–1953. <http://www.vldb.org/pvldb/vol6/p1942-debrabant.pdf>.

- J. DeBrabant, J. Arulraj, A. Pavlo, M. Stonebraker, S. B. Zdonik, and S. Dulloor. 2014. A prolegomenon on OLTP database systems for non-volatile memory. In *Proc. 5th International Workshop on Accelerating Data Management Systems Using Modern Processor and Storage Architectures*, pp. 57–63. [http://www.adms-conf.org/2014/adms14\\_debrabant.pdf](http://www.adms-conf.org/2014/adms14_debrabant.pdf).
- D. Deng, R. C. Fernandez, Z. Abedjan, S. Wang, M. Stonebraker, A. K. Elmagarmid, I. F. Ilyas, S. Madden, M. Ouzzani, and N. Tang. 2017a. The data civilizer system. In *Proc. 8th Biennial Conference on Innovative Data Systems Research*. <http://cidrdb.org/cidr2017/papers/p44-deng-cidr17.pdf>. 293
- D. Deng, A. Kim, S. Madden, and M. Stonebraker. 2017b. SILKMOTH: an efficient method for finding related sets with maximum matching constraints. *CoRR*, abs/1704.04738. <http://arxiv.org/abs/1704.04738>.
- D. J. DeWitt, R. H. Katz, F. Olken, L. D. Shapiro, M. Stonebraker, and D. A. Wood. 1984. Implementation techniques for main memory database systems. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 1–8. DOI: [10.1145/602259.602261](https://doi.org/10.1145/602259.602261). 111
- D. J. DeWitt and M. Stonebraker. January 2008. MapReduce: A major step backwards. *The Database Column*. [http://homes.cs.washington.edu/~billhowe/mapreduce\\_a\\_major\\_step\\_backwards.html](http://homes.cs.washington.edu/~billhowe/mapreduce_a_major_step_backwards.html). Accessed April 8, 2018. 50, 114, 136, 184, 209
- D. J. DeWitt, I. F. Ilyas, J. F. Naughton, and M. Stonebraker. 2013. We are drowning in a sea of least publishable units (lpus). In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 921–922. DOI: [10.1145/2463676.2465345](https://doi.org/10.1145/2463676.2465345).
- P. Dobbins, T. Dohzen, C. Grant, J. Hammer, M. Jones, D. Oliver, M. Pamuk, J. Shin, and M. Stonebraker. 2007. Morpheus 2.0: A data transformation management system. In *Proc. 3rd International Workshop on Database Interoperability*.
- T. Dohzen, M. Pamuk, J. Hammer, and M. Stonebraker. 2006. Data integration through transform reuse in the Morpheus project. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 736–738. DOI: [10.1145/1142473.1142571](https://doi.org/10.1145/1142473.1142571).
- J. Dozier, M. Stonebraker, and J. Frew. 1994. Sequoia 2000: A next-generation information system for the study of global change. In *Proc. 13th IEEE Symposium Mass Storage Systems*, pp. 47–56. DOI: [10.1109/MASS.1994.373028](https://doi.org/10.1109/MASS.1994.373028).
- J. Duggan and M. Stonebraker. 2014. Incremental elasticity for array databases. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 409–420. DOI: [10.1145/2588555.2588569](https://doi.org/10.1145/2588555.2588569).
- J. Duggan, A. J. Elmore, M. Stonebraker, M. Balazinska, B. Howe, J. Kepner, S. Madden, D. Maier, T. Mattson, and S. B. Zdonik. 2015a. The BigDAWG polystore system. *ACM SIGMOD Record*, 44(2): 11–16. DOI: [10.1145/2814710.2814713](https://doi.org/10.1145/2814710.2814713). 284
- J. Duggan, O. Papaemmanouil, L. Battle, and M. Stonebraker. 2015b. Skew-aware join optimization for array databases. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 123–135. DOI: [10.1145/2723372.2723709](https://doi.org/10.1145/2723372.2723709).

- A. Dziedzic, J. Duggan, A. J. Elmore, V. Gadepally, and M. Stonebraker. 2015. BigDAWG: a polystore for diverse interactive applications. *Data Systems for Interactive Analysis Workshop*.
- A. Dziedzic, A. J. Elmore, and M. Stonebraker. 2016. Data transformation and migration in polystores. In *Proc. 2016 IEEE High Performance Extreme Computing Conference*, pp. 1–6. DOI: [10.1109/HPEC.2016.7761594](https://doi.org/10.1109/HPEC.2016.7761594). 372
- A. J. Elmore, J. Duggan, M. Stonebraker, M. Balazinska, U. Çetintemel, V. Gadepally, J. Heer, B. Howe, J. Kepner, T. Kraska, S. Madden, D. Maier, T. G. Mattson, S. Papadopoulos, J. Parkhurst, N. Tatbul, M. Vartak, and S. Zdonik. 2015. A demonstration of the BigDAWG polystore system. *Proc. VLDB Endowment*, 8(12): 1908–1911. <http://www.vldb.org/pvldb/vol8/p1908-Elmore.pdf>. 287, 371
- R. S. Epstein and M. Stonebraker. 1980. Analysis of distributed data base processing strategies. In *Proc. 6th International Conference on Very Data Bases*, pp. 92–101.
- R. S. Epstein, M. Stonebraker, and E. Wong. 1978. Distributed query processing in a relational data base system. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 169–180. DOI: [10.1145/509252.509292](https://doi.org/10.1145/509252.509292). 198
- R. C. Fernandez, Z. Abedjan, S. Madden, and M. Stonebraker. 2016. Towards large-scale data discovery: position paper. In *Proc. 3rd International Workshop on Exploratory Search in Databases and the Web*, pp. 3–5. DOI: [10.1145/2948674.2948675](https://doi.org/10.1145/2948674.2948675).
- R. C. Fernandez, D. Deng, E. Mansour, A. A. Qahtan, W. Tao, Z. Abedjan, A. K. Elmagarmid, I. F. Ilyas, S. Madden, M. Ouzzani, M. Stonebraker, and N. Tang. 2017b. A demo of the data civilizer system. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 1639–1642. DOI: [10.1145/3035918.3058740](https://doi.org/10.1145/3035918.3058740).
- R. C. Fernandez, D. Deng, E. Mansour, A. A. Qahtan, W. Tao, Z. Abedjan, A. Elmagarmid, I. F. Ilyas, S. Madden, M. Ouzzani, M. Stonebraker, and N. Tang. 2017a. A demo of the data civilizer system. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 1636–1642. 293
- R. C. Fernandez, Z. Abedjan, F. Koko, G. Yuan, S. Madden, and M. Stonebraker. 2018a. Aurum: A data discovery system. In *Proc. 34th International Conference on Data Engineering*, pp. 1001–1012.
- R. C. Fernandez, E. Mansour, A. Qahtan, A. Elmagarmid, I. Ilyas, S. Madden, M. Ouzzani, M. Stonebraker, and N. Tang. 2018b. Seeping semantics: Linking datasets using word embeddings for data discovery. In *Proc. 34th International Conference on Data Engineering*, pp. 989–1000.
- V. Gadepally, P. Chen, J. Duggan, A. J. Elmore, B. Haynes, J. Kepner, S. Madden, T. Mattson, and M. Stonebraker. 2016a. The BigDAWG polystore system and architecture. In *Proc. 2016 IEEE High Performance Extreme Computing Conference*, pp. 1–6. DOI: [10.1109/HPEC.2016.7761636](https://doi.org/10.1109/HPEC.2016.7761636). 287, 373
- V. Gadepally, P. Chen, J. Duggan, A. J. Elmore, B. Haynes, J. Kepner, S. Madden, T. Mattson, and M. Stonebraker. 2016b. The BigDAWG polystore system and architecture. *CoRR*, abs/1609.07548. <http://arxiv.org/abs/1609.07548>.



- V. Gadepally, P. Chen, J. Duggan, A. Elmore, B. Haynes, J. Kepner, S. Madden, T. Mattson, and M. Stonebraker. 2016c. The BigDAWG polystore system and architecture. In *Proc. 2016 IEEE High Performance Extreme Computing Conference*, pp. 1–6. DOI: [10.1109/HPEC.2016.7761636](https://doi.org/10.1109/HPEC.2016.7761636).
- V. Gadepally, J. Duggan, A. J. Elmore, J. Kepner, S. Madden, T. Mattson, and M. Stonebraker. 2016d. The BigDAWG architecture. *CoRR*, abs/1602.08791. <http://arxiv.org/abs/1602.08791>.
- A. Go, M. Stonebraker, and C. Williams. 1975. An approach to implementing a geo-data system. In *Proc. Workshop on Data Bases for Interactive Design*, pp. 67–77.
- J. Gray, H. Schek, M. Stonebraker, and J. D. Ullman. 2003. The Lowell report. In *Proc. ACM SIGMOD International Conference on Management of Data*, p. 680. DOI: [10.1145/872757.872873](https://doi.org/10.1145/872757.872873). 92
- M. N. Gubanov and M. Stonebraker. 2013. Bootstrapping synonym resolution at web scale. In *Proc. DIMACS/CCICADA Workshop on Big Data Integration*.
- M. N. Gubanov and M. Stonebraker. 2014. Large-scale semantic profile extraction. In *Proc. 17th International Conference on Extending Database Technology*, pp. 644–647. DOI: [10.5441/002/edbt.2014.64](https://doi.org/10.5441/002/edbt.2014.64).
- M. N. Gubanov, M. Stonebraker, and D. Bruckner. 2014. Text and structured data fusion in data tamer at scale. In *Proc. 30th International Conference on Data Engineering*, pp. 1258–1261. DOI: [10.1109/ICDE.2014.6816755](https://doi.org/10.1109/ICDE.2014.6816755).
- A. M. Gupta, V. Gadepally, and M. Stonebraker. 2016. Cross-engine query execution in federated database systems. In *Proc. 2016 IEEE High Performance Extreme Computing Conference*, pp. 1–6. DOI: [10.1109/HPEC.2016.7761648](https://doi.org/10.1109/HPEC.2016.7761648). 373
- A. Guttman and M. Stonebraker. 1982. Using a relational database management system for computer aided design data. *Quarterly Bulletin IEEE Technical Committee on Data Engineering*, 5(2): 21–28. <http://sites.computer.org/debull/82JUN-CD.pdf>. 201
- J. Hammer, M. Stonebraker, and O. Topsakal. 2005. THALIA: test harness for the assessment of legacy information integration approaches. In *Proc. 21st International Conference on Data Engineering*, pp. 485–486. DOI: [10.1109/ICDE.2005.140](https://doi.org/10.1109/ICDE.2005.140).
- R. Harding, D. V. Aken, A. Pavlo, and M. Stonebraker. 2017. An evaluation of distributed concurrency control. *Proc. VLDB Endowment*, 10(5): 553–564. DOI: [10.14778/3055540.3055548](https://doi.org/10.14778/3055540.3055548).
- S. Harizopoulos, D. J. Abadi, S. Madden, and M. Stonebraker. 2008. OLTP through the looking glass, and what we found there. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 981–992. DOI: [10.1145/1376616.1376713](https://doi.org/10.1145/1376616.1376713). 152, 246, 251, 346
- P. B. Hawthorn and M. Stonebraker. 1979. Performance analysis of a relational data base management system. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 1–12. DOI: [10.1145/582095.582097](https://doi.org/10.1145/582095.582097).
- G. Held and M. Stonebraker. 1975. Storage structures and access methods in the relational data base management system INGRES. In *Proc. ACM Pacific 75—Data: Its Use, Organization and Management*, pp. 26–33. 194



- G. Held and M. Stonebraker. 1978. B-trees re-examined. *Communications of the ACM*, 21(2): 139–143. DOI: [10.1145/359340.359348](https://doi.org/10.1145/359340.359348). 90, 197
- G. Held, M. Stonebraker, and E. Wong. 1975. INGRES: A relational data base system. In *National Computer Conference*, pp. 409–416. DOI: [10.1145/1499949.1500029](https://doi.org/10.1145/1499949.1500029). 102, 397
- J. M. Hellerstein and M. Stonebraker. 1993. Predicate migration: Optimizing queries with expensive predicates. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 267–276. DOI: [10.1145/170035.170078](https://doi.org/10.1145/170035.170078).
- J. M. Hellerstein and M. Stonebraker. 2005. *Readings in Database Systems*, 4. MIT Press. ISBN 978-0-262-69314-1. <http://mitpress.mit.edu/books/readings-database-systems>.
- J. M. Hellerstein, M. Stonebraker, and R. Caccia. 1999. Independent, open enterprise data integration. *Quarterly Bulletin IEEE Technical Committee on Data Engineering*, 22(1): 43–49. <http://sites.computer.org/debull/99mar/cohera.ps>.
- J. M. Hellerstein, M. Stonebraker, and J. R. Hamilton. 2007. Architecture of a database system. *Foundations and Trends in Databases*, 1(2): 141–259. DOI: [10.1561/19000000002](https://doi.org/10.1561/19000000002).
- W. Hong and M. Stonebraker. 1991. Optimization of parallel query execution plans in XPRS. In *Proc. 1st International Conference on Parallel and Distributed Information Systems*, pp. 218–225. DOI: [10.1109/PDIS.1991.183106](https://doi.org/10.1109/PDIS.1991.183106).
- W. Hong and M. Stonebraker. 1993. Optimization of parallel query execution plans in XPRS. *Distributed and Parallel Databases*, 1(1): 9–32. DOI: [10.1007/BF01277518](https://doi.org/10.1007/BF01277518).
- J. Hwang, M. Balazinska, A. Rasin, U. Çetintemel, M. Stonebraker, and S. B. Zdonik. 2005. High-availability algorithms for distributed stream processing. In *Proc. 21st International Conference on Data Engineering*, pp. 779–790. DOI: [10.1109/ICDE.2005.72](https://doi.org/10.1109/ICDE.2005.72). 228, 230, 325
- A. Jhingran and M. Stonebraker. 1990. Alternatives in complex object representation: A performance perspective. In *Proc. 6th International Conference on Data Engineering*, pp. 94–102. DOI: [10.1109/ICDE.1990.113458](https://doi.org/10.1109/ICDE.1990.113458).
- A. Jindal, P. Rawlani, E. Wu, S. Madden, A. Deshpande, and M. Stonebraker. 2014. VERTEXICA: your relational friend for graph analytics! *Proc. VLDB Endowment*, 7(13): 1669–1672. <http://www.vldb.org/pvldb/vol7/p1669-jindal.pdf>.
- R. Kallman, H. Kimura, J. Natkins, A. Pavlo, A. Rasin, S. B. Zdonik, E. P. C. Jones, S. Madden, M. Stonebraker, Y. Zhang, J. Hugg, and D. J. Abadi. 2008. H-store: a high-performance, distributed main memory transaction processing system. *Proc. VLDB Endowment*, 1(2): 1496–1499. DOI: [10.14778/1454159.1454211](https://doi.org/10.14778/1454159.1454211). 247, 249, 341
- R. H. Katz, J. K. Ousterhout, D. A. Patterson, and M. Stonebraker. 1988. A project on high performance I/O subsystems. *Quarterly Bulletin IEEE Technical Committee on Data Engineering*, 11(1): 40–47. <http://sites.computer.org/debull/88MAR-CD.pdf>.
- J. T. Kohl, C. Staelin, and M. Stonebraker. 1993a. Highlight: Using a log-structured file system for tertiary storage management. In *Proc. of the Usenix Winter 1993 Technical Conference*, pp. 435–448.

- J. T. Kohl, M. Stonebraker, and C. Staelin. 1993b. Highlight: a file system for tertiary storage. In *Proc. 12th IEEE Symposium Mass Storage Systems*, pp. 157–161. DOI: [10.1109/MASS.1993.289765](https://doi.org/10.1109/MASS.1993.289765).
- C. P. Kolovson and M. Stonebraker. 1989. Indexing techniques for historical databases. In *Proc. 5th International Conference on Data Engineering*, pp. 127–137. DOI: [10.1109/ICDE.1989.47208](https://doi.org/10.1109/ICDE.1989.47208).
- C. P. Kolovson and M. Stonebraker. 1991. Segment indexes: Dynamic indexing techniques for multi-dimensional interval data. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 138–147. DOI: [10.1145/115790.115807](https://doi.org/10.1145/115790.115807).
- R. A. Kowalski, D. B. Lenat, E. Soloway, M. Stonebraker, and A. Walker. 1988. Knowledge management—panel report. In *Proc. 2nd International Conference on Expert Database Systems*, pp. 63–69.
- A. Kumar and M. Stonebraker. 1987a. The effect of join selectivities on optimal nesting order. *ACM SIGMOD Record*, 16(1): 28–41. DOI: [10.1145/24820.24822](https://doi.org/10.1145/24820.24822).
- A. Kumar and M. Stonebraker. 1987b. Performance evaluation of an operating system transaction manager. In *Proc. 13th International Conference on Very Large Data Bases*, pp. 473–481. <http://www.vldb.org/conf/1987/P473.pdf>.
- A. Kumar and M. Stonebraker. 1988. Semantics based transaction management techniques for replicated data. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 117–125. DOI: [10.1145/50202.50215](https://doi.org/10.1145/50202.50215).
- A. Kumar and M. Stonebraker. 1989. Performance considerations for an operating system transaction manager. *IEEE Transactions on Software Engineering*, 15(6): 705–714. DOI: [10.1109/32.24724](https://doi.org/10.1109/32.24724).
- R. Kung, E. N. Hanson, Y. E. Ioannidis, T. K. Sellis, L. D. Shapiro, and M. Stonebraker. 1984. Heuristic search in data base systems. In *Proc. 1st International Workshop on Expert Database Systems*, pp. 537–548.
- C. A. Lynch and M. Stonebraker. 1988. Extended user-defined indexing with application to textual databases. In *Proc. 14th International Conference on Very Large Data Bases*, pp. 306–317. <http://www.vldb.org/conf/1988/P306.pdf>.
- N. Malviya, A. Weisberg, S. Madden, and M. Stonebraker. 2014. Rethinking main memory OLTP recovery. In *Proc. 30th International Conference on Data Engineering*, pp. 604–615. DOI: [10.1109/ICDE.2014.6816685](https://doi.org/10.1109/ICDE.2014.6816685).
- E. Mansour, D. Deng, A. Qahtan, R. C. Fernandez, Wenbo, Z. Abedjan, A. Elmagarmid, I. Ilyas, S. Madden, M. Ouzzani, M. Stonebraker, and N. Tang. 2018. Building data civilizer pipelines with an advanced workflow engine. In *Proc. 34th International Conference on Data Engineering*, pp. 1593–1596.
- T. Mattson, D. A. Bader, J. W. Berry, A. Buluç, J. Dongarra, C. Faloutsos, J. Feo, J. R. Gilbert, J. Gonzalez, B. Hendrickson, J. Kepner, C. E. Leiserson, A. Lumsdaine, D. A. Padua, S. Poole, S. P. Reinhardt, M. Stonebraker, S. Wallach, and A. Yoo. 2013. Standards for graph algorithm primitives. In *Proc. 2013 IEEE High Performance Extreme Computing Conference*, pp. 1–2. DOI: [10.1109/HPEC.2013.6670338](https://doi.org/10.1109/HPEC.2013.6670338).

- T. Mattson, D. A. Bader, J. W. Berry, A. Buluç, J. J. Dongarra, C. Faloutsos, J. Feo, J. R. Gilbert, J. Gonzalez, B. Hendrickson, J. Kepner, C. E. Leiserson, A. Lumsdaine, D. A. Padua, S. W. Poole, S. P. Reinhardt, M. Stonebraker, S. Wallach, and A. Yoo. 2014. Standards for graph algorithm primitives. *CoRR*, abs/1408.0393. DOI: [10.1109/HPEC.2013.6670338](https://doi.org/10.1109/HPEC.2013.6670338).
- N. H. McDonald and M. Stonebraker. 1975. CUPID - the friendly query language. In *Proc. ACM Pacific 75—Data: Its Use, Organization and Management*, pp. 127–131.
- J. Meehan, N. Tatbul, S. B. Zdonik, C. Aslantas, U. Çetintemel, J. Du, T. Kraska, S. Madden, D. Maier, A. Pavlo, M. Stonebraker, K. Tufte, and H. Wang. 2015a. S-store: Streaming meets transaction processing. *CoRR*, abs/1503.01143. DOI: [10.14778/2831360.2831367](https://doi.org/10.14778/2831360.2831367). 234
- J. Meehan, N. Tatbul, S. Zdonik, C. Aslantas, U. Çetintemel, J. Du, T. Kraska, S. Madden, D. Maier, A. Pavlo, M. Stonebraker, K. Tufte, and H. Wang. 2015b. S-store: Streaming meets transaction processing. *Proc. VLDB Endowment*, 8(13): 2134–2145. DOI: [10.14778/2831360.2831367](https://doi.org/10.14778/2831360.2831367). 234, 288, 331, 374
- J. Morcos, Z. Abedjan, I. F. Ilyas, M. Ouzzani, P. Papotti, and M. Stonebraker. 2015. Dataxformer: An interactive data transformation tool. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 883–888. DOI: [10.1145/2723372.2735366](https://doi.org/10.1145/2723372.2735366). 296
- B. Muthuswamy, L. Kerschberg, C. Zaniolo, M. Stonebraker, D. S. P. Jr., and M. Jarke. 1985. Architectures for expert-DBMS (panel). In *Proc. 1985 ACM Annual Conference on the Range of Computing: Mid-80's*, pp. 424–426. DOI: [10.1145/320435.320555](https://doi.org/10.1145/320435.320555).
- K. O'Brien, V. Gadepally, J. Duggan, A. Dziedzic, A. J. Elmore, J. Kepner, S. Madden, T. Mattson, Z. She, and M. Stonebraker. 2017. BigDAWG polystore release and demonstration. *CoRR*, abs/1701.05799. <http://arxiv.org/abs/1701.05799>.
- V. E. Ogle and M. Stonebraker. 1995. Chabot: Retrieval from a relational database of images. *IEEE Computer*, 28(9): 40–48. DOI: [10.1109/2.410150](https://doi.org/10.1109/2.410150).
- M. A. Olson, W. Hong, M. Ubell, and M. Stonebraker. 1996. Query processing in a parallel object-relational database system. *Quarterly Bulletin IEEE Technical Committee on Data Engineering*, 19(4): 3–10. <http://sites.computer.org/debull/96DEC-CD.pdf>.
- C. Olston, M. Stonebraker, A. Aiken, and J. M. Hellerstein. 1998a. VIQING: visual interactive querying. In *Proc. 1998 IEEE Symposium on Visual Languages*, pp. 162–169. DOI: [10.1109/VL.1998.706159](https://doi.org/10.1109/VL.1998.706159).
- C. Olston, A. Woodruff, A. Aiken, M. Chu, V. Ercegovac, M. Lin, M. Spalding, and M. Stonebraker. 1998b. Datasplash. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 550–552. DOI: [10.1145/276304.276377](https://doi.org/10.1145/276304.276377).
- J. Ong, D. Fogg, and M. Stonebraker. 1984. Implementation of data abstraction in the relational database system ingres. *ACM SIGMOD Record*, 14(1): 1–14. DOI: [10.1145/984540.984541](https://doi.org/10.1145/984540.984541). 201, 202, 206
- R. Overmyer and M. Stonebraker. 1982. Implementation of a time expert in a data base system. *ACM SIGMOD Record*, 12(3): 51–60. DOI: [10.1145/984505.984509](https://doi.org/10.1145/984505.984509).

- A. Pavlo, E. Paulson, A. Rasin, D. J. Abadi, D. J. DeWitt, S. Madden, and M. Stonebraker. 2009. A comparison of approaches to large-scale data analysis. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 165–178. DOI: [10.1145/1559845.1559865](https://doi.org/10.1145/1559845.1559865).
- H. Pirk, S. Madden, and M. Stonebraker. 2015. By their fruits shall ye know them: A data analyst’s perspective on massively parallel system design. In *Proc. 11th Workshop on Data Management on New Hardware*, pp. 5:1–5:6. DOI: [10.1145/2771937.2771944](https://doi.org/10.1145/2771937.2771944).
- G. Planthaber, M. Stonebraker, and J. Frew. 2012. Earthdb: scalable analysis of MODIS data using SciDB. In *Proc. 1st ACM SIGSPATIAL International Workshop on Analytics for Big Geospatial Data*, pp. 11–19. DOI: [10.1145/2447481.2447483](https://doi.org/10.1145/2447481.2447483).
- S. Potamianos and M. Stonebraker. 1996. The POSTGRES rules system. In *Active Database Systems: Triggers and Rules For Advanced Database Processing*, pp. 43–61. Morgan Kaufmann. 91, 168
- C. Ré, D. Agrawal, M. Balazinska, M. Cafarella, M. Jordan, T. Kraska, and R. Ramakrishnan. 2015. Machine learning and databases: The sound of things to come or a cacophony of hype? In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 283–284.
- D. R. Ries and M. Stonebraker. 1977a. Effects of locking granularity in a database management system. *ACM Transactions on Database Systems*, 2(3): 233–246. DOI: [10.1145/320557.320566](https://doi.org/10.1145/320557.320566). 91, 198
- D. R. Ries and M. Stonebraker. 1977b. A study of the effects of locking granularity in a database management system (abstract). In *Proc. ACM SIGMOD International Conference on Management of Data*, p. 121. DOI: [10.1145/509404.509422](https://doi.org/10.1145/509404.509422). 91
- D. R. Ries and M. Stonebraker. 1979. Locking granularity revisited. *ACM Transactions on Database Systems*, 4(2): 210–227. <http://doi.acm.org/10.1145/320071.320078>. DOI: [10.1145/320071.320078](https://doi.org/10.1145/320071.320078). 91
- L. A. Rowe and M. Stonebraker. 1981. Architecture of future data base systems. *ACM SIGMOD Record*, 11(1): 30–44. DOI: [10.1145/984471.984473](https://doi.org/10.1145/984471.984473).
- L. A. Rowe and M. Stonebraker. 1986. The commercial INGRES epilogue. In M. Stonebraker, editor, *The INGRES Papers: Anatomy of a Relational Database System*, pp. 63–82. Addison-Wesley.
- L. A. Rowe and M. Stonebraker. 1987. The POSTGRES data model. In *Proc. 13th International Conference on Very Large Data Bases*, pp. 83–96. <http://www.vldb.org/conf/1987/P083.pdf>. 258
- L. A. Rowe and M. Stonebraker. 1990. The POSTGRES data model. In A. F. Cardenas and D. McLeod, editors, *Research Foundations in Object-Oriented and Semantic Database Systems*, pp. 91–110. Prentice Hall.
- S. Sarawagi and M. Stonebraker. 1994. Efficient organization of large multidimensional arrays. In *Proc. 10th International Conference on Data Engineering*, pp. 328–336. DOI: [10.1109/ICDE.1994.283048](https://doi.org/10.1109/ICDE.1994.283048).

- S. Sarawagi and M. Stonebraker. 1996. Reordering query execution in tertiary memory databases. In *Proc. 22th International Conference on Very Large Data Bases*. <http://www.vldb.org/conf/1996/P156.pdf>.
- G. A. Schloss and M. Stonebraker. 1990. Highly redundant management of distributed data. In *Proc. Workshop on the Management of Replicated Data*, pp. 91–92.
- A. Seering, P. Cudré-Mauroux, S. Madden, and M. Stonebraker. 2012. Efficient versioning for scientific array databases. In *Proc. 28th International Conference on Data Engineering*, pp. 1013–1024. DOI: [10.1109/ICDE.2012.102](https://doi.org/10.1109/ICDE.2012.102).
- L. J. Seligman, N. J. Belkin, E. J. Neuhold, M. Stonebraker, and G. Wiederhold. 1995. Metrics for accessing heterogeneous data: Is there any hope? (panel). In *Proc. 21th International Conference on Very Large Data Bases*, p. 633. <http://www.vldb.org/conf/1995/P633.pdf>.
- M. I. Seltzer and M. Stonebraker. 1990. Transaction support in read optimized and write optimized file systems. In *Proc. 16th International Conference on Very Large Data Bases*, pp. 174–185. <http://www.vldb.org/conf/1990/P174.pdf>.
- M. I. Seltzer and M. Stonebraker. 1991. Read optimized file system designs: A performance evaluation. In *Proc. 7th International Conference on Data Engineering*, pp. 602–611. DOI: [10.1109/ICDE.1991.131509](https://doi.org/10.1109/ICDE.1991.131509).
- M. Serafini, R. Taft, A. J. Elmore, A. Pavlo, A. Aboulmaga, and M. Stonebraker. 2016. Clay: Fine-grained adaptive partitioning for general database schemas. *Proc. VLDB Endowment*, 10(4): 445–456. DOI: [10.14778/3025111.3025125](https://doi.org/10.14778/3025111.3025125).
- J. Sidell, P. M. Aoki, A. Sah, C. Staelin, M. Stonebraker, and A. Yu. 1996. Data replication in mariposa. In *Proc. 12th International Conference on Data Engineering*, pp. 485–494. DOI: [10.1109/ICDE.1996.492198](https://doi.org/10.1109/ICDE.1996.492198).
- A. Silberschatz, M. Stonebraker, and J. D. Ullman. 1990. Database systems: Achievements and opportunities—the “Lagunita” report of the NSF invitational workshop on the future of database system research held in Palo Alto, CA, February 22–23, 1990. *ACM SIGMOD Record*, 19(4): 6–22. DOI: [10.1145/122058.122059](https://doi.org/10.1145/122058.122059). 92
- A. Silberschatz, M. Stonebraker, and J. D. Ullman. 1991. Database systems: Achievements and opportunities. *Communications of the ACM*, 34(10): 110–120. DOI: [10.1145/125223.125272](https://doi.org/10.1145/125223.125272).
- A. Silberschatz, M. Stonebraker, and J. D. Ullman. 1996. Database research: Achievements and opportunities into the 21st century. *ACM SIGMOD Record*, 25(1): 52–63. DOI: [10.1145/381854.381886](https://doi.org/10.1145/381854.381886).
- D. Skeen and M. Stonebraker. 1981. A formal model of crash recovery in a distributed system. In *Proc. 5th Berkeley Workshop on Distributed Data Management and Computer Networks*, pp. 129–142.
- D. Skeen and M. Stonebraker. 1983. A formal model of crash recovery in a distributed system. *IEEE Transactions on Software Engineering*, 9(3): 219–228. DOI: [10.1109/TSE.1983.236608](https://doi.org/10.1109/TSE.1983.236608). 199

- M. Stonebraker and R. Cattell. 2011. 10 rules for scalable performance in “simple operation” datastores. *Communications of the ACM*, 54(6): 72–80. DOI: [10.1145/1953122.1953144](https://doi.org/10.1145/1953122.1953144).
- M. Stonebraker and U. Çetintemel. 2005. “One size fits all”: An idea whose time has come and gone (abstract). In *Proc. 21st International Conference on Data Engineering*, pp. 2–11. DOI: [10.1109/ICDE.2005.1](https://doi.org/10.1109/ICDE.2005.1). [50](#), [92](#), [103](#), [131](#), [152](#), [367](#), [401](#)
- M. Stonebraker and D. J. DeWitt. 2008. A tribute to Jim Gray. *Communications of the ACM*, 51(11): 54–57. DOI: [10.1145/1400214.1400230](https://doi.org/10.1145/1400214.1400230).
- M. Stonebraker and A. Guttman. 1984. Using a relational database management system for computer aided design data—an update. *Quarterly Bulletin IEEE Technical Committee on Data Engineering*, 7(2): 56–60. <http://sites.computer.org/debull/84JUN-CD.pdf>.
- M. Stonebraker and A. Guttman. 1984. R-trees: a dynamic index structure for spatial searching. In *Proc. of the 1984 ACM SIGMOD International Conference on Management of Data (SIGMOD '84)*, pp. 47–57. ACM, New York. DOI: [10.1145/602259.602266](https://doi.org/10.1145/602259.602266). [201](#)
- M. Stonebraker and M. A. Hearst. 1988. Future trends in expert data base systems. In *Proc. 2nd International Conference on Expert Database Systems*, pp. 3–20. [395](#)
- M. Stonebraker and G. Held. 1975. Networks, hierarchies and relations in data base management systems. In *Proc. ACM Pacific 75—Data: Its Use, Organization and Management*, pp. 1–9.
- M. Stonebraker and J. M. Hellerstein, editors. 1998. *Readings in Database Systems*, 3. Morgan Kaufmann.
- M. Stonebraker and J. M. Hellerstein. 2001. Content integration for e-business. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 552–560. DOI: [10.1145/375663.375739](https://doi.org/10.1145/375663.375739).
- M. Stonebraker and J. Hong. 2009. Saying good-bye to DBMSS, designing effective interfaces. *Communications of the ACM*, 52(9): 12–13. DOI: [10.1145/1562164.1562169](https://doi.org/10.1145/1562164.1562169).
- M. Stonebraker and J. Hong. 2012. Researchers’ big data crisis; understanding design and functionality. *Communications of the ACM*, 55(2): 10–11. DOI: [10.1145/2076450.2076453](https://doi.org/10.1145/2076450.2076453).
- M. Stonebraker and J. Kalash. 1982. TIMBER: A sophisticated relation browser (invited paper). In *Proc. 8th International Conference on Very Data Bases*, pp. 1–10. <http://www.vldb.org/conf/1982/P001.pdf>.
- M. Stonebraker and K. Keller. 1980. Embedding expert knowledge and hypothetical data bases into a data base system. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 58–66. DOI: [10.1145/582250.582261](https://doi.org/10.1145/582250.582261). [200](#)
- M. Stonebraker and G. Kemnitz. 1991. The Postgres next generation database management system. *Communications of the ACM*, 34(10): 78–92. DOI: [10.1145/125223.125262](https://doi.org/10.1145/125223.125262). [168](#), [206](#), [213](#)
- M. Stonebraker and A. Kumar. 1986. Operating system support for data management. *Quarterly Bulletin IEEE Technical Committee on Data Engineering*, 9(3): 43–50. <http://sites.computer.org/debull/86SEP-CD.pdf>. [47](#)



- M. Stonebraker and D. Moore. 1996. *Object-Relational DBMSs: The Next Great Wave*. Morgan Kaufmann. 111
- M. Stonebraker and E. J. Neuhold. 1977. A distributed database version of INGRES. In *Proc. 2nd Berkeley Workshop on Distributed Data Management and Computer Networks*, pp. 19–36. 109, 198, 199
- M. Stonebraker and M. A. Olson. 1993. Large object support in POSTGRES. In *Proc. 9th International Conference on Data Engineering*, pp. 355–362. DOI: [10.1109/ICDE.1993.344046](https://doi.org/10.1109/ICDE.1993.344046).
- M. Stonebraker and J. Robertson. 2013. Big data is “buzzword du jour,” CS academics “have the best job”. *Communications of the ACM*, 56(9): 10–11. DOI: [10.1145/2500468.2500471](https://doi.org/10.1145/2500468.2500471).
- M. Stonebraker and L. A. Rowe. 1977. Observations on data manipulation languages and their embedding in general purpose programming languages. In *Proc. 3rd International Conference on Very Data Bases*, pp. 128–143.
- M. Stonebraker and L. A. Rowe. 1984. Database portals: A new application program interface. In *Proc. 10th International Conference on Very Large Data Bases*, pp. 3–13. <http://www.vldb.org/conf/1984/P003.pdf>.
- M. Stonebraker and L. A. Rowe. 1986. The design of Postgres. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 340–355. DOI: [10.1145/16894.16888](https://doi.org/10.1145/16894.16888). 149, 203, 206
- M. Stonebraker and P. Rubinstein. 1976. The INGRES protection system. In *Proc. 1976 ACM Annual Conference*, pp. 80–84. DOI: [10.1145/800191.805536](https://doi.org/10.1145/800191.805536). 398
- M. Stonebraker and G. A. Schloss. 1990. Distributed RAID—A new multiple copy algorithm. In *Proc. 6th International Conference on Data Engineering*, pp. 430–437. DOI: [10.1109/ICDE.1990.113496](https://doi.org/10.1109/ICDE.1990.113496).
- M. Stonebraker and A. Weisberg. 2013. The VoltDB main memory DBMS. *Quarterly Bulletin IEEE Technical Committee on Data Engineering*, 36(2): 21–27. <http://sites.computer.org/debull/A13june/VoltDB1.pdf>.
- M. Stonebraker and E. Wong. 1974b. Access control in a relational data base management system by query modification. In *Proc. 1974 ACM Annual Conference, Volume 1*, pp. 180–186. DOI: [10.1145/800182.810400](https://doi.org/10.1145/800182.810400). 45
- M. Stonebraker, P. Rubinstein, R. Conway, D. Strip, H. R. Hartson, D. K. Hsiao, and E. B. Fernandez. 1976a. SIGBDP (paper session). In *Proc. 1976 ACM Annual Conference*, p. 79. DOI: [10.1145/800191.805535](https://doi.org/10.1145/800191.805535).
- M. Stonebraker, E. Wong, P. Kreps, and G. Held. 1976b. The design and implementation of INGRES. *ACM Transactions on Database Systems*, 1(3): 189–222. DOI: [10.1145/320473.320476](https://doi.org/10.1145/320473.320476). 47, 148, 398
- M. Stonebraker, R. R. Johnson, and S. Rosenberg. 1982a. A rules system for a relational data base management system. In *Proc. 2nd International Conference on Databases: Improving Database Usability and Responsiveness*, pp. 323–335. 91, 202

- M. Stonebraker, J. Woodfill, J. Ranstrom, M. C. Murphy, J. Kalash, M. J. Carey, and K. Arnold. 1982b. Performance analysis of distributed data base systems. *Quarterly Bulletin IEEE Technical Committee on Data Engineering*, 5(4): 58–65. <http://sites.computer.org/debull/82DEC-CD.pdf>.
- M. Stonebraker, W. B. Rubenstein, and A. Guttman. 1983a. Application of abstract data types and abstract indices to CAD data bases. In *Engineering Design Applications*, pp. 107–113.
- M. Stonebraker, H. Stettner, N. Lynn, J. Kalash, and A. Guttman. 1983b. Document processing in a relational database system. *ACM Transactions on Information Systems*, 1(2): 143–158. DOI: [10.1145/357431.357433](https://doi.org/10.1145/357431.357433).
- M. Stonebraker, J. Woodfill, and E. Andersen. 1983c. Implementation of rules in relational data base systems. *Quarterly Bulletin IEEE Technical Committee on Data Engineering*, 6(4): 65–74. <http://sites.computer.org/debull/83DEC-CD.pdf>. 91, 202
- M. Stonebraker, J. Woodfill, J. Ranstrom, J. Kalash, K. Arnold, and E. Andersen. 1983d. Performance analysis of distributed data base systems. In *Proc. 2nd Symposium on Reliable Distributed Systems*, pp. 135–138.
- M. Stonebraker, J. Woodfill, J. Ranstrom, M. C. Murphy, M. Meyer, and E. Allman. 1983e. Performance enhancements to a relational database system. *ACM Transactions on Database Systems*, 8(2): 167–185. DOI: [10.1145/319983.319984](https://doi.org/10.1145/319983.319984).
- M. Stonebraker, E. Anderson, E. N. Hanson, and W. B. Rubenstein. 1984a. Quel as a data type. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 208–214. DOI: [10.1145/602259.602287](https://doi.org/10.1145/602259.602287). 208
- M. Stonebraker, J. Woodfill, J. Ranstrom, J. Kalash, K. Arnold, and E. Andersen. 1984b. Performance analysis of distributed data base systems. *Performance Evaluation*, 4(3): 220. DOI: [10.1016/0166-5316\(84\)90036-1](https://doi.org/10.1016/0166-5316(84)90036-1).
- M. Stonebraker, D. DuBourdieu, and W. Edwards. 1985. Problems in supporting data base transactions in an operating system transaction manager. *Operating Systems Review*, 19(1): 6–14. DOI: [10.1145/1041490.1041491](https://doi.org/10.1145/1041490.1041491).
- M. Stonebraker, T. K. Sellis, and E. N. Hanson. 1986. An analysis of rule indexing implementations in data base systems. In *Proc. 1st International Conference on Expert Database Systems*, pp. 465–476. 91
- M. Stonebraker, J. Anton, and E. N. Hanson. 1987a. Extending a database system with procedures. *ACM Transactions on Database Systems*, 12(3): 350–376. DOI: [10.1145/27629.27631](https://doi.org/10.1145/27629.27631).
- M. Stonebraker, J. Anton, and M. Hirohama. 1987b. Extendability in POSTGRES. *Quarterly Bulletin IEEE Technical Committee on Data Engineering*, 10(2): 16–23. <http://sites.computer.org/debull/87JUN-CD.pdf>.
- M. Stonebraker, E. N. Hanson, and C. Hong. 1987c. The design of the postgres rules system. In *Proc. 3th International Conference on Data Engineering*, pp. 365–374. DOI: [10.1109/ICDE.1987.7272402](https://doi.org/10.1109/ICDE.1987.7272402). 91



- M. Stonebraker, E. N. Hanson, and S. Potamianos. 1988a. The POSTGRES rule manager. *IEEE Transactions on Software Engineering*, 14(7): 897–907. DOI: [10.1109/32.42733](https://doi.org/10.1109/32.42733). 91, 168
- M. Stonebraker, R. H. Katz, D. A. Patterson, and J. K. Ousterhout. 1988b. The design of XPRS. In *Proc. 14th International Conference on Very Large Data Bases*, pp. 318–330. <http://www.vldb.org/conf/1988/P318.pdf>.
- M. Stonebraker, M. A. Hearst, and S. Potamianos. 1989. A commentary on the POSTGRES rule system. *ACM SIGMOD Record*, 18(3): 5–11. DOI: [10.1145/71031.71032](https://doi.org/10.1145/71031.71032). 91, 168, 395
- M. Stonebraker, A. Jhingran, J. Goh, and S. Potamianos. 1990a. On rules, procedures, caching and views in data base systems. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 281–290. DOI: [10.1145/93597.98737](https://doi.org/10.1145/93597.98737).
- M. Stonebraker, L. A. Rowe, and M. Hirohama. 1990b. The implementation of postgres. *IEEE Transactions on Knowledge and Data Engineering*, 2(1): 125–142. DOI: [10.1109/69.50912](https://doi.org/10.1109/69.50912). 47, 168
- M. Stonebraker, L. A. Rowe, B. G. Lindsay, J. Gray, M. J. Carey, and D. Beech. 1990e. Third generation data base system manifesto—The committee for advanced DBMS function. In *Proc. ACM SIGMOD International Conference on Management of Data*, p. 396.
- M. Stonebraker, L. A. Rowe, B. G. Lindsay, J. Gray, M. J. Carey, M. L. Brodie, P. A. Bernstein, and D. Beech. 1990c. Third-generation database system manifesto—The committee for advanced DBMS function. *ACM SIGMOD Record*, 19(3): 31–44. DOI: [10.1145/101077.390001](https://doi.org/10.1145/101077.390001). 91
- M. Stonebraker, L. A. Rowe, B. G. Lindsay, J. Gray, M. J. Carey, M. L. Brodie, P. A. Bernstein, and D. Beech. 1990d. Third-generation database system manifesto—The committee for advanced DBMS function. In *Proc. IFIP TC2/WG 2.6 Working Conference on Object-Oriented Databases: Analysis, Design & Construction*, pp. 495–511. 91
- M. Stonebraker, R. Agrawal, U. Dayal, E. J. Neuhold, and A. Reuter. 1993a. DBMS research at a crossroads: The Vienna update. In *Proc. 19th International Conference on Very Large Data Bases*, pp. 688–692. <http://www.vldb.org/conf/1993/P688.pdf>.
- M. Stonebraker, J. Chen, N. Nathan, C. Paxson, A. Su, and J. Wu. 1993b. Tioga: A database-oriented visualization tool. In *Proceedings IEEE Conference Visualization*, pp. 86–93. DOI: [10.1109/VISUAL.1993.398855](https://doi.org/10.1109/VISUAL.1993.398855). 393
- M. Stonebraker, J. Chen, N. Nathan, C. Paxson, and J. Wu. 1993c. Tioga: Providing data management support for scientific visualization applications. In *Proc. 19th International Conference on Very Large Data Bases*, pp. 25–38. <http://www.vldb.org/conf/1993/P025.pdf>. 393
- M. Stonebraker, J. Frew, and J. Dozier. 1993d. The SEQUOIA 2000 project. In *Proc. 3rd International Symposium Advances in Spatial Databases*, pp. 397–412. DOI: [10.1007/3-540-56869-7\\_22](https://doi.org/10.1007/3-540-56869-7_22).

- M. Stonebraker, J. Frew, K. Gardels, and J. Meredith. 1993e. The Sequoia 2000 benchmark. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 2–11. DOI: [10.1145/170035.170038](https://doi.org/10.1145/170035.170038).
- M. Stonebraker, P. M. Aoki, R. Devine, W. Litwin, and M. A. Olson. 1994a. Mariposa: A new architecture for distributed data. In *Proc. 10th International Conference on Data Engineering*, pp. 54–65. DOI: [10.1109/ICDE.1994.283004](https://doi.org/10.1109/ICDE.1994.283004). 401
- M. Stonebraker, R. Devine, M. Kornacker, W. Litwin, A. Pfeffer, A. Sah, and C. Staelin. 1994b. An economic paradigm for query processing and data migration in Mariposa. In *Proc. 3rd International Conference on Parallel and Distributed Information Systems*, pp. 58–67. DOI: [10.1109/PDIS.1994.331732](https://doi.org/10.1109/PDIS.1994.331732).
- M. Stonebraker, P. M. Aoki, W. Litwin, A. Pfeffer, A. Sah, J. Sidell, C. Staelin, and A. Yu. 1996. Mariposa: A wide-area distributed database system. *VLDB Journal*, 5(1): 48–63. DOI: [10.1007/s007780050015](https://doi.org/10.1007/s007780050015).
- M. Stonebraker, P. Brown, and M. Herbach. 1998a. Interoperability, distributed applications and distributed databases: The virtual table interface. *Quarterly Bulletin IEEE Technical Committee on Data Engineering*, 21(3): 25–33. <http://sites.computer.org/debull/98sept/informix.ps>.
- M. Stonebraker, P. Brown, and D. Moore. 1998b. *Object-Relational DBMSs*, 2. Morgan Kaufmann.
- M. Stonebraker, D. J. Abadi, A. Batkin, X. Chen, M. Cherniack, M. Ferreira, E. Lau, A. Lin, S. Madden, E. J. O’Neil, P. E. O’Neil, A. Rasin, N. Tran, and S. B. Zdonik. 2005a. C-store: A column-oriented DBMS. In *Proc. 31st International Conference on Very Large Data Bases*, pp. 553–564. <http://www.vldb2005.org/program/paper/thu/p553-stonebraker.pdf>. 104, 132, 151, 238, 242, 258, 333, 335, 402
- M. Stonebraker, U. Çetintemel, and S. B. Zdonik. 2005b. The 8 requirements of real-time stream processing. *ACM SIGMOD Record*, 34(4): 42–47. DOI: [10.1145/1107499.1107504](https://doi.org/10.1145/1107499.1107504). 282
- M. Stonebraker, C. Bear, U. Çetintemel, M. Cherniack, T. Ge, N. Hachem, S. Harizopoulos, J. Lifter, J. Rogers, and S. B. Zdonik. 2007a. One size fits all? Part 2: Benchmarking studies. In *Proc. 3rd Biennial Conference on Innovative Data Systems Research*, pp. 173–184. <http://www.cidrdb.org/cidr2007/papers/cidr07p20.pdf>. 103, 282
- M. Stonebraker, S. Madden, D. J. Abadi, S. Harizopoulos, N. Hachem, and P. Helland. 2007b. The end of an architectural era (it’s time for a complete rewrite). In *Proc. 33rd International Conference on Very Large Data Bases*, pp. 1150–1160. <http://www.vldb.org/conf/2007/papers/industrial/p1150-stonebraker.pdf>. 247, 341, 344
- M. Stonebraker, J. Becla, D. J. DeWitt, K. Lim, D. Maier, O. Ratzesberger, and S. B. Zdonik. 2009. Requirements for science data bases and SciDB. In *Proc. 4th Biennial Conference on Innovative Data Systems Research*. [http://www-db.cs.wisc.edu/cidr/cidr2009/Paper\\_26.pdf](http://www-db.cs.wisc.edu/cidr/cidr2009/Paper_26.pdf). 257

- M. Stonebraker, D. J. Abadi, D. J. DeWitt, S. Madden, E. Paulson, A. Pavlo, and A. Rasin. 2010. Mapreduce and parallel DBMS: friends or foes? *Communications of the ACM*, 53(1): 64–71. DOI: [10.1145/1629175.1629197](https://doi.org/10.1145/1629175.1629197). 50, 136, 251
- M. Stonebraker, P. Brown, A. Poliakov, and S. Raman. 2011. The architecture of SciDB. In *Proc. 23rd International Conference on Scientific and Statistical Database Management*, pp. 1–16. DOI: [10.1007/978-3-642-22351-8\\_1](https://doi.org/10.1007/978-3-642-22351-8_1).
- M. Stonebraker, A. Ailamaki, J. Kepner, and A. S. Szalay. 2012. The future of scientific data bases. In *Proc. 28th International Conference on Data Engineering*, pp. 7–8. DOI: [10.1109/ICDE.2012.151](https://doi.org/10.1109/ICDE.2012.151).
- M. Stonebraker, P. Brown, D. Zhang, and J. Becla. 2013a. SciDB: A database management system for applications with complex analytics. *Computing in Science and Engineering*, 15(3): 54–62. DOI: [10.1109/MCSE.2013.19](https://doi.org/10.1109/MCSE.2013.19).
- M. Stonebraker, D. Bruckner, I. F. Ilyas, G. Beskales, M. Cherniack, S. B. Zdonik, A. Pagan, and S. Xu. 2013b. Data curation at scale: The data tamer system. In *Proc. 6th Biennial Conference on Innovative Data Systems Research*. [http://www.cidrdb.org/cidr2013/Papers/CIDR13\\_Paper28.pdf](http://www.cidrdb.org/cidr2013/Papers/CIDR13_Paper28.pdf). 105, 150, 269, 297, 357, 358
- M. Stonebraker, J. Duggan, L. Battle, and O. Papaemmanouil. 2013c. SciDB DBMS research at M.I.T. *Quarterly Bulletin IEEE Technical Committee on Data Engineering*, 36(4): 21–30. <http://sites.computer.org/debull/A13dec/p21.pdf>.
- M. Stonebraker, S. Madden, and P. Dubey. 2013d. Intel “big data” science and technology center vision and execution plan. *ACM SIGMOD Record*, 42(1): 44–49. DOI: [10.1145/2481528.2481537](https://doi.org/10.1145/2481528.2481537).
- M. Stonebraker, A. Pavlo, R. Taft, and M. L. Brodie. 2014. Enterprise database applications and the cloud: A difficult road ahead. In *Proc. 7th IEEE International Conference on Cloud Computing*, pp. 1–6. DOI: [10.1109/IC2E.2014.97](https://doi.org/10.1109/IC2E.2014.97).
- M. Stonebraker, D. Deng, and M. L. Brodie. 2016. Database decay and how to avoid it. In *Proc. 2016 IEEE International Conference on Big Data*, pp. 7–16. DOI: [10.1109/BigData.2016.7840584](https://doi.org/10.1109/BigData.2016.7840584).
- M. Stonebraker, D. Deng, and M. L. Brodie. 2017. Application-database co-evolution: A new design and development paradigm. In *North East Database Day*, pp. 1–3.
- M. Stonebraker. 1972a. Retrieval efficiency using combined indexes. In *Proc. 1972 ACM-SIGFIDET Workshop on Data Description, Access and Control*, pp. 243–256.
- M. Stonebraker. 1972b. A simplification of forrester's model of an urban area. *IEEE Transactions on Systems, Man, and Cybernetics*, 2(4): 468–472. DOI: [10.1109/TSMC.1972.4309156](https://doi.org/10.1109/TSMC.1972.4309156).
- M. Stonebraker. 1974a. The choice of partial inversions and combined indices. *International Journal Parallel Programming*, 3(2): 167–188. DOI: [10.1007/BF00976642](https://doi.org/10.1007/BF00976642).
- M. Stonebraker. 1974b. A functional view of data independence. In *Proc. 1974 ACM SIGMOD Workshop on Data Description, Access and Control*, pp. 63–81. DOI: [10.1145/800296.811505](https://doi.org/10.1145/800296.811505). 404, 405

- M. Stonebraker. 1975. Getting started in INGRES—A tutorial, Memorandum No. ERL-M518, Electronics Research Laboratory, College of Engineering, UC Berkeley. [196](#)
- M. Stonebraker. 1975. Implementation of integrity constraints and views by query modification. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 65–78. DOI: [10.1145/500080.500091](#). [45](#), [90](#)
- M. Stonebraker. 1976a. Proposal for a network INGRES. In *Proc. 1st Berkeley Workshop on Distributed Data Management and Computer Networks*, p. 132.
- M. Stonebraker. 1976b. The data base management system INGRES. In *Proc. 1st Berkeley Workshop on Distributed Data Management and Computer Networks*, p. 336. [195](#)
- M. Stonebraker. 1976c. A comparison of the use of links and secondary indices in a relational data base system. In *Proc. 2nd International Conference on Software Engineering*, pp. 527–531. <http://dl.acm.org/citation.cfm?id=807727>.
- M. Stonebraker. 1978. Concurrency control and consistency of multiple copies of data in distributed INGRES. In *Proc. 3rd Berkeley Workshop on Distributed Data Management and Computer Networks*, pp. 235–258. [90](#), [398](#)
- M. Stonebraker. May 1979a. Muffin: A distributed database machine. Technical Report ERL Technical Report UCB/ERL M79/28, University of California at Berkeley. [151](#)
- M. Stonebraker. 1979b. Concurrency control and consistency of multiple copies of data in distributed INGRES. *IEEE Transactions on Software Engineering*, 5(3): 188–194. DOI: [10.1109/TSE.1979.234180](#). [398](#)
- M. Stonebraker. 1980. Retrospection on a database system. *ACM Transactions on Database Systems*, 5(2): 225–240. DOI: [10.1145/320141.320158](#).
- M. Stonebraker. 1981a. Operating system support for database management. *Communications of the ACM*, 24(7): 412–418. DOI: [10.1145/358699.358703](#).
- M. Stonebraker. 1981b. Chairman's column. *ACM SIGMOD Record*, 11(3): i–iv.
- M. Stonebraker. 1981c. Chairman's column. *ACM SIGMOD Record*, 11(4): 2–4.
- M. Stonebraker. 1981d. Chairman's column. *ACM SIGMOD Record*, 12(1): 1–3.
- M. Stonebraker. 1981e. In memory of Kevin Whitney. *ACM SIGMOD Record*, 12(1): 7.
- M. Stonebraker. 1981f. Chairman's column. *ACM SIGMOD Record*, 11(1): 1–4.
- M. Stonebraker. 1981g. Hypothetical data bases as views. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 224–229. DOI: [10.1145/582318.582352](#).
- M. Stonebraker. 1982a. Chairman's column. *ACM SIGMOD Record*, 12(3): 2–4.
- M. Stonebraker. 1982b. Letter to Peter Denning (two VLDB conferences). *ACM SIGMOD Record*, 12(3): 6–7.
- M. Stonebraker. 1982c. Chairman's column. *ACM SIGMOD Record*, 12(4): a–c.
- M. Stonebraker. 1982d. Chairman's column. *ACM SIGMOD Record*, 13(1): 2–3&4.
- M. Stonebraker. 1982e. Adding semantic knowledge to a relational database system. In M. L. Brodie, M. John, and S. J. W., editors, *On Conceptual Modelling*, pp. 333–352. Springer. DOI: [10.1007/978-1-4612-5196-5\\_12](#).

- M. Stonebraker. 1982f. A database perspective. In M. L. Brodie, M. John, and S. J. W., editors, *On Conceptual Modelling*, pp. 457–458. Springer. DOI: [10.1007/978-1-4612-5196-5\\_18](https://doi.org/10.1007/978-1-4612-5196-5_18).
- M. Stonebraker. 1983a. DBMS and AI: is there any common point of view? In *Proc. ACM SIGMOD International Conference on Management of Data*, p. 134. DOI: [10.1145/582192.582215](https://doi.org/10.1145/582192.582215). 201, 205
- M. Stonebraker. April 1983b. Chairman's column. *ACM SIGMOD Record*, 13(3): 1–3.
- M. Stonebraker. January 1983c. Chairman's column. *ACM SIGMOD Record*, 13(2): 1–3.
- M. Stonebraker. 1984. Virtual memory transaction management. *Operating Systems Review*, 18(2): 8–16. DOI: [10.1145/850755.850757](https://doi.org/10.1145/850755.850757). 203
- M. Stonebraker. 1985a. Triggers and inference in data base systems. In *Proc. 1985 ACM Annual Conference on the Range of Computing: Mid-80's Perspective*, p. 426. DOI: [10.1145/320435.323372](https://doi.org/10.1145/320435.323372).
- M. Stonebraker. 1985b. Triggers and inference in database systems. In M. L. Brodie and J. Mylopoulos, editors, *On Knowledge Base Management Systems*, pp. 297–314. Springer. 202
- M. Stonebraker. 1985c. Expert database systems/bases de données et systèmes experts. In *Journées Bases de Données Avancées*.
- M. Stonebraker. 1985d. The case for shared nothing. In *Proc. International Workshop on High-Performance Transaction Systems*, p. 0. 91
- M. Stonebraker. 1985e. Tips on benchmarking data base systems. *Quarterly Bulletin IEEE Technical Committee on Data Engineering*, 8(1): 10–18. <http://sites.computer.org/debull/85MAR-CD.pdf>.
- M. Stonebraker, editor. 1986a. *The INGRES Papers: Anatomy of a Relational Database System*. Addison-Wesley.
- M. Stonebraker. 1986b. Inclusion of new types in relational data base systems. In *Proc. 2nd International Conference on Data Engineering*, pp. 262–269. DOI: [10.1109/ICDE.1986.7266230](https://doi.org/10.1109/ICDE.1986.7266230). 88, 202, 258
- M. Stonebraker. 1986c. Object management in Postgres using procedures. In *Proc. International Workshop on Object-Oriented Database Systems*, pp. 66–72. <http://dl.acm.org/citation.cfm?id=318840>. 45, 88, 399
- M. Stonebraker. 1986d. The case for shared nothing. *Quarterly Bulletin IEEE Technical Committee on Data Engineering*, 9(1): 4–9. <http://sites.computer.org/debull/86MAR-CD.pdf>. 91, 216
- M. Stonebraker. 1986e. Design of relational systems (introduction to section 1). In M. Stonebraker, editor, *The INGRES Papers: Anatomy of a Relational Database System*, pp. 1–3. Addison-Wesley.
- M. Stonebraker. 1986f. Supporting studies on relational systems (introduction to section 2). In M. Stonebraker, editor, *The INGRES Papers*, pp. 83–85. Addison-Wesley.

- M. Stonebraker. 1986g. Distributed database systems (introduction to section 3). In M. Stonebraker, editor, *The INGRES Papers: Anatomy of a Relational Database System*, pp. 183–186. Addison-Wesley.
- M. Stonebraker. 1986h. The design and implementation of distributed INGRES. In M. Stonebraker, editor, *The INGRES Papers: Anatomy of a Relational Database System*, pp. 187–196. Addison-Wesley.
- M. Stonebraker. 1986i. User interfaces for database systems (introduction to section 4). In M. Stonebraker, editor, *The INGRES Papers: Anatomy of a Relational Database System*, pp. 243–245. Addison-Wesley.
- M. Stonebraker. 1986j. Extended semantics for the relational model (introduction to section 5). In M. Stonebraker, editor, *The INGRES Papers: Anatomy of a Relational Database System*, pp. 313–316. Addison-Wesley.
- M. Stonebraker. 1986k. Database design (introduction to section 6). In M. Stonebraker, editor, *The INGRES Papers: Anatomy of a Relational Database System*, pp. 393–394. Addison-Wesley.
- M. Stonebraker. 1986l. Object management in a relational data base system. In *Digest of Papers - COMPCON*, pp. 336–341.
- M. Stonebraker. 1987. The design of the POSTGRES storage system. In *Proc. 13th International Conference on Very Large Data Bases*, pp. 289–300. <http://www.vldb.org/conf/1987/P289.pdf>. 168, 214, 258
- M. Stonebraker, editor. 1988a. *Readings in Database Systems*. Morgan Kaufmann.
- M. Stonebraker. 1988b. Future trends in data base systems. In *Proc. 4th International Conference on Data Engineering*, pp. 222–231. DOI: [10.1109/ICDE.1988.105464](https://doi.org/10.1109/ICDE.1988.105464).
- M. Stonebraker. 1989a. The case for partial indexes. *ACM SIGMOD Record*, 18(4): 4–11. DOI: [10.1145/74120.74121](https://doi.org/10.1145/74120.74121).
- M. Stonebraker. 1989b. Future trends in database systems. *IEEE Transactions on Knowledge and Data Engineering*, 1(1): 33–44. DOI: [10.1109/69.43402](https://doi.org/10.1109/69.43402).
- M. Stonebraker. 1990a. The third-generation database manifesto: A brief retrospection. In *Proc. IFIP TC2/WG 2.6 Working Conference on Object-Oriented Databases: Analysis, Design & Construction*, pp. 71–72.
- M. Stonebraker. 1990b. Architecture of future data base systems. *Quarterly Bulletin IEEE Technical Committee on Data Engineering*, 13(4): 18–23. <http://sites.computer.org/debull/90DEC-CD.pdf>.
- M. Stonebraker. 1990c. Data base research at Berkeley. *ACM SIGMOD Record*, 19(4): 113–118. DOI: [10.1145/122058.122072](https://doi.org/10.1145/122058.122072).
- M. Stonebraker. 1990d. Introduction to the special issue on database prototype systems. *IEEE Transactions on Knowledge and Data Engineering*, 2(1): 1–3. DOI: [10.1109/TKDE.1990.10000](https://doi.org/10.1109/TKDE.1990.10000).
- M. Stonebraker. 1990e. The Postgres DBMS. In *Proc. ACM SIGMOD International Conference on Management of Data*, p. 394.

- M. Stonebraker. 1991a. Managing persistent objects in a multi-level store. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 2–11. DOI: [10.1145/115790.115791](https://doi.org/10.1145/115790.115791).
- M. Stonebraker. 1991b. Object management in Postgres using procedures. In K. R. Dittrich, U. Dayal, and A. P. Buchmann, editors, *On Object-Oriented Database System*, pp. 53–64. Springer. DOI: [http://10.1007/978-3-642-84374-7\\_5](https://doi.org/10.1007/978-3-642-84374-7_5).
- M. Stonebraker. 1971. The reduction of large scale Markov models for random chains. Ph.D. Dissertation. University of Michigan, Ann Arbor, MI. AAI7123885. 43
- M. Stonebraker. 1992a. The integration of rule systems and database systems. *IEEE Transactions on Knowledge and Data Engineering*, 4(5): 415–423. DOI: [10.1109/69.166984](https://doi.org/10.1109/69.166984). 91
- M. Stonebraker, editor. 1992b. *Proceedings of the 1992 ACM SIGMOD International Conference on Management of Data*. ACM Press.
- M. Stonebraker. 1993a. The SEQUOIA 2000 project. *Quarterly Bulletin IEEE Technical Committee on Data Engineering*, 16(1): 24–28. <http://sites.computer.org/debull/93MAR-CD.pdf>.
- M. Stonebraker. 1993b. Are we polishing a round ball? (panel abstract). In *Proc. 9th International Conference on Data Engineering*, p. 606.
- M. Stonebraker. 1993c. The miro DBMS. In *Proc. ACM SIGMOD International Conference on Management of Data*, p. 439. DOI: [10.1145/170035.170124](https://doi.org/10.1145/170035.170124). 314
- M. Stonebraker. 1994a. SEQUOIA 2000: A reflection of the first three years. In *Proc. 7th International Working Conference on Scientific and Statistical Database Management*, pp. 108–116. DOI: [10.1109/SSDM.1994.336956](https://doi.org/10.1109/SSDM.1994.336956).
- M. Stonebraker, editor. 1994b. *Readings in Database Systems*, 2. Morgan Kaufmann.
- M. Stonebraker. 1994c. Legacy systems—the Achilles heel of downsizing (panel). In *Proc. 3rd International Conference on Parallel and Distributed Information Systems*, p. 108.
- M. Stonebraker. 1994d. In memory of Bob Kooi (1951–1993). *ACM SIGMOD Record*, 23(1): 3. DOI: [10.1145/181550.181551](https://doi.org/10.1145/181550.181551).
- M. Stonebraker. 1995. An overview of the Sequoia 2000 project. *Digital Technical Journal*, 7(3). <http://www.hpl.hp.com/hpjournal/dtj/vol7num3/vol7num3art3.pdf>. 215, 255
- M. Stonebraker. 1998. Are we working on the right problems? (panel). In *Proc. ACM SIGMOD International Conference on Management of Data*, p. 496. DOI: [10.1145/276304.276348](https://doi.org/10.1145/276304.276348).
- M. Stonebraker. 2002. Too much middleware. *ACM SIGMOD Record*, 31(1): 97–106. DOI: [10.1145/507338.507362](https://doi.org/10.1145/507338.507362). 91
- M. Stonebraker. 2003. Visionary: A next generation visualization system for databases. In *Proc. ACM SIGMOD International Conference on Management of Data*, p. 635. <http://www.acm.org/sigmod/sigmod03/eproceedings/papers/ind00.pdf>.
- M. Stonebraker. 2004. Outrageous ideas and/or thoughts while shaving. In *Proc. 20th International Conference on Data Engineering*, p. 869. DOI: [10.1109/ICDE.2004.1320096](https://doi.org/10.1109/ICDE.2004.1320096).



- M. Stonebraker. 2008a. Why did Jim Gray win the Turing Award? *ACM SIGMOD Record*, 37(2): 33–34. DOI: [10.1145/1379387.1379398](https://doi.org/10.1145/1379387.1379398).
- M. Stonebraker. 2008b. Technical perspective—one size fits all: An idea whose time has come and gone. *Communications of the ACM*, 51(12): 76. DOI: [10.1145/1409360.1409379](https://doi.org/10.1145/1409360.1409379). 92
- M. Stonebraker. 2009a. Stream processing. In L. Liu and M. T. Özsu, editors. *Encyclopedia of Database Systems*, pp. 2837–2838. Springer. DOI: [10.1007/978-0-387-39940-9\\_371](https://doi.org/10.1007/978-0-387-39940-9_371).
- M. Stonebraker. 2009b. A new direction for tpc? In *Proc. 1st TPC Technology Conference on Performance Evaluation and Benchmarking*, pp. 11–17. DOI: [10.1007/978-3-642-10424-4\\_2](https://doi.org/10.1007/978-3-642-10424-4_2).
- M. Stonebraker. 2010a. SQL databases v. nosql databases. *Communications of the ACM*, 53(4): 10–11. DOI: [10.1145/1721654.1721659](https://doi.org/10.1145/1721654.1721659). 50
- M. Stonebraker. 2010b. In search of database consistency. *Communications of the ACM*, 53(10): 8–9. DOI: [10.1145/1831407.1831411](https://doi.org/10.1145/1831407.1831411).
- M. Stonebraker. 2011a. Stonebraker on data warehouses. *Communications of the ACM*, 54(5): 10–11. DOI: [10.1145/1941487.1941491](https://doi.org/10.1145/1941487.1941491).
- M. Stonebraker. 2011b. Stonebraker on nosql and enterprises. *Communications of the ACM*, 54(8): 10–11. DOI: [10.1145/1978542.1978546](https://doi.org/10.1145/1978542.1978546). 50
- M. Stonebraker. 2012a. SciDB: An open-source DBMS for scientific data. *ERCIM News*, 2012(89). <http://ercim-news.ercim.eu/en89/special/scidb-an-open-source-dbms-for-scientific-data>.
- M. Stonebraker. 2012b. New opportunities for new SQL. *Communications of the ACM*, 55(11): 10–11. DOI: [10.1145/2366316.2366319](https://doi.org/10.1145/2366316.2366319).
- M. Stonebraker. 2013. We are under attack; by the least publishable unit. In *Proc. 6th Biennial Conference on Innovative Data Systems Research*. [http://www.cidrdb.org/cidr2013/Talks/CIDR13\\_Gongshow16.ppt](http://www.cidrdb.org/cidr2013/Talks/CIDR13_Gongshow16.ppt). 273
- M. Stonebraker. 2015a. Turing lecture. In *Proc. Federated Computing Research Conference*, pp. 2–2. DOI: [10.1145/2820468.2820471](https://doi.org/10.1145/2820468.2820471).
- M. Stonebraker. 2015b. What it's like to win the Turing Award. *Communications of the ACM*, 58(11): 11. [xxxi](#), [xxxiii](#)
- M. Stonebraker. 2015c. The Case for Polystores. ACM SIGMOD Blog, <http://wp.sigmod.org/?p=1629>. 370, 371
- M. Stonebraker. 2016. The land sharks are on the squawk box. *Communications of the ACM*, 59(2): 74–83. DOI: [10.1145/2869958](https://doi.org/10.1145/2869958). 50, 129, 139, 260, 319
- M. Stonebraker. 2018. My top ten fears about the DBMS field. In *Proc. 34th International Conference on Data Engineering*, pp. 24–28.
- M. Sullivan and M. Stonebraker. 1991. Using write protected data structures to improve software fault tolerance in highly available database management systems. In *Proc. 17th International Conference on Very Large Data Bases*, pp. 171–180. <http://www.vldb.org/conf/1991/P171.pdf>.



- R. Taft, E. Mansour, M. Serafini, J. Duggan, A. J. Elmore, A. Aboulmaga, A. Pavlo, and M. Stonebraker. 2014a. E-store: Fine-grained elastic partitioning for distributed transaction processing. *Proc. VLDB Endowment*, 8(3): 245–256. <http://www.vldb.org/pvldb/vol8/p245-taft.pdf>. 188, 251
- R. Taft, M. Vartak, N. R. Satish, N. Sundaram, S. Madden, and M. Stonebraker. 2014b. Genbase: a complex analytics genomics benchmark. In *Proc. ACM SIGMOD International Conference on Management of Data*, pp. 177–188. DOI: [10.1145/2588555.2595633](https://doi.org/10.1145/2588555.2595633).
- R. Taft, W. Lang, J. Duggan, A. J. Elmore, M. Stonebraker, and D. J. DeWitt. 2016. Step: Scalable tenant placement for managing database-as-a-service deployments. In *Proc. 7th ACM Symposium on Cloud Computing*, pp. 388–400. DOI: [10.1145/2987550.2987575](https://doi.org/10.1145/2987550.2987575).
- R. Taft, N. El-Sayed, M. Serafini, Y. Lu, A. Aboulmaga, M. Stonebraker, R. Mayerhofer, and F. Andrade. 2018. P-Store: an elastic database system with predictive provisioning. In *Proc. ACM SIGMOD International Conference on Management of Data*. 188
- W. Tao, D. Deng, and M. Stonebraker. 2017. Approximate string joins with abbreviations. *Proc. VLDB Endowment*, 11(1): 53–65.
- N. Tatbul, U. Çetintemel, S. B. Zdonik, M. Cherniack, and M. Stonebraker. 2003. Load shedding in a data stream manager. In *Proc. 29th International Conference on Very Large Data Bases*, pp. 309–320. <http://www.vldb.org/conf/2003/papers/S10P03.pdf>. 228, 229
- N. Tatbul, S. Zdonik, J. Meehan, C. Aslantas, M. Stonebraker, K. Tufte, C. Giossi, and H. Quach. 2015. Handling shared, mutable state in stream processing with correctness guarantees. *Quarterly Bulletin IEEE Technical Committee on Data Engineering*, 38(4): 94–104. <http://sites.computer.org/debull/A15dec/p94.pdf>.
- T. J. Teorey, J. W. DeHeus, R. Gerritsen, H. L. Morgan, J. F. Spitzer, and M. Stonebraker. 1976. SIGMOD (paper session). In *Proc. 1976 ACM Annual Conference*, p. 275. DOI: [10.1145/800191.805596](https://doi.org/10.1145/800191.805596).
- M. S. Tuttle, S. H. Brown, K. E. Campbell, J. S. Carter, K. Keck, M. J. Lincoln, S. J. Nelson, and M. Stonebraker. 2001a. The semantic web as “perfection seeking”: A view from drug terminology. In *Proc. 1st Semantic Web Working Symposium*, pp. 5–16. <http://www.semanticweb.org/SWWS/program/full/paper49.pdf>.
- M. S. Tuttle, S. H. Brown, K. E. Campbell, J. S. Carter, K. Keck, M. J. Lincoln, S. J. Nelson, and M. Stonebraker. 2001b. The semantic web as “perfection seeking”: A view from drug terminology. In I. F. Cruz, S. Decker, J. Euzenat, and D. L. McGuinness, editors, *The Emerging Semantic Web, Selected Papers from the 1st Semantic Web Working Symposium*, volume 75 of *Frontiers in Artificial Intelligence and Applications*. IOS Press.
- J. Widom, A. Bosworth, B. Lindsey, M. Stonebraker, and D. Suciu. 2000. Of XML and databases (panel session): Where’s the beef? In *Proc. ACM SIGMOD International Conference on Management of Data*, p. 576. DOI: [10.1145/335191.335476](https://doi.org/10.1145/335191.335476).

- M. W. Wilkins, R. Berlin, T. Payne, and G. Wiederhold. 1985. Relational and entity-relationship model databases and specialized design files in vlsi design. In *Proc. 22nd ACM/IEEE Design Automation Conference*, pp. 410–416.
- J. Woodfill and M. Stonebraker. 1983. An implementation of hypothetical relations. In *Proc. 9th International Conference on Very Data Bases*, pp. 157–166. <http://www.vldb.org/conf/1983/P157.pdf>.
- A. Woodruff and M. Stonebraker. 1995. Buffering of intermediate results in dataflow diagrams. In *Proc. IEEE Symposium on Visual Languages*, p. 187. DOI: [10.1109/VL.1995.520808](https://doi.org/10.1109/VL.1995.520808).
- A. Woodruff and M. Stonebraker. 1997. Supporting fine-grained data lineage in a database visualization environment. In *Proc. 13th International Conference on Data Engineering*, pp. 91–102. DOI: [10.1109/ICDE.1997.581742](https://doi.org/10.1109/ICDE.1997.581742).
- A. Woodruff, P. Wisnovsky, C. Taylor, M. Stonebraker, C. Paxson, J. Chen, and A. Aiken. 1994. Zooming and tunneling in Tioga: Supporting navigation in multimedia space. In *Proc. IEEE Symposium on Visual Languages*, pp. 191–193. DOI: [10.1109/VL.1994.363622](https://doi.org/10.1109/VL.1994.363622).
- A. Woodruff, A. Su, M. Stonebraker, C. Paxson, J. Chen, A. Aiken, P. Wisnovsky, and C. Taylor. 1995. Navigation and coordination primitives for multidimensional visual browsers. In *Proc. IFIP WG 2.6 3rd Working Conference Visual Database Systems*, pp. 360–371. DOI: [10.1007/978-0-387-34905-3\\_23](https://doi.org/10.1007/978-0-387-34905-3_23).
- A. Woodruff, J. A. Landay, and M. Stonebraker. 1998a. Goal-directed zoom. In *CHI '98 Conference Summary on Human Factors in Computing Systems*, pp. 305–306. DOI: [10.1145/286498.286781](https://doi.org/10.1145/286498.286781).
- A. Woodruff, J. A. Landay, and M. Stonebraker. 1998b. Constant density visualizations of non-uniform distributions of data. In *Proc. 11th Annual ACM Symposium on User Interface Software and Technology*, pp. 19–28. DOI: [10.1145/288392.288397](https://doi.org/10.1145/288392.288397).
- A. Woodruff, J. A. Landay, and M. Stonebraker. 1998c. Constant information density in zoomable interfaces. In *Proc. Working Conference on Advanced Visual Interfaces*, pp. 57–65. DOI: [10.1145/948496.948505](https://doi.org/10.1145/948496.948505).
- A. Woodruff, J. A. Landay, and M. Stonebraker. 1999. VIDA: (visual information density adjuster). In *CHI '99 Extended Abstracts on Human Factors in Computing Systems*, pp. 19–20. DOI: [10.1145/632716.632730](https://doi.org/10.1145/632716.632730).
- A. Woodruff, C. Olston, A. Aiken, M. Chu, V. Ercegovic, M. Lin, M. Spalding, and M. Stonebraker. 2001. Datasplash: A direct manipulation environment for programming semantic zoom visualizations of tabular data. *Journal of Visual Languages and Computing*, 12(5): 551–571. DOI: [10.1006/jvlc.2001.0219](https://doi.org/10.1006/jvlc.2001.0219).
- E. Wu, S. Madden, and M. Stonebraker. 2012. A demonstration of dbwipes: Clean as you query. *Proc. VLDB Endowment*, 5(12): 1894–1897. DOI: [10.14778/2367502.2367531](https://doi.org/10.14778/2367502.2367531).
- E. Wu, S. Madden, and M. Stonebraker. 2013. Subzero: A fine-grained lineage system for scientific databases. In *Proc. 29th International Conference on Data Engineering*, pp. 865–876. DOI: [10.1109/ICDE.2013.6544881](https://doi.org/10.1109/ICDE.2013.6544881).

- X. Yu, G. Bezerra, A. Pavlo, S. Devadas, and M. Stonebraker. 2014. Staring into the abyss: An evaluation of concurrency control with one thousand cores. *Proc. VLDB Endowment*, 8(3): 209–220. <http://www.vldb.org/pvldb/vol8/p209-yu.pdf>.
- K. Yu, V. Gadepally, and M. Stonebraker. 2017. Database engine integration and performance analysis of the BigDAWG polystore system. *High Performance Extreme Computing Conference (HPEC)*. IEEE, 2017. DOI: [10.1109/HPEC.2017.8091081](https://doi.org/10.1109/HPEC.2017.8091081). 376
- S. B. Zdonik, M. Stonebraker, M. Cherniack, U. Çetintemel, M. Balazinska, and H. Balakrishnan. 2003. The aurora and medusa projects. *Quarterly Bulletin IEEE Technical Committee on Data Engineering*, 26(1): 3–10. <http://sites.computer.org/debull/A03mar/zdonik.ps>. 228, 324



## References

- D. Abadi, Y. Ahmad, M. Balazinska, U. Çetintemel, M. Cherniack, J.-H. Hwang, W. Lindner, A. Maskey, A. Rasin, E. Ryzkina, N. Tatbul, Y. Xing, and S. Zdonik. 2005. The design of the Borealis stream processing engine. *Proc. of the 2nd Biennial Conference on Innovative Data Systems Research (CIDR'05)*, Asilomar, CA, January. 228
- Z. Abedjan, L. Golab, and F. Naumann. August 2015. Profiling relational data: a survey. *The VLDB Journal*, 24(4): 557–581. DOI: DOI: [10.1007/s00778-015-0389-y](https://doi.org/10.1007/s00778-015-0389-y). 297
- ACM. 2015a. Announcement: Michael Stonebraker, Pioneer in Database Systems Architecture, Receives 2014 ACM Turing Award. [http://amturing.acm.org/award\\_winners/stonebraker\\_1172121.cfm](http://amturing.acm.org/award_winners/stonebraker_1172121.cfm). Accessed February 5, 2018.
- ACM. March 2015b. Press Release: MIT's Stonebraker Brought Relational Database Systems from Concept to Commercial Success, Set the Research Agenda for the Multibillion-Dollar Database Field for Decades. [http://sigmodrecord.org/publications/sigmodRecord/1503/pdfs/04\\_announcements\\_Stonebraker.pdf](http://sigmodrecord.org/publications/sigmodRecord/1503/pdfs/04_announcements_Stonebraker.pdf). Accessed February 5, 2018.
- ACM. 2016. A.M. Turing Award Citation and Biography. [http://amturing.acm.org/award\\_winners/stonebraker\\_1172121.cfm](http://amturing.acm.org/award_winners/stonebraker_1172121.cfm). Accessed September 24, 2018. xxxi
- Y. Ahmad, B. Berg, U. Çetintemel, M. Humphrey, J. Hwang, A. Jhingran, A. Maskey, O. Papaemmanouil, A. Rasin, N. Tatbul, W. Xing, Y. Xing, and S. Zdonik. June 2005. Distributed operation in the Borealis Stream Processing Engine. Demonstration, *ACM SIGMOD International Conference on Management of Data (SIGMOD'05)*. Baltimore, MD. Best Demonstration Award. 230, 325
- M. M. Astrahan, M. W. Blasgen, D. D. Chamberlin, K. P. Eswaran, J. N. Gray, P. P. Griffiths, W. F. King, R. A. Lorie, P. R. McJones, J. W. Mehl, G. R. Putzolu, I. L. Traiger, B. W. Wade, and V. Watson. 1976. System R: relational approach to database management. *ACM Transactions on Database Systems*, 1(2): 97–137. DOI: [10.1145/320455.320457](https://doi.org/10.1145/320455.320457). 397
- P. Bailis, E. Gan, S. Madden, D. Narayanan, K. Rong, and S. Suri. 2017. Macrobase: Prioritizing attention in fast data. *Proc. of the 2017 ACM International Conference on Management of Data*. ACM. DOI: [10.1145/3035918.3035928](https://doi.org/10.1145/3035918.3035928). 374
- Berkeley Software Distribution. n.d. In Wikipedia. [http://en.wikipedia.org/wiki/Berkeley\\_Software\\_Distribution](http://en.wikipedia.org/wiki/Berkeley_Software_Distribution). Last accessed March 1, 2018. 109

- G. Beskales, I.F. Ilyas, L. Golab, and A. Galiullin. 2013. On the relative trust between inconsistent data and inaccurate constraints. *Proc. of the IEEE International Conference on Data Engineering, ICDE 2013*, pp. 541–552. Australia. DOI: [10.1109/ICDE.2013.6544854](https://doi.org/10.1109/ICDE.2013.6544854). 270
- L. S. Blackford, J. Choi, A. Cleary, E. D’Azevedo, J. Demmel, I. Dhillon, J. Dongarra, S. Hammarling, G. Henry, A. Petitet, K. Stanley, D. Walker, R. C. Whaley. 2017. *ScaLAPACK Users’ Guide*. Society for Industrial and Applied Mathematics <http://netlib.org/scalapack/slug/index.html>. Last accessed December 31, 2017. 258
- D. Bitton, D. J. DeWitt, and C. Turbyfill. 1983. Benchmarking database systems—a systematic approach. Computer Sciences Technical Report #526, University of Wisconsin. <http://minds.wisconsin.edu/handle/1793/58490>, 111
- P. A. Boncz, M. L. Kersten, and S. Manegold. December 2008. Breaking the memory wall in MonetDB. *Communications of the ACM*, 51(12): 77–85. DOI: [10.1145/1409360.1409380](https://doi.org/10.1145/1409360.1409380). 151
- M. L. Brodie. June 2015. Understanding data science: an emerging discipline for data-intensive discovery. In S. Cutt, editor, *Getting Data Right: Tackling the Challenges of Big Data Volume and Variety*. O’Reilly Media, Sebastopol, CA. 291
- Brown University, Department of Computer Science. Fall 2002. Next generation stream-based applications. *Conduit Magazine*, 11(2). [https://cs.brown.edu/about/conduit/conduit\\_v11n2.pdf](https://cs.brown.edu/about/conduit/conduit_v11n2.pdf). Last accessed May 14, 2018. 322
- BSD licenses. n.d. In Wikipedia. [http://en.wikipedia.org/wiki/BSD\\_licenses](http://en.wikipedia.org/wiki/BSD_licenses). Last accessed March 1, 2018. 109
- M. Cafarella and C. Ré. April 2018. The last decade of database research and its blindingly bright future. or Database Research: A love song. DAWN Project, Stanford University. <http://dawn.cs.stanford.edu/2018/04/11/db-community/>. 6
- M. J. Carey, D. J. DeWitt, M. J. Franklin, N. E. Hall, M. L. McAuliffe, J. F. Naughton, D. T. Schuh, M. H. Solomon, C. K. Tan, O. G. Tsatalos, S. J. White, and M. J. Zwilling. 1994. Shoring up persistent applications. *Proc. of the 1994 ACM SIGMOD international conference on Management of data (SIGMOD ’94)*, 383–394. DOI: [10.1145/191839.191915](https://doi.org/10.1145/191839.191915). 152
- M. J. Carey, D. J. DeWitt, M. J. Franklin, N.E. Hall, M. L. McAuliffe, J. F. Naughton, D. T. Schuh, M. H. Solomon, C. K. Tan, O. G. Tsatalos, S. J. White, and M. J. Zwilling. 1994. Shoring up persistent applications. In *Proc. of the 1994 ACM SIGMOD International Conference on Management of Data (SIGMOD ’94)*, pp. 383–394. DOI: [10.1145/191839.191915](https://doi.org/10.1145/191839.191915). 336
- M. J. Carey, L. M. Haas, P. M. Schwarz, M. Arya, W. E. Cody, R. Fagin, M. Flickner, A. W. Luniewski, W. Niblack, and D. Petkovic. 1995. Towards heterogeneous multimedia information systems: The garlic approach. In *Research Issues in Data Engineering, 1995: Distributed Object Management, Proceedings*, pp. 124–131. IEEE. DOI: [10.1109/RIDE.1995.378736](https://doi.org/10.1109/RIDE.1995.378736). 284
- CERN. <http://home.cern/about/computing>. Last accessed December 31, 2017.
- D. D. Chamberlin and R. F. Boyce. 1974. SEQUEL: A structured English query language. In *Proc. of the 1974 ACM SIGFIDET (now SIGMOD) Workshop on Data Description*,

- Access and Control (SIGFIDET '74)*, pp. 249–264. ACM, New York. DOI: [10.1145/800296.811515](https://doi.org/10.1145/800296.811515). 404, 407
- D. D. Chamberlin, M. M. Astrahan, K. P. Eswaran, P. P. Griffiths, R. A. Lorie, J. W. Mehl, P. Reisner, and B. W. Wade. 1976. SEQUEL 2: a unified approach to data definition, manipulation, and control. *IBM Journal of Research and Development*, 20(6): 560–575. DOI: [10.1147/rd.206.0560](https://doi.org/10.1147/rd.206.0560). 398
- S. Chandrasekaran, O. Cooper, A. Deshpande, M.J. Franklin, J.M. Hellerstein, W. Hong, S. Krishnamurthy, S. Madden, V. Raman, F. Reiss, and M. Shah. 2003. TelegraphCQ: Continuous dataflow processing for an uncertain world. *Proc. of the 2003 ACM SIGMOD International Conference on Management of Data (SIGMOD '03)*, pp. 668–668. ACM, New York. DOI: [10.1145/872757.872857](https://doi.org/10.1145/872757.872857). 231
- J. Chen, D.J. DeWitt, F. Tian, and Y. Wang. 2000. NiagaraCQ: A scalable continuous query system for Internet databases. *Proc. of the 2000 ACM SIGMOD International Conference on Management of Data (SIGMOD '00)*, pp. 379–390. ACM, New York. DOI: [10.1145/342009.335432](https://doi.org/10.1145/342009.335432). 231
- M. Cherniack, H. Balakrishnan, M. Balazinska, D. Carney, U. Çetintemel, Y. Xing, and S. Zdonik. 2003. Scalable distributed stream processing. *Proc. of the First Biennial Conference on Innovative Database Systems (CIDR'03)*, Asilomar, CA, January. 228
- C. M. Christensen. 1997. *The Innovator's Dilemma: When New Technologies Cause Great Firms to Fail*. Harvard Business School Press, Boston, MA. 100
- X. Chu, I. F. Ilyas, and P. Papotti. 2013a. Holistic data cleaning: Putting violations into context. *Proc. of the IEEE International Conference on Data Engineering, ICDE 2013*, pp. 458–469. Australia. DOI: [10.1109/ICDE.2013.6544847](https://doi.org/10.1109/ICDE.2013.6544847). 270, 297
- X. Chu, I. F. Ilyas, and P. Papotti. 2013b. Discovering denial constraints. *Proc. of the VLDB Endowment, PVLDB 6(13)*: 1498–1509. DOI: [10.14778/2536258.2536262](https://doi.org/10.14778/2536258.2536262). 270
- X. Chu, J. Morcos, I. F. Ilyas, M. Ouzzani, P. Papotti, N. Tang, and Y. Ye. 2015. Katara: A data cleaning system powered by knowledge bases and crowdsourcing. In *Proc. of the 2015 ACM SIGMOD International Conference on Management of Data (SIGMOD '15)*, pp. 1247–1261. ACM, New York. DOI: [10.1145/2723372.2749431](https://doi.org/10.1145/2723372.2749431). 297
- P. J. A. Cock, C. J. Fields, N. Goto, M. L. Heuer, and P. M. Rice. 2009. The Sanger FASTQ file format for sequences with quality scores, and the Solexa/Illumina FASTQ variants. *Nucleic Acids Research* 38.6: 1767–1771. DOI: [10.1093/nar/gkp1137](https://doi.org/10.1093/nar/gkp1137). 374
- E. F. Codd. June 1970. A relational model of data for large shared data banks. *Communications of the ACM*, 13(6): 377–387. DOI: [10.1145/362384.362685](https://doi.org/10.1145/362384.362685). 42, 98, 166, 397, 404, 405, 407
- M. Collins. 2016. Thomson Reuters uses Tamr to deliver better connected content at a fraction of the time and cost of legacy approaches. Tamr blog, July 28. <https://www.tamr.com/video/thomson-reuters-uses-tamr-deliver-better-connected-content-fraction-time-cost-legacy-approaches/>. Last accessed January 24, 2018. 275

- G. Copeland and D. Maier. 1984. Making smalltalk a database system. *Proc. of the 1984 ACM SIGMOD International Conference on Management of Data (SIGMOD '84)*, pp. 316–325. ACM, New York. DOI: [10.1145/602259.602300](https://doi.org/10.1145/602259.602300). 111
- C. Cranor, T. Johnson, V. Shkapenyuk, and O. Spatscheck. 2003. Gigascope: A stream database for network applications. *Proc. of the 2003 ACM SIGMOD International Conference on Management of Data (SIGMOD '03)*, pp. 647–651. ACM, New York. DOI: [10.1145/872757.872838](https://doi.org/10.1145/872757.872838). 231
- A. Crotty, A. Galakatos, K. Dursun, T. Kraska, U. Cetintemel, and S. Zdonik. 2015. Tupleware: “Big Data, Big Analytics, Small Clusters.” *CIDR*. DOI: [10.1.1.696.32](https://doi.org/10.1.1.696.32). 374
- M. Dallachiesa, A. Ebaid, A. Eldawi, A. Elmagarmid, I. F. Ilyas, M. Ouzzani, and N. Tang. 2013. NADEEF, a commodity data cleaning system. *Proc. of the 2013 ACM SIGMOD Conference on Management of Data*, pp. 541–552. New York. <http://dx.doi.org/10.1145/2463676.2465327>. 270, 297
- T. Dasu and J. M. Loh. 2012. Statistical distortion: Consequences of data cleaning. *PVLDB*, 5(11): 1674–1683. DOI: [10.14778/2350229.2350279](https://doi.org/10.14778/2350229.2350279). 297
- C. J. Date and E. F. Codd. 1975. The relational and network approaches: Comparison of the application programming interfaces. In *Proc. of the 1974 ACM SIGFIDET (now SIGMOD) Workshop on Data Description, Access and Control: Data Models: Data-Structure-Set Versus Relational (SIGFIDET '74)*, pp. 83–113. ACM, New York. DOI: [10.1145/800297.811534](https://doi.org/10.1145/800297.811534). 405
- D. J. DeWitt. 1979a. Direct a multiprocessor organization for supporting relational database management systems. *IEEE Transactions of Computers*, 28(6), 395–406. DOI: [10.1109/TC.1979.1675379](https://doi.org/10.1109/TC.1979.1675379). 109
- D. J. DeWitt. 1979b. Query execution in DIRECT. In *Proc. of the 1979 ACM SIGMOD International Conference on Management of Data (SIGMOD '79)*, pp. 13–22. ACM, New York. DOI: [10.1145/582095.582098](https://doi.org/10.1145/582095.582098). 109
- D. J. DeWitt, R. H. Gerber, G. Graefe, M. L. Heytens, K. B. Kumar, and M. Muralikrishna. 1986. GAMMA—a high performance dataflow database machine. *Proc. of the 12th International Conference on Very Large Data Bases (VLDB '86)*, W. W. Chu, G. Gardarin, S. Ohsuga, and Y. Kambayashi, editors, pp. 228–237. Morgan Kaufmann Publishers Inc., San Francisco, CA. 111
- D. J. DeWitt, S. Ghandeharizadeh, D. A. Schneider, A. Bricker, H.-I. Hsiao, and R. Rasmussen. March 1990. The Gamma database machine project. *IEEE Transactions on Knowledge and Data Engineering*, 2(1): 44–62. DOI: [10.1109/69.50905](https://doi.org/10.1109/69.50905). 151, 400
- D. DeWitt and J. Gray. June 1992. Parallel database systems: the future of high performance database systems. *Communications of the ACM*, 35(6): 85–98. DOI: [10.1145/129888.129894](https://doi.org/10.1145/129888.129894). 199
- D. J. DeWitt, A. Halverson, R. Nehme, S. Shankar, J. Aguilar-Saborit, A. Avanes, M. Flasz, and J. Gramling. 2013. Split query processing in polybase. *Proc. of the 2013 ACM SIGMOD International Conference on Management of Data (SIGMOD '13)*, pp. 1255–1266. ACM, New York. 284



- C. Diaconu, C. Freedman, E. Ismert, P-A. Larson, P. Mittal, R. Stonecipher, N. Verma, and M. Zwillig. 2013. Hekaton: SQL server's memory-optimized OLTP engine. In *Proc. of the 2013 ACM SIGMOD International Conference on Management of Data (SIGMOD '13)*, pp. 1243–1254. ACM, New York. DOI: [10.1145/2463676.2463710](https://doi.org/10.1145/2463676.2463710).
- K. P. Eswaran, J. N. Gray, R. A. Lorie, and I. L. Traiger. November 1976. The notions of consistency and predicate locks in a database system. *Communications of the ACM*, 19(11): 624–633. DOI: [10.1145/360363.360369](https://doi.org/10.1145/360363.360369). 114
- W. Fan, J. Li, S. Ma, N. Tang, and W. Yu. April 2012. Towards certain fixes with editing rules and master data. *The VLDB Journal*, 21(2): 213–238. DOI: [10.1007/s00778-011-0253-7](https://doi.org/10.1007/s00778-011-0253-7). 297
- D. Fogg. September 1982. Implementation of domain abstraction in the relational database system INGRES. Master of Science Report, Dept. of Electrical Engineering and Computer Sciences, University of California, Berkeley, CA. 201
- T. Flory, A. Robbin, and M. David. May 1988. Creating SIPP longitudinal analysis files using a relational database management system. CDE Working Paper No. 88-32, Institute for Research on Poverty, University of Wisconsin-Madison, Madison, WI. 197
- V. Gadepally, J. Kepner, W. Arcand, D. Bestor, B. Bergeron, C. Byun, L. Edwards, M. Hubbell, P. Michaleas, J. Mullen, A. Prout, A. Rosa, C. Yee, and A. Reuther. 2015. D4M: Bringing associative arrays to database engines. *High Performance Extreme Computing Conference (HPEC)*. IEEE, 2015. DOI: [10.1109/HPEC.2015.7322472](https://doi.org/10.1109/HPEC.2015.7322472). 370
- V. Gadepally, K. O'Brien, A. Dziedzic, A. Elmore, J. Kepner, S. Madden, T. Mattson, J. Rogers, Z. She, and M. Stonebraker. September 2017. BigDAWG Version 0.1. *IEEE High Performance Extreme*. DOI: [10.1109/HPEC.2017.8091077](https://doi.org/10.1109/HPEC.2017.8091077). 288, 369
- J. Gantz and D. Reinsel. 2013. The Digital Universe in 2020: Big Data, Bigger Digital Shadows, and Biggest Growth in the Far East—United States, IDC, February. 5
- L. Gerhardt, C. H. Faham, and Y. Yao. 2015. Accelerating scientific analysis with SciDB. *Journal of Physics: Conference Series*, 664(7). 268
- B. Grad. 2007. Oral history of Michael Stonebraker, Transcription. Recorded: August 23, 2007. Computer History Museum, Moultonborough, NH. <http://archive.computerhistory.org/resources/access/text/2012/12/102635858-05-01-acc.pdf>. Last accessed April 8, 2018. 42, 43, 44, 98
- A. Guttman. 1984. R-trees: a dynamic index structure for spatial searching. In *Proc. of the 1984 ACM SIGMOD International Conference on Management of Data (SIGMOD '84)*, pp. 47–57. ACM, New York. DOI: [10.1145/602259.602266](https://doi.org/10.1145/602259.602266). 205
- L. M. Haas, J. C. Freytag, G. M. Lohman, and H. Pirahesh. 1989. Extensible query processing in starburst. In *Proc. of the 1989 ACM SIGMOD International Conference on Management of Data (SIGMOD '89)*, pp. 377–388. ACM, New York. DOI: [10.1145/67544.66962](https://doi.org/10.1145/67544.66962). 399
- D. Halperin, V. Teixeira de Almeida, L. L. Choo, S. Chu, P. Koutris, D. Moritz, J. Ortiz, V. Ruamviboonsuk, J. Wang, A. Whitaker. 2014. Demonstration of the Myria big data management service. *Proc. of the 2014 ACM SIGMOD International Conference*

- on Management of Data (SIGMOD '14)*, p. 881–884. ACM, New York. DOI: [10.1145/2588555.2594530](https://doi.org/10.1145/2588555.2594530). 284, 370
- B. Haynes, A. Cheung, and M. Balazinska. 2016. PipeGen: Data pipe generator for hybrid analytics. *Proc. of the Seventh ACM Symposium on Cloud Computing (SoCC '16)*, M. K. Aguilera, B. Cooper, and Y. Diao, editors, pp. 470–483. ACM, New York. DOI: [10.1145/2987550.2987567](https://doi.org/10.1145/2987550.2987567). 287
- M. A. Hearst. 2009. *Search user interfaces*. Cambridge University Press, New York. 394
- J. M. Hellerstein, J. F. Naughton, and A. Pfeffer. 1995. Generalized search trees for database systems. In *Proc. of the 21th International Conference on Very Large Data Bases (VLDB '95)*, pp. 562–573. Morgan Kaufmann Publishers Inc., San Francisco, CA. <http://dl.acm.org/citation.cfm?id=645921.673145>. 210
- J. M. Hellerstein, E. Koutsoupias, D. P. Miranker, C. H. Papadimitriou, V. Samoladas. 2002. On a model of indexability and its bounds for range queries, *Journal of the ACM (JACM)*, 49.1: 35–55. DOI: [10.1145/505241.505244](https://doi.org/10.1145/505241.505244). 210
- IBM. 1997. Special Issue on IBM's S/390 Parallel Sysplex Cluster. *IBM Systems Journal*, 36(2). 400
- S. Idreos, F. Groffen, N. Nes, S. Manegold, S. K. Mullender, and M. L. Kersten. 2012. MonetDB: two decades of research in column-oriented database architectures. *IEEE Data Engineering Bulletin*, 35(1): 40–45. 258
- N. Jain, S. Mishra, A. Srinivasan, J. Gehrke, J. Widom, H. Balakrishnan, U. Çetintemel, M. Cherniack, R. Tibbetts, and S. Zdonik. 2008. Towards a streaming SQL standard. *Proc. VLDB Endowment*, pp. 1379–1390. August 1–2. DOI: [10.14778/1454159.1454179](https://doi.org/10.14778/1454159.1454179). 229
- A. E. W. Johnson, T. J. Pollard, L. Shen, L. H. Lehman, M. Feng, M. Ghassemi, B. E. Moody, P. Szolovits, L. A. G. Celi, and R. G. Mark. 2016. MIMIC-III, a freely accessible critical care database. *Scientific Data* 3: 160035 DOI: [10.1038/sdata.2016.35](https://doi.org/10.1038/sdata.2016.35). 370
- V. Josifovski, P. Schwarz, L. Haas, and E. Lin. 2002. Garlic: a new flavor of federated query processing for DB2. In *Proc. of the 2002 ACM SIGMOD International Conference on Management of Data (SIGMOD '02)*, pp. 524–532. ACM, New York. DOI: [10.1145/564691.564751](https://doi.org/10.1145/564691.564751). 401
- J. W. Josten, C. Mohan, I. Narang, and J. Z. Teng. 1997. DB2's use of the coupling facility for data sharing. *IBM Systems Journal*, 36(2): 327–351. DOI: [10.1147/sj.362.0327](https://doi.org/10.1147/sj.362.0327). 400
- S. Kandel, A. Paepcke, J. Hellerstein, and J. Heer. 2011. Wrangler: Interactive visual specification of data transformation scripts. In *Proc. of the SIGCHI Conference on Human Factors in Computing Systems (CHI '11)*, pp. 3363–3372. ACM, New York. DOI: [10.1145/1978942.1979444](https://doi.org/10.1145/1978942.1979444). 297
- R. Katz. editor. June 1982. Special issue on design data management. *IEEE Database Engineering Newsletter*, 5(2). 200
- J. Kepner, V. Gadepally, D. Hutchison, H. Jensen, T. Mattson, S. Samsi, and A. Reuther. 2016. Associative array model of SQL, NoSQL, and NewSQL Databases. *IEEE High*

- Performance Extreme Computing Conference (HPEC) 2016*, Waltham, MA, September 13–15. DOI: [10.1109/HPEC.2016.7761647](https://doi.org/10.1109/HPEC.2016.7761647). 289
- V. Kevin and M. Whitney. 1974. Relational data management implementation techniques. In *Proc. of the 1974 ACM SIGFIDET (now SIGMOD) Workshop on Data Description, Access and Control (SIGFIDET '74)*, pp. 321–350. ACM, New York. DOI: [10.1145/800296.811519](https://doi.org/10.1145/800296.811519) 404
- Z. Khayyat, I.F. Ilyas, A. Jindal, S. Madden, M. Ouzzani, P. Papotti, J.-A. Quiané-Ruiz, N. Tang, and S. Yin. 2015. Bigdancing: A system for big data cleansing. In *Proc. of the 2015 ACM SIGMOD International Conference on Management of Data (SIGMOD '15)*, pp. 1215–1230. ACM, New York. DOI: [10.1145/2723372.2747646](https://doi.org/10.1145/2723372.2747646). 297
- R. Kimball and M. Ross. 2013. *The Data Warehouse Toolkit*. John Wiley & Sons, Inc. <https://www.kimballgroup.com/data-warehouse-business-intelligence-resources/books/>. Last accessed March 2, 2018. 337
- M. Kornacker, C. Mohan, and J.M. Hellerstein. 1997. Concurrency and recovery in generalized search trees. In *Proc. of the 1997 ACM SIGMOD International Conference on Management of Data (SIGMOD '97)*, pp. 62–72. ACM, New York. DOI: [10.1145/253260.253272](https://doi.org/10.1145/253260.253272). 210
- A. Lamb, M. Fuller, R. Varadarajan, N. Tran, B. Vandiver, L. Doshi, and C. Bear. August 2012. The Vertica Analytic Database: C-Store 7 years later. *Proc. VLDB Endowment*, 5(12): 1790–1801. DOI: [10.14778/2367502.2367518](https://doi.org/10.14778/2367502.2367518). 333, 336
- L. Lamport. 2001. Paxos Made Simple. <http://lamport.azurewebsites.net/pubs/paxos-simple.pdf>. Last accessed December 31, 2017. 258
- D. Laney. 2001. 3D data management: controlling data volume, variety and velocity. META Group Research, February 6. <https://blogs.gartner.com/doug-laney/files/2012/01/ad949-3D-Data-Management-Controlling-Data-Volume-Velocity-and-Variety.pdf>. Last accessed April 22, 2018. 357
- P.-A. Larson, C. Clinciu, E.N. Hanson, A. Oks, S.L. Price, S. Rangarajan, A. Surna, and Q. Zhou. 2011. SQL server column store indexes. In *Proceedings of the 2011 ACM SIGMOD International Conference on Management of Data (SIGMOD '11)*, pp. 1177–1184. ACM, New York. DOI: [10.1145/1989323.1989448](https://doi.org/10.1145/1989323.1989448).
- J. LeFevre, J. Sankaranarayanan, H. Hacigumus, J. Tatemura, N. Polyzotis, and M. J. Carey. 2014. MISO: Souping up big data query processing with a multistore system. *Proc. of the 2014 ACM SIGMOD International Conference on Management of Data (SIGMOD '14)*, pp. 1591–1602. ACM, New York. DOI: [10.1145/2588555.2588568](https://doi.org/10.1145/2588555.2588568). 284
- B. G. Lindsay. 1987. A retrospective of R\*: a distributed database management system. In *Proc. of the IEEE*, 75(5): 668–673. DOI: [10.1109/PROC.1987.13780](https://doi.org/10.1109/PROC.1987.13780). 400
- B. Liskov and S.N. Zilles. 1974. Programming with abstract data types. *SIGPLAN Notices*, 9(4): 50–59. DOI: [10.1145/942572.807045](https://doi.org/10.1145/942572.807045). 88
- S. Marcin and A. Csillaghy. 2016. Running scientific algorithms as array database operators: Bringing the processing power to the data. *2016 IEEE International Conference on Big Data*. pp. 3187–3193. DOI: [10.1109/BigData.2016.7840974](https://doi.org/10.1109/BigData.2016.7840974). 350

- T. Mattson, V. Gadepally, Z. She, A. Dziedzic, and J. Parkhurst. 2017. Demonstrating the BigDAWG polystore system for ocean metagenomic analysis. *CIDR'17 Chaminade*, CA. <http://cidrdb.org/cidr2017/papers/p120-mattson-cidr17.pdf>. 288, 374
- J. Meehan, C. Aslantas, S. Zdonik, N. Tatbul, and J. Du. 2017. Data ingestion for the connected world. *Conference on Innovative Data Systems Research (CIDR'17)*, Chaminade, CA, January. 376
- A. Metaxides, W. B. Helgeson, R. E. Seth, G. C. Bryson, M. A. Coane, D. G. Dodd, C. P. Earnest, R. W. Engles, L. N. Harper, P. A. Hartley, D. J. Hopkin, J. D. Joyce, S. C. Knapp, J. R. Lucking, J. M. Muro, M. P. Persily, M. A. Ramm, J. F. Russell, R. F. Schubert, J. R. Sidlo, M. M. Smith, and G. T. Werner. April 1971. *Data Base Task Group Report to the CODASYL Programming Language Committee*. ACM, New York. 43
- C. Mohan, D. Haderle, B. Lindsay, H. Pirahesh, and P. Schwarz. 1992. ARIES: a transaction recovery method supporting fine-granularity locking and partial rollbacks using write-ahead logging. *ACM Transactions on Database Systems*, 17(1), 94–162. DOI: [10.1145/128765.128770](https://doi.org/10.1145/128765.128770). 402
- R. Motwani, J. Widom, A. Arasu B. Babcock, S. Babu, M. Datar, G. Manku, C. Olston, J. Rosenstein, and R. Varma. 2003. Query processing, approximation, and resource management in a data stream management system. *Proc. of the First Biennial Conference on Innovative Data Systems Research (CIDR)*, January. 229, 231
- A. Oloso, K-S Kuo, T. Clune, P. Brown, A. Poliakov, H. Yu. 2016. Implementing connected component labeling as a user defined operator for SciDB. *Proc. of 2016 IEEE International Conference on Big Data (Big Data)*. Washington, DC. DOI: [10.1109/BigData.2016.7840945](https://doi.org/10.1109/BigData.2016.7840945). 263, 350
- M. A. Olson. 1993. The design and implementation of the inversion file system. *USENIX Winter*. <http://www.usenix.org/conference/usenix-winter-1993-conference/presentation/design-and-implementation-inversion-file-syste>. Last accessed January 22, 2018. 215
- J. C. Ong. 1982. Implementation of abstract data types in the relational database system INGRES, Master of Science Report, Dept. of Electrical Engineering and Computer Sciences, University of California, Berkeley, CA, September 1982. 201
- A. Palmer. 2013. Culture matters: Facebook CIO talks about how well Vertica, Facebook people mesh. *Koa Labs Blog*, December 20. <http://koablog.wordpress.com/2013/12/20/culture-matters-facebook-cio-talks-about-how-well-vertica-facebook-people-mesh>. Last accessed March 14, 2018. 132, 133
- A. Palmer. 2015a. The simple truth: happy people, healthy company. *Tamr Blog*, March 23. <http://www.tamr.com/the-simple-truth-happy-people-healthy-company/>. Last accessed March 14, 2018. 138
- A. Palmer. 2015b. Where the red book meets the unicorn, *Xconomy*, June 22. <http://www.xconomy.com/boston/2015/06/22/where-the-red-book-meets-the-unicorn/> Last accessed March 14, 2018. 130
- A. Pavlo and M. Aslett. September 2016. What's really new with NewSQL? *ACM SIGMOD Record*, 45(2): 45–55. DOI: [10.1145/3003665.3003674](https://doi.org/10.1145/3003665.3003674). 246

- G. Press. 2016. Cleaning big data: most time-consuming, least enjoyable data science task, survey says. *Forbes*, May 23. <https://www.forbes.com/sites/gilpress/2016/03/23/data-preparation-most-time-consuming-least-enjoyable-data-science-task-survey-says/#79e14e326f63>. 357
- N. Prokoshyna, J. Szlichta, F. Chiang, R. J. Miller, and D. Srivastava. 2015. Combining quantitative and logical data cleaning. *PVLDB*, 9(4): 300–311. DOI: [10.14778/2856318.2856325](https://doi.org/10.14778/2856318.2856325). 297
- E. Ryvkina, A. S. Maskey, M. Cherniack, and S. Zdonik. 2006. Revision processing in a stream processing engine: a high-level design. *Proc. of the 22nd International Conference on Data Engineering (ICDE'06)*, pp. 141–. Atlanta, GA, April. IEEE Computer Society, Washington, DC. DOI: [10.1109/ICDE.2006.130](https://doi.org/10.1109/ICDE.2006.130). 228
- C. Saracco and D. Haderle. 2013. The history and growth of IBM's DB2. *IEEE Annals of the History of Computing*, 35(2): 54–66. DOI: [10.1109/MAHC.2012.55](https://doi.org/10.1109/MAHC.2012.55). 398
- N. Savage. May 2015. Forging relationships. *Communications of the ACM*, 58(6): 22–23. DOI: [10.1145/2754956](https://doi.org/10.1145/2754956).
- M. C. Schatz and B. Langmead. 2013. The DNA data deluge. *IEEE Spectrum Magazine*. <https://spectrum.ieee.org/biomedical/devices/the-dna-data-deluge>. 354
- Z. She, S. Ravishankar, and J. Duggan. 2016. BigDAWG polystore query optimization through semantic equivalences. *High Performance Extreme Computing Conference (HPEC)*. IEEE, 2016. DOI: [10.1109/HPEC.2016.7761584](https://doi.org/10.1109/HPEC.2016.7761584). 373
- SIGFIDET panel discussion. 1974. In *Proc. of the 1974 ACM SIGFIDET (now SIGMOD) Workshop on Data Description, Access and Control: Data Models: Data-Structure-Set Versus Relational (SIGFIDET '74)*, pp. 121–144. ACM, New York. DOI: [10.1145/800297.811534](https://doi.org/10.1145/800297.811534). 404
- R. Snodgrass. December 1982. Monitoring distributed systems: a relational approach. Ph.D. Dissertation, Computer Science Department, Carnegie Mellon University, Pittsburgh, PA, 197
- A. Szalay. June 2008. The Sloan digital sky survey and beyond. *ACM SIGMOD Record*, 37(2): 61–66. 255
- Tamr. 2017. Tamr awarded patent for enterprise-scale data unification system. Tamr blog. February 9 2017. <https://www.tamr.com/tamr-awarded-patent-enterprise-scale-data-unification-system-2/>. Last accessed January 24, 2018. 275
- R. Tan, R. Chirkova, V. Gadepally, and T. Mattson. 2017. Enabling query processing across heterogeneous data models: A survey. *IEEE Big Data Workshop: Methods to Manage Heterogeneous Big Data and Polystore Databases*, Boston, MA. DOI: [10.1109/BigData.2017.8258302](https://doi.org/10.1109/BigData.2017.8258302). 284, 376
- N. Tatbul and S. Zdonik. 2006. Window-aware Load Shedding for Aggregation Queries over Data Streams. In *Proc. of the 32nd International Conference on Very Large Databases (VLDB'06)*, Seoul, Korea. 228, 229

- N. Tatbul, U. Çetintemel, and S. Zdonik. 2007. "Staying FIT: Efficient Load Shedding Techniques for Distributed Stream Processing." *International Conference on Very Large Data Bases (VLDB'07)*, Vienna, Austria. 228, 229
- R. P. van de Riet. 1986. Expert database systems. In *Future Generation Computer Systems*, 2(3): 191–199, DOI: [10.1016/0167-739X\(86\)90015-4](https://doi.org/10.1016/0167-739X(86)90015-4). 407
- M. Vartak, S. Rahman, S. Madden, A. Parameswaran, and N. Polyzotis. September 2015. Seedb: Efficient data-driven visualization recommendations to support visual analytics. *PVLDB*, 8(13): 2182–2193. DOI: [10.14778/2831360.2831371](https://doi.org/10.14778/2831360.2831371). 297
- B. Wallace. June 9, 1986. Data base tool links to remote sites. *Network World*. [http://books.google.com/books?id=aBwEAAAAMBAJ&pg=PA49&lpg=PA49&dq=ingres+star&source=bl&ots=FSMIR4thMj&sig=S1fzaaOT5CHRq4cwbLFEQp4UYCs&hl=en&sa=X&ved=0ahUKEwj1J\\_NttvZAhUG82MKHco2CfAQ6AEIYZAP#v=onepage&q=ingres%20star&f=false](http://books.google.com/books?id=aBwEAAAAMBAJ&pg=PA49&lpg=PA49&dq=ingres+star&source=bl&ots=FSMIR4thMj&sig=S1fzaaOT5CHRq4cwbLFEQp4UYCs&hl=en&sa=X&ved=0ahUKEwj1J_NttvZAhUG82MKHco2CfAQ6AEIYZAP#v=onepage&q=ingres%20star&f=false). Last accessed March 14, 2018.305
- J. Wang and N. J. Tang. 2014. Towards dependable data repairing with fixing rules. In *Proc. of the 2014 ACM SIGMOD International Conference on Management of Data (SIGMOD '14)*, pp. 457–468. ACM, New York. DOI: [10.1145/2588555.2610494](https://doi.org/10.1145/2588555.2610494). 297
- E. Wong and K. Youssefi. September 1976. Decomposition—a strategy for query processing. *ACM Transactions on Database Systems*, 1(3): 223–241. DOI: [10.1145/320473.320479](https://doi.org/10.1145/320473.320479). 196
- E. Wu and S. Madden. 2013. Scorpion: Explaining away outliers in aggregate queries. *PVLDB*, 6(8): 553–564. DOI: [10.14778/2536354.2536356](https://doi.org/10.14778/2536354.2536356). 297
- Y. Xing, S. Zdonik, and J.-H. Hwang. April 2005. Dynamic load distribution in the Borealis Stream Processor. *Proc. of the 21st International Conference on Data Engineering (ICDE'05)*, Tokyo, Japan. DOI: [10.1109/ICDE.2005.53](https://doi.org/10.1109/ICDE.2005.53). 228, 230, 325

# Index

Page numbers with ‘f’ indicate figures; page numbers with ‘n’ indicate footnotes.

- 50-year perspective of Stonebraker
  - 1970 fall, University of Michigan, [107](#)–[108](#)
  - 1976 fall, Wisconsin, [108](#)–[111](#)
  - 1983 fall, Berkeley, [111](#)
  - 1988–1995, [111](#)–[112](#)
  - 2000, Project Sequoia, [112](#)–[113](#)
  - 2003, CIDR Conference launch, [113](#)
  - 2005, MIT sabbatical, [113](#)–[114](#)
  - 2008, MapReduce, [114](#)
  - 2014, Turing Award, [114](#)
  - 2016, MIT, [115](#)
  - 2017, encounter, [115](#)
- 1000 Genomes Browser, [267](#)
- 1 Million Veterans program, [354](#)
- Abadi, Daniel J., [56f](#)
  - C-Store project perspective article, [235](#)–[244](#)
  - end of architectural era, [463](#)–[489](#)
  - H-Store prototype, [247](#), [249](#)
  - OLTP databases, [411](#)–[439](#)
  - Vertica Systems, [334](#)
- Abstract data types (ADTs)
  - Ingres, [88](#), [148](#)
  - Ingres prototype, [202](#), [203f](#)
  - Postgres, [88](#), [149](#), [206](#)–[212](#), [316](#), [523](#)
- “Access Control in a Relational Database Management System By Query Modification” (Stonebraker and Wong), [45](#)
- Access methods
  - Ingres, [575](#)–[585](#), [577f](#), [582f](#)
  - Postgres, [534](#)–[535](#)
- Access Methods Interface (AMI), [582](#)–[585](#)
- ACM Software System Award, [44](#)
- Action enterprise, [310](#)
- Active databases in Postgres, [212](#)–[213](#)
- Active-passive replication in OLTP, [434](#)–[435](#)
- Addamark system, [494](#)
- Address space limitations in Ingres, [574](#)–[575](#)
- Administration files in Ingres, [576](#)
- ADMINS system, [404](#)
- ADT-Ingres, [206](#)
- Adult supervision for startup companies, [122](#)–[123](#)
- Advanced H-Store strategy, [478](#)
- Affero GPL license, [262](#)
- Aggregation operators in C-Store, [509](#)
- Aggregation systems in one size fits all, [451](#)–[452](#)
- AI systems
  - Ingres-oriented, [202](#)
  - machine learning, [66](#)



- Algorithmic complexity in Tamr, 359–361
- Allman, Eric, 99, 111, 195
- Allocator fragmentation in VoltDB, 342–343
- AllofUs program, 354
- Anchor tuples in H-Store, 474
- Anchored projections in C-Store, 497
- Anton, Jeff
  - Miro team, 314f
  - Postgres productionization, 315
- AntsDBMS, 478
- Anurag, Maskey, on Aurora project, 324f
- Aoki, Paul, 312–313n
- Apache Hadoop project, 4–5, 171
  - criticism of, 69, 184
  - open source impact, 171
- Apache HAWQ, 222
- Apache Spark, 402
- Application logic in one size fits all, 452–454, 458
- Area greater than (AGT) operator in Postgres, 525
- ARIES (Algorithms for Recovery and isolation Exploiting Semantics), 346
- Arizona State University, 149
- Array Functional Language (AFL) for SciDB, 267, 353–354
- Arrays
  - Postgres, 537
  - SciDB, 260–262, 261f
- AS/400 platform, 380
- Aster Data
  - founding, 222
  - Postgres parallelization, 209
- Audio data management, 265
- Aurora codelines and stream processing systems, 321–326, 324f
- Aurora project
  - founding, 46
  - origins, 225–227
  - research story, 387–391
  - StreamBase based on, 89
  - systems, 227–231
- Aurum system
  - Data Civilizer, 189
  - description, 299
- AUX directory in Ingres, 576
- Availability
  - OLTP design, 468–470
  - one size fits all, 454–455
- AVL trees, 436
- AWS Redshift, 222
- B-trees and B-tree indexes, 90
  - C-Store, 503
  - commercial Ingres codeline, 307
  - description, 493
  - OLTP, 429, 435–436
  - and Postgres, 210
- “B-Trees Re-examined” (Stonebraker and Held), 196–197
- Bachman, Charles
  - relational-CODASYL debate, 404, 406
  - Turing Award, 3, 48, 87
- Bailis, Peter, 159
- Balakrishnan, Hari
  - C-Store project, 238
  - StreamBase, 232, 329
- Balazinska, Magdalena
  - Aurora/Borealis/StreamBase reunion, 332f
  - Borealis project, 186
  - stream processing era article, 225–234
- Bates-Haus, Nikolaus, Tamr codeline article, 357–366, 365f
- Batkin, Adam, C-Store seminal work, 491–518
- Battle, Leilani, 159
- Bear, Chuck, 334–336
- Beaumont, Chris, 353
- Begoli, Edmon, 370
- Berkeley Municipal Court system, 148
- Berkeley Software Distribution (BSD) license
  - and Ingres, 166–167
  - origins, 165, 398
- Berkeley years
  - 1983 fall, 111

- technical contributions, 185–186
- BerkeleyDB in StreamBase, 453
- Berman, Richard, 99
- Bernstein, Philip A.
  - on leadership and advocacy, 87–92
  - relational database Ph.D, 403
- Berube, Paul, 374
- Beskales, George
  - Data Tamer project, 269–270, 273
  - Tamr co-founder, 360, 365f
  - Tamr company, 120, 274
- Betts, Ryan, 249
- BFGoodrich company, 97–98
- Bicycle story, 15–16
  - Anacortes, 16, 17f
  - Battle Lake, 26–27
  - Carrington, 26
  - difficult part, 79
  - Drake, 22–24, 23f
  - Ellicottville, 28
  - Luddington, 28
  - Marias Pass, 21, 21f
  - as metaphor for building system software, 32–35
  - Moultonborough, 35
  - Sutton, 31–32
  - Troy, 30
  - Winthrop, 18
  - Wollaston Beach, 31–32, 31f
- Big Data era
  - characteristics, 5–6
  - and Postgres, 209
  - stream processing in, 233–234
  - volume, velocity, and variety, 357
- BigDAWG codeline
  - future, 376
  - introduction, 367–370
  - milestones, 369, 369f
  - origins, 370–371, 371f
  - public demonstration, 371–373, 372f
  - refining, 373–374, 374f
  - release, 375–376
- BigDAWG polystore system
  - conclusion, 288–289
  - data movement, 286–287
  - description, 189
  - development, 284–285
  - ISTC, 279–280
  - one size does not fit all, 282–284
  - origins, 280–282, 281–282f
  - perspective on, 63–64
  - query modeling and optimization, 285–286
  - releases and demos, 287–288, 288f
- Biller, Steve, 374
- BIN directory in Ingres, 576
- Biobank program, 354
- Bioinformatics market, 263–267
- Bitstrings in C-Store, 509
- Blob storage in VoltDB, 343
- Blocking, 359
- Bochkov, Dmitry, 337
- Bohrbugs, 34
- Borealis codelines for stream processing systems, 321–326
- Borealis project
  - Aurora project move to, 325
  - origins, 225–227
  - systems, 227–231, 230f
- Bottleneck studies, 436
- Bowes, Ari, on Miro team, 314f
- Boyce, Bruce, 406
- Brooks, Fred, 223
- Brown, Paul
  - quad chart, 111
  - SciDB codeline, 352
  - scientific data management article, 253–268
- Bruckner, Daniel
  - Data Tamer project, 269–270, 273
  - Tamr company, 274, 365f
- Buffer management
  - OLTP, 414
  - Shore, 421, 423
- Buffer Pool Manager in Shore, 419
- Bulk copy of data in QUEL, 566
- Business acumen on startup company teams, 122–123

- Business-to-consumer (B2C) space, [150](#)
- Butterworth, Paul, commercial Ingres
  - codeline article, [303–310](#)
- C language in Postgres codelines, [311–312](#)
- C-Store project, [104](#)
  - column-oriented database, [132](#)
  - COMMIT statements, [506](#)
  - concat operators, [509](#)
  - covering sets of projections, [498](#)
- deletes, [503](#)
  - launch, [46–47](#)
  - one size doesn't fit all era, [187](#)
  - performance, [151](#)
- primary keys, [496](#)
  - prototype, [121](#)
  - Vertica Systems based on, [89](#)
- C-Store project perspective
  - Abadi and computer science, [235–238](#)
  - building, [240–242](#), [241f](#)
  - idea, evolution, and impact, [238–240](#)
  - Vertica Systems founding, [242–244](#)
- C-Store seminal work
  - abstract, [491–492](#)
  - conclusion, [516–517](#)
  - data model, [496–500](#), [499f](#)
  - introduction, [492–496](#), [495f](#)
  - performance, [511–515](#)
  - query execution, [509–510](#)
  - query optimization, [510](#)
  - related work, [515–516](#)
  - RS column store, [500–501](#)
  - storage management, [502–503](#)
  - tuple movers, [508](#)
  - updates and transactions, [503–508](#), [505f](#)
  - WS column store, [502](#), [507–508](#)
- Cache-conscious B-trees in OLTP, [435–436](#)
- Caches in rules systems, [546](#)
- Cafarella, Michael, [6](#)
- Caltech Mead-Conway VLSI design era, [201](#)
- Career flowchart, [54–55](#)
- Carey, Michael J., [56f](#)
  - data storage capabilities, [308](#)
  - Ingres later years article, [193–204](#)
  - Ingres project contributions, [185–186](#)
- Carnes, Donna, on Miro team, [314f](#)
- Carney, Don
  - Aurora project, [324f](#)
  - StreamBase Systems, [232](#)
- Carter, Fred, commercial Ingres codeline
  - article, [303–310](#)
- CASSM project, [109](#)
- Catalog relations in Ingres, [576](#)
- Çetintemel, Ugur
  - Aurora/Borealis/StreamBase reunion, [332f](#)
  - Aurora project, [323](#), [324f](#)
  - one size does not fit all, [282–284](#)
  - one size fits all seminal work, [441–462](#)
  - StreamBase Systems, [232](#)
- Chamberlin, Don
  - IBM Database, [382](#), [384](#)
  - relational-CODASYL debate, [406](#)
  - XQuery language, [208](#)
- Chen, Jolly
  - Postgres conversion, [149](#)
  - Postgres parser, [218](#)
  - PostgreSQL, [170](#)
  - SQLization project, [317–318](#)
- Chen, Peinan, [373](#)
- Cherniack, Mitch, [103](#)
  - Aurora/Borealis/StreamBase reunion, [332f](#)
  - Aurora project, [235](#), [323](#), [324f](#)
  - C-Store seminal work, [491–518](#)
  - expert sourcing, [273](#)
  - StreamBase Systems, [232](#)
  - Tamr project, [120](#)
  - Vertica Systems, [334](#)
- Chicken Test in Postgres productionization, [315](#)
- Chisholm, Sally (Penny), [373–374](#)
- Chisholm Laboratory data for BigDAWG
  - polystore system, [287](#), [374](#)
- Christiansen, Clayton, [27](#), [61](#), [100](#)
- CitusDB, [222–223](#)
- Classes, transaction, [471](#), [480–481](#)

- Climate change and Project Sequoia, [112–113](#)
- CLOS (Common LISP Object System), [522](#), [530–531](#)
- CLOSER function in Ingres, [585](#)
- Cloud Spanner, [70–71](#)
- Cloudera
  - and MapReduce, [69–70](#)
  - open source impact, [171](#)
- Cluster computing in OLTP, [416](#)
- CODASYL (Conference on Data Systems Languages)
  - Codd report, [43](#)
  - database standard proposal, [87](#), [97–98](#)
- Codd, Edgar (Ted)
  - Ingres development, [100](#)
  - Ingres inspired by, [43–44](#), [148](#)
  - Ingres platform, [101](#)
  - matrix algebra, [258](#)
  - relational-CODASYL debate, [404–406](#)
  - scientific data management, [263](#)
  - SCM SIGFIDET conference, [403](#)
  - Stonebraker influenced by, [42–44](#)
  - Turing Award, [3](#), [48](#)
- Cohera Corporation, [89](#), [401](#)
- Collected works of Stonebraker, [607–633](#)
- Column-oriented database, [132](#)
- Column store architecture, [239–240](#), [242](#), [492–493](#)
- Comeau, Al, [384](#)
- Commercial Ingres codeline, [303–304](#)
  - conclusion, [309](#)
  - open source Ingres, [309–310](#)
  - product production, [305–306](#)
  - research to commercial efforts, [304–305](#)
  - storage structures, [306–308](#)
  - user-defined types, [308–309](#)
- Commercialization
  - C-Store project, [242–244](#)
  - impact for, [237–238](#)
  - Postgres, [220–223](#)
- Commuting members, [473](#)
- Compaction in VoltDB, [342–344](#)
- Companies founded by Stonebraker, [48–49](#)
- Company control in startup company
  - guidelines, [127](#)
- Compatibility problem in one size fits all, [442](#)
- Complex objects
  - Ingres, [202–203](#)
  - Postgres, [207–208](#)
- Complexity
  - avoiding, [236–237](#)
  - rules systems, [541–543](#)
  - Vertica Systems, [338–339](#)
- Compression methods in C-Store project, [240–242](#), [241f](#)
- Computer Information and Control Engineering (CICE) program, [43](#)
- Computer science degrees
  - need for, [81–82](#)
  - University of Michigan, [107–108](#)
- Computer Systems Research Group (CSRG), [165–167](#)
- Concepts limitation in Postgres, [522](#)
- Concurrency control
  - C-Store, [505–506](#)
  - Distributed Ingres, [200](#)
  - H-Store, [478](#)
  - Ingres, [588–591](#)
  - OLTP, [433](#)
  - Postgres, [214n](#)
- CondorDB, [245](#)
- Conference on Innovative Data Systems Research (CIDR)
  - BigDAWG polystore system, [288](#)
  - creation, [50](#), [92](#)
  - launch, [113](#)
  - success of, [76](#)
- Connection points in Aurora, [229](#)
- Consistency
  - OLTP, [435](#)
  - QUEL, [567](#)
- Constrained tree applications (CTAs), [472](#)
- Constrained tree schema in H-Store project, [247](#)
- Constructed types in Postgres, [523](#)
- Control flow in Ingres, [575](#)

- Convey, Christian, on Aurora project, 324f
- Copyright for Ingres project, 109–110, 166
- CORBA, 112
- Correct primitives in one size fits all, 451–452
- Cost problem in one size fits all, 442
- CREATE command in Ingres, 600
- Credit, assigning, 394–395
- Cueball (partner), 36
- Current epochs in C-Store, 504
- Customers
  - forgotten, 158–159
  - startup companies, 122, 126
  - StreamBase, 327–328
  - Vertica Systems, 334–335, 339–340
- Cycles in OLTP, 432
- Data blades in Illustra, 88
- Data Civilizer, 291–292
  - conclusion, 300
  - data cleaning challenge, 296–298
  - data discovery challenge, 298–300
  - data transformation challenge, 295–296
  - description, 189
  - design, 294–295
  - life of analysts, 292–293
  - mung work automation, 64
  - need for, 292
  - purpose, 293–294
- Data cleaning challenge in Data Civilizer, 296–298
- Data discovery
  - Aurora project, 388
  - Data Civilizer, 298–300
- Data-driven discovery paradigm, 2
- Data ingest rates DBMS limitation, 226
- Data Language/ALPHA, 570
- Data Management Technology Kairometer, 3, 8
- Data movement in BigDAWG polystore system, 286–287
- Data structures in Ingres, 575–585, 577f, 582f
- Data Tamer project, 269–270
  - creation, 105
  - customers, 65
  - description, 47
  - idea and prototype, 270–273, 271f
  - ideas source, 152
  - lessons, 276–277
  - research contributions, 188
  - startup, 90
  - Tamr company, 273–275
- Data transformation challenge in Data Civilizer, 295–296
- Data unification at scale. *See* Data Tamer project
- Data warehouses
  - ideas source, 151
  - multiple data copies, 494
  - one size fits all, 443–445, 444–445f, 455–456
  - rise of, 401
  - schemas, 484
- Database designer in H-Store, 475
- Database Management Systems (DBMSs)
  - description, 3
- Databases, brief history of, 1–6
- DataBlades, 221, 317
- DataBlitz system, 436
- DATADIR directory in Ingres, 576
- Datalog rule systems, 212
- Date, Chris
  - referential integrity paper, 18
  - SCM SIGFIDET conference, 403
- David, Martin, 197
- DB2/400 system, 380
- Db2 for i system, 380
- Db2 for VSE&VM, 380
- Db2 for z/OS system, 380
- DB2/MVS system, 380
- DBA relations in Ingres, 576–577
- DBC project, 109
- Deadlocks in Ingres, 590
- Declarative query language, Ingres as, 196
- DECOMP program, 591–597

- Decompress operators in C-Store, [509](#)
- Deep storage technologies in Postgres, [215–216](#)
- Deferred update and recovery in Ingres, [600–602](#)
- DELETE function in Ingres, [584](#)
- Deleted record vectors (DRVs) in C-Store, [504](#)
- Densepack values in storage, [493](#), [501](#)
- Depth of transaction classes, [475](#)
- “Design and Implementation of Ingres” (Stonebraker), [47](#)
- Design problems in Postgres, [549–550](#)
- DeWitt, David J.
  - 50-year perspective article, [107–115](#)
  - CIDR, [92](#)
  - CondorDB version, [245](#)
  - Gamma, [151](#), [400](#)
  - H-Store project, [246](#)
  - Hadoop criticism, [184](#)
  - one size doesn’t fit all era, [188](#)
  - publications, [76](#)
  - Shore prototype, [152](#)
  - Vertica Systems, [136](#), [336](#)
- DIRECT project, [109–111](#)
- Disk orientation DBMS limitation, [226](#)
- Disk persistence in VoltDB, [346–347](#)
- Distinct values in C-Store, [500–501](#)
- Distributed COMMIT processing in C-Store, [506](#)
- Distributed computing, [400](#)
- Distributed databases, [150](#)
- Distributed Ingres, [198–200](#), [199f](#), [305](#)
- Docker tool for BigDAWG, [375](#)
- Document data management in Ingres, [201](#)
- Dozier, Jeff, [112](#)
- Du, Jiang, [376](#)
- Dynamic loading in Postgres, [551–552](#)
- Dynamic locking in H-Store, [477](#)
- Dziedzic, Adam, [373](#)
- Early years and education, [42–43](#)
- Elephants, [4](#), [94](#)
- Ellicott, Andy, [134](#)
- Ellison, Larry
  - Oracle claims, [149](#)
  - SQL language support, [196](#)
- Elmore, Aaron J., BigDAWG polystore system article, [279–289](#)
- Emberson, Richard, on Miro team, [314f](#)
- EMP1 (friend), [27](#)
- Encoding schemes in C-Store, [500–501](#)
- “End of an Architectural Era: It’s Time for a Complete Rewrite” paper (Stonebraker), [247](#)
- End of architectural era seminal work
  - abstract, [463–464](#)
  - H-Store, [473–479](#), [479f](#)
  - introduction, [464–466](#)
  - OLTP design considerations, [466–470](#)
  - one size does not fit all comments, [483–486](#)
  - performance, [479–483](#), [479f](#)
  - summary and future work, [486–487](#)
  - transaction, processing and environment assumptions, [470–473](#)
- End of epochs in C-Store, [504](#)
- End-to-end system Data Civilizer design, [294–295](#)
- EnterpriseDB, [222](#)
- Entrepreneur-Turned-Shark (friend), [28](#)
- Epochs in C-Store, [504–505](#), [505f](#)
- Epstein, Bob
  - BSD license, [166](#)
  - Distributed Ingres, [198](#)
  - Ingres source code, [110–111](#)
  - stored procedures, [91](#)
  - venture capitalist contact, [140](#)
- EQUEL language for Ingres
  - comments, [570–571](#)
  - invocation from, [573–574](#)
  - overview, [568–570](#)
- Erwin, Christina, on Aurora project, [324f](#)
- ETL toolkit, [358](#)
- Exceptions in rules systems, [539](#)
- Excessive formalism, [162](#)

- Expanding fields, failure to cope with, 155–158, 156f
- Expansive execution in BigDAWG polystore system, 285
- Experimental results for OLTP, 428–432
- Expert database systems, Ingres-oriented, 202
- Expert sourcing, 273
- Explicit parallelism in Tamr, 363
- Extremely Large Data Bases (XLDB)
  - conference and workshop, 256–257
- Fabry, Bob, 165
- Factoring in one size fits all, 458–459, 459f
- Failover
  - OLTP design, 469
  - one size fits all, 454–455
- Failures
  - consequences, 160–163
  - expanding fields, 155–158, 156f
  - forgotten customers, 158–159
  - paper deluge, 159–160
  - summary, 164
- Fast path feature in Postgres, 218, 316, 530–531
- Federation architecture, 401
- Female Ph.D.s graduated, 394
- Fernandez, Raul Castro
  - Aurum research story, 387–391
  - Data Civilizer, 189
  - Data Civilizer article, 291–300
- Festschrift, 143–144
- Files
  - Ingres, 576–578, 577f
  - UNIX environment, 571
- Financial-feed processing in one size fits all, 446–447, 448f
- FIND function in Ingres, 583–584
- First customers for startup companies, 126
- First International Conference on Expert Database Systems, 407–408
- Fogg, Dennis, 206
- Ford, Jim, 99
- Foreign keys in C-Store, 496
- Foreign-order values in C-Store, 501
- Forgotten customers, 158–159
- Fork-lift upgrades in OLTP design, 468
- Fournier, Marc, 318
- Franklin, Michael J., 56f
  - papers rejected by, 113
  - Telegraph Team, 325
- Frew, Jim, 112
- Freytag, Christoph, 173
- Functional ideas, 184
- Functions
  - Postgres, 208–209, 211, 523–527
  - POSTQUEL, 535–536, 555–556
- Gadepally, Vijay
  - BigDAWG codeline article, 367–376
  - BigDAWG releases, 287
- Galvez, Eddie
  - Aurora project, 324f
  - StreamBase Systems, 232, 330
- Gamma project, 111, 151, 400
- Garlic project, 401
- Gates, Bill, 215n
- GEM language, 217
- Generalized Search Tree (GiST) interface, 210–211
- Genomic data for SciDB, 354–355
- Geographic Information Systems (GIS), 148
- GET function in Ingres, 583
- Gigascope project, 231
- Global Biobank Engine, 267
- Go, Angela, 148
- Goby company
  - B2C space, 150
  - Data Tamer, 152
  - startup, 90
- Google technologies, 70–71
- Goss, Jeff, 385f
- Gosselin, Dave, 352–353
- Governor's Academy, 43
- GPUs, 67
- Graduate students, 82
- Graphical user interfaces (GUIs) in
  - prototypes, 121

- Grassy Brook company
  - founding, [140–141](#)
  - quad chart, [230–232](#), [231f](#)
- Gray, Jim, [36](#)
  - 2002 SIGMOD conference, [113](#)
  - CIDR, [92](#)
  - Project Sequoia, [112](#)
  - scientific data management, [255](#), [259](#)
  - System R project, [100](#)
  - Tandem Computers, [101](#)
  - Turing Award, [3](#), [48](#), [114](#)
- Great Relational-CODASYL Debate, [403–406](#)
- Greenplum startup, [209](#), [221–222](#)
- Grid computing in OLTP design, [468](#)
- Gupta, Ankush, [373](#)
- Guttman, Antonin
  - Ingres CAD management features, [201–202](#)
  - R-Tree index structure, [205](#), [210](#)
- H-Store, [245–246](#)
  - basic strategy, [478](#)
  - buddies, [475](#)
  - conclusion, [251](#)
  - database designer, [475](#)
  - description, [47](#)
- execution supervisor, [475](#)
  - founding, [104](#)
- general transactions, [247](#)
  - ideas source, [151–152](#)
  - one size doesn't fit all era, [187](#)
  - performance, [479–483](#), [479f](#)
  - prototypes, [247–250](#)
  - query execution, [474–475](#)
  - system architecture, [246–247](#), [473–474](#)
  - transaction classes, [471–473](#)
  - transaction management, replication and recovery, [476–478](#)
  - VoltDB and PayPal, [250](#)
  - VoltDB based on, [89](#)
  - VoltDB executor in, [75](#)
  - VoltDB split, [251](#)
- Hachem, Nabil
  - Data Civilizer, [295](#)
  - end of architectural era seminal work, [463–489](#)
- Haderle, Don, recollections article, [397–402](#)
- Hadoop, [4–5](#)
  - criticism of, [69](#), [184](#)
  - open source impact, [171](#)
- Hadoop Distributed File System (HDFS ), [70](#)
- Hagen, Dale, [384](#), [385f](#)
- Hamilton, James
  - on 2014 ACM Turing Award, [93–95](#)
  - IBM relational database code bases article, [379–385](#)
  - on server costs, [67–68](#)
- Hammer, Joachim, [90](#)
- Hanson, Eric, [212](#)
- Harizopoulos, Stavros
  - end of architectural era seminal work, [463–489](#)
  - H-Store prototype, [247](#)
  - OLTP databases, [411–439](#)
- Harris, Herschel, [385f](#)
- HASH structure in commercial Ingres codeline, [306–307](#)
- Hatoun, Matt, on Aurora project, [324f](#)
- Hawthorn, Paula
  - Illustra, [221](#)
  - Miro team, [314f](#)
  - Postgres productionization, [315](#)
- Hearst, Marti, student perspective article, [393–396](#)
- Heath, Bobbi
  - H-Store prototype, [249](#)
  - StreamBase Systems, [232](#)
- Hedges, Richard, [384](#)
- Heisenbugs, [34](#)
- Held, Gerald
  - “B-Trees Re-examined,” [194–197](#)
  - Ingres implementation seminal work, [561–605](#)
  - relational database industry birth, [97–106](#)
- Helland, Pat, end of architectural era seminal work, [463–489](#)
- Hellerstein, Joseph M.
  - Data Tamer, [152](#)



- Hellerstein, Joseph M. (*continued*)
  - Postgres codelines, 313
  - Postgres description, 186
  - Postgres perspective article, 205–224
  - Postgres project, 102
  - Tamr project, 120
  - Telegraph Team, 325
- Heroku provider, 219–220
- High availability
  - OLTP, 417, 468–470
  - one size fits all, 454–455
- High water mark (HWM) in C-Store, 504–505, 505f
- Hill, Faith, 294
- Hints in Postgres, 525
- HiPac project, 213
- Hirohama, Michael, Postgres implementation seminal work, 519–559
- Historical mode queries, 495
- Hive executor, 114
- Hobbib, Bill, 322, 326
- Hong, Wei
  - Illustra, 206, 221
  - Miro team, 314f
  - Postgres and Illustra codelines article, 311–319
  - Postgres conversion, 149
  - XPRS architecture, 216–217
- Horizontica version, 249, 331
- Hot standby in OLTP design, 469
- Howe, Bill, 370
- HTCondor project, 245
- Hugg, John
  - H-Store prototype, 249–250
  - VoltDB codeline article, 341–348
- Huras, Matt, 383–384, 385f
- Hwang, Jeong-Hyon, on Aurora project, 324f
- Hypothetical relations in Ingres, 201
- IBM
  - IMS database, 88
  - SQL language, 102
- IBM relational database code bases
  - four code bases, 379–381
  - future, 384–385
  - portable code base, 381–384
- IEEE International Conference on Data Engineering 2015 talk, 35
- Illustra codelines, 311
  - overview, 313–317
  - Postgres and SQL, 315–316
  - Postgres productionization, 315
- Illustra Information Technologies, Inc.
  - open source impact, 171
  - Oracle competitor, 102–103
  - Postgres, 112
  - Postgres commercial adaptations, 220–221
  - startup, 206
- Ilyas, Ihab
  - Data Tamer project article, 269–277
  - Tamr co-founder, 120, 188, 360, 365f
- Implementation efficiency in rules systems, 544
- “Implementation of Integrity Constraints and Views By Query Modification” (Stonebraker), 45
- “Implementation of Postgres” (Stonebraker), 47
- Implementing rules, 91
- IMS database, 88
- In-QTel, 233
- Inbound vs. outbound processing in one size fits all, 449–450f, 449–451
- INDEX catalog in Ingres, 580
- Indexes
  - C-Store, 498, 499f, 501, 509
  - commercial Ingres, 306–307
  - Postgres, 525, 550
  - primary and secondary, 493
  - VoltDB, 343
- Industrial Liaison Program (ILP), 121
- Industry involvement, 46
- InfiniBand, 67
- Informix
  - Illustra integrated into, 35
  - Illustra purchase, 102–103, 112
  - startups bought by, 66

- Informix Universal Server, [317](#)
- Ingres implementation seminal work, [561–562](#)
  - Access Methods Interface, [582–585](#)
  - conclusion, [602–603](#)
  - concurrency control, [588–591](#)
  - data structures and access methods, [575–585](#), [577f](#), [582f](#)
  - deferred update and recovery, [600–602](#)
  - EQUEL, [568–571](#)
  - file structure, [576–578](#), [577f](#)
  - future extensions, [603](#)
  - introduction, [562–563](#)
  - invocation from EQUEL, [573–574](#)
  - invocation from UNIX, [572–573](#)
  - performance, [602](#)
  - Process 2, [585–591](#)
  - Process 3, [591–599](#)
  - Process 4, [599–602](#)
  - process structure, [571–575](#), [572f](#), [574f](#)
  - QUEL and utility commands, [563–568](#)
  - query modification, [585–588](#)
  - storage structures, [580–582](#), [582f](#)
  - system catalogs, [578–580](#)
  - UNIX environment, [571–572](#)
  - user feedback, [602–603](#)
- Ingres later years, [193–194](#)
  - contributions, [195–198](#)
  - Distributed Ingres, [198–200](#), [199f](#)
  - relational DBMS, [194–198](#), [197f](#)
  - support domains, [201–204](#)
- Ingres project
  - ATTRIBUTE catalog, [579–580](#)
  - Berkeley years, [185–186](#)
  - birth of, [43–45](#)
  - and BSD, [166–167](#)
  - commercial codeline. *See* Commercial Ingres codeline
  - competition, [100–103](#)
- COPY command, [600](#)
  - copyright, [109–110](#)
  - decomposition of queries, [591–597](#)
  - distributed, [150](#)
  - ideas source, [147–148](#)
  - impact, [3](#)
  - leadership and advocacy, [87–90](#)
  - open source, [167–168](#), [309–310](#)
  - platform, [101](#)
  - Postgres design helped by, [27–28](#)
  - process structure, [110–111](#)
  - target platform, [108](#)
  - team, [99–100](#)
  - timing, [98–99](#)
  - Wisconsin, fall 1976, [108–111](#)
- Ingres Star, [305](#)
- Inheritance in Postgres, [524](#), [529](#)
- Innovators Dilemma (Christiansen), [27](#), [61](#), [100](#)
- INSERT function in Ingres, [584](#)
- Insertion vectors (IVs) in C-Store, [504](#)
- Inserts in C-Store, [503](#)
- Instructions vs. cycles in OLTP, [432](#)
- INTEGRITY catalog in Ingres, [580](#)
- Integrity control in QUEL, [567](#)
- Intel Science and Technology Center (ISTC)
  - for Big Data, [279](#), [367](#)
- Intellectual property, [126](#)
- Inter-snapshot log in VoltDB, [346](#)
- Intermediate H-Store strategy, [478](#)
- Inversion file system, [215](#)
- Irrelevant theories, [161–162](#)
- ISAM (indexed sequential access method), [306–307](#)
- Islands in BigDAWG polystore system, [284](#)
- Join indexes in C-Store, [498](#), [499f](#), [501](#), [509](#)
- Join operators in C-Store, [509](#)
- Jones, Anita, [193](#)
- Jones, Evan, [248–249](#)
- Joy, Bill, [109](#), [165](#)
- JSON data model, [208](#)
- K-safe systems, [494](#), [499–500](#)
- Katz, Randy, [216](#)
- KDB system, [494](#)
- Kelley, Gary, [151](#)
- Kepner, Jeremy, [370](#)
- Kerschberg, Larry, [407](#)

- Kersten, Martin, [75](#), [151](#)
- Keyed storage structure in Ingres, [581](#)–[582](#), [582f](#)
- Keys
  - C-Store, [496](#)–[498](#)
  - Postgres, [550](#)
- Kimura, Hideaki, [247](#)–[249](#)
- Kinchen, Jason, SciDB codeline article, [349](#)–[355](#)
- KISS adage (Keep it Simple, Stupid) adage, [153](#)
- Knowledge management in Postgres, [524](#)
- Knudsen, Eliot, [360](#)
- Kooi, Bob, [306](#)
- Kraska, Tim, [67](#)
- Kreps, Peter
  - Ingres implementation seminal work, [561](#)–[605](#)
  - Ingres team, [99](#)
- Land Sharks, [15](#)–[16](#), [29](#)
- Langer, Robert, [137](#)
- Language constructs as DBMS limitation, [226](#)
- Language support in Postgres, [217](#)–[218](#)
- Large objects in Postgres, [537](#)
- Large Synoptic Survey Telescope, [256](#)
- Latching
  - H-Store, [483](#)
  - OLTP, [414](#)
  - Shore, [419](#)–[420](#), [423](#)
- Latency in VoltDB, [344](#)–[348](#)
- Lau, Edmond, C-Store seminal work, [491](#)–[518](#)
- Lawande, Shilpa
  - Vertica Systems, [134](#), [136](#)
  - Vertica Systems codeline article, [333](#)–[340](#)
- Leadership, partnership approach to, [131](#)–[132](#)
- Leadership and advocacy, [87](#)
  - advocacy, [91](#)–[92](#)
  - mechanisms, [90](#)–[92](#)
  - systems, [87](#)–[90](#)
- Least Publishable Units (LPUs)
  - grain optimization, [389](#)
  - problems from, [76](#), [159](#), [161](#)
- Leibensperger, Mike, [354](#)
- Lew, Ivan, [385f](#)
- Lewis, Bryan, [350](#), [353](#)
- licenses, BSD, [165](#), [398](#)
- Lighthouse customers for startup companies, [122](#)
- Lightstone, Sam, [385f](#)
- Lindsay, Bruce, [382](#), [382](#)–[384](#)
- Lineage support in scientific data management, [266](#)
- Linear Road benchmark, [326](#), [446](#)
- Liquidation preference in startup company guidelines, [124](#)–[125](#)
- Liskov, Barbara, [88](#), [206](#)
- Lisman, John, [235](#)
- LISP for Postgres, [311](#)–[312](#), [553](#)–[554](#)
- Liu, Jason, at Tamr, [365f](#)
- Lock Manager in Shore, [419](#)
- Locking
  - C-Store, [503](#), [505](#)–[506](#)
  - H-Store, [477](#), [483](#)
  - Ingres, [590](#)–[591](#)
  - OLTP, [414](#)
  - performance evaluation of, [90](#)–[91](#)
  - Shore, [419](#)–[421](#), [423](#)
- Log-centric storage and recovery in Postgres, [213](#)–[215](#)
- Log Manager in Shore, [419](#)
- Logical data model in C-Store, [496](#)
- Logless databases, [413](#)
- Logos and T-shirts for projects, [395](#), [395f](#)
- Logs and logging
  - H-Store, [482](#)
  - OLTP, [414](#), [417](#)
  - for recovery purposes, [21](#)–[22](#)
  - redo, [470](#)
  - Shore, [419](#)–[421](#), [423](#)
  - undo, [471](#)
  - VoltDB, [346](#)
- Lohman, Guy, [382](#), [384](#)
- Lorie, Raymond, [113](#)

- Low water mark (LWM) in C-Store, [504](#)
- LSM-tree concept, [495](#)
- Lucking, J. R., [404–405](#)
- MacAIMS Data Management System, [404](#)
- MacDonald, Nancy, [99](#)
- Machine learning, [66](#)
- Madden, Samuel, [56f](#)
  - BigDAWG, [370](#)
  - C-Store project, [238](#), [240](#)
  - C-Store seminal work, [491–518](#)
  - end of architectural era seminal work, [463–489](#)
  - Festschrift, [143](#)
  - H-Store prototype, [247](#)
  - ISTC, [279](#), [281](#)
  - OLTP databases seminal work, [411–439](#)
  - research contributions article, [183–189](#)
  - Vertica Systems, [334](#)
- Madden, Samuel, on Stonebraker
  - academic career and birth of Ingres, [43–45](#)
  - advocacy, [50](#)
  - awards and honors, [49](#)
  - companies founded, [48–49](#)
  - early years and education, [42–43](#)
  - industry, MIT, and new millennium, [46–47](#)
  - legacy, [47–48](#)
  - personal life, [50](#)
  - post-Ingres years, [45–46](#)
  - synopsis, [41–42](#)
- MADlib library, [221–222](#)
- Mahony, Colin, [134](#)
- Maier, David, [176](#), [325](#)
- Main memory
  - OLTP design, [466–467](#)
  - studies, [436](#)
- “Making Smalltalk a Database System” (Copeland and Maier), [111](#)
- MapReduce
  - blog post, [114](#)
  - criticism of, [5](#), [68–70](#), [136](#)
  - and Postgres, [209](#)
- Mariposa system
  - description, [88–89](#)
  - federation architecture, [401](#)
  - prototype, [150](#)
- Mark, Roger, [370](#)
- Marketing problem in one size fits all, [442](#)
- MARS system, [436](#)
- Mask operators in C-Store, [509](#)
- Mattson, Tim, BigDAWG polystore system
  - article, [279–289](#)
- McCline, Matt, [215n](#)
- McKnight, Kathy, [384](#)
- McPherson, John, [382](#), [384](#)
- McQueston, James, [349](#), [351–353](#), [355](#)
- MDM (master data management), [358](#)
- Meehan, John, [376](#)
- Memory
  - OLTP design, [466–467](#)
  - studies, [436](#)
- Memory resident databases in OLTP, [416](#)
- Merck databases, [64–65](#)
- Meredith, Jeff
  - Illustra, [206](#), [221](#)
  - Miro team, [314f](#)
  - Postgres, [312](#), [314](#)
- Merge-out process, [495](#), [508](#)
- Message transport in one size fits all, [458](#)
- Method and System for Large Scale Data Curation patent, [275](#)
- MIMIC (Multiparameter Intelligent Monitoring in Intensive Care)
  - dataset, [370–371](#), [371f](#)
- Miro team, [314f](#)
- Miso system, [284](#)
- “Mission to Planet Earth” (MTPE) effort, [112–113](#)
- Mistakes in startup company guidelines, [128](#)
- MIT
  - 2005 sabbatical, [113–114](#)
  - 2016, [115](#)
  - Aurora and StreamBase projects, [46](#)
  - Industrial Liaison Program, [121](#)
  - research contributions, [186](#)

- MIT CSAIL, [103](#)
- MODIFY command in OVQP, [599–600](#)
- Mohan, C., [382](#), [384](#)
- Mom (friend), [27](#)
- MonetDB project, [113](#), [151](#)
- Morgenthaler, Gary, on Miro team, [314f](#)
- Morpheus project
  - description, [47](#)
  - prototype, [150](#)
  - startup, [90](#)
- Morris, Barry
  - Aurora/Borealis/StreamBase reunion, [332f](#)
  - StreamBase Systems, [232](#)
- Mucklo, Matthew, [376](#)
- Muffin parallel databases, [151](#)
- MUFFIN prototype, [200](#)
- Multi-core support in OLTP, [434](#)
- Multi-threading in OLTP design, [467](#)
- Multilingual access in Postgres, [522](#)
- Myria project, [284](#)
- MySQL, [137](#)
  
- Nakerud, Jon, [149](#)
- NASA “Mission to Planet Earth” effort, [112–113](#)
- National Science Foundation (NSF)
  - proposal success rate, [74](#)
  - RANN program, [147–148](#)
- Naughton, Jeff, [113](#)
- Naumann, Felix, RDBMS genealogy article, [173–179](#)
- Navigational era, [3](#)
- Naylor, Arch, [107–108](#)
- Nested queries in Postgres, [528](#)
- Netezza startup, [221](#)
- Network Time Protocol (NTP) for VoltDB, [344–345](#)
- New Order transactions in OLTP, [426](#), [430–432](#), [431–432f](#)
- NewSQL architecture, [246](#), [282f](#)
- No-overwrite storage manager in Postgres, [548](#)
- Non-volatile RAM, [67](#)
- Nonkeyed storage structure in Ingres, [581](#)
- Nonrepeatable errors, [34](#)
- Normal functions in Postgres, [524](#)
- NoSQL systems, [246](#)
- Novartis Institute for Biomedical Research (NIBR), [121](#), [152](#), [361](#)
  
- Object identifiers (OIDs) in Postgres, [532](#)
- Object Management Extension in
  - commercial Ingres, [308](#)
- Object management in Postgres implementation, [520](#)
- “Object Management in Postgres Using Procedures” (Stonebraker), [45–46](#)
- Object-orientation in Postgres, [206–207](#), [531–532](#)
- Object-Oriented Databases (OODBs), [111–112](#), [217–218](#)
- Object-Relational DBMSs: Tracking the Next Great Wave (Stonebraker and Brown), [111](#)
- Object-Relational model, [186](#)
- O’Brien, Kyle, [375](#)
- O’Connell, Claire, [360](#)
- Olson, Mike
  - Illustra, [221](#)
  - Inversion file system, [215](#)
  - open source article, [165–171](#)
  - Postgres B-tree implementation, [214n](#)
  - Postgres codelines, [313](#)
- OLTP (Online Transaction Processing)
  - applications in H-Store project, [246–247](#)
- OLTP (Online Transaction Processing)
  - databases seminal work
    - abstract, [411–412](#)
    - alternative DBMS architectures, [413](#)
    - cache-conscious B-trees, [435–436](#)
    - conclusion, [436–437](#)
    - concurrency control, [433](#)
    - contributions and paper organization, [414–415](#)
    - experimental results, [428–432](#)
    - future engines, [433–436](#)

- instructions vs. cycles, [432](#)
- introduction, [412–415](#)
- multi-core support, [434](#)
- New Order transactions, [430–432](#), [431–432f](#)
- overheads, [414](#)
- payment, [429–430](#), [429f](#)
- performance study, [424–432](#)
- related work, [436](#)
- replication management, [434–435](#)
- results, [414–415](#), [415f](#)
- setup and measurement methodology, [427](#)
- Shore, [418–424](#)
- throughput, [428](#)
- trends, [416–418](#)
- weak consistency, [435](#)
- OLTP (Online Transaction Processing)
  - design considerations
  - grid computing and fork-lift upgrades, [468](#)
  - high availability, [468–470](#)
  - knobs, [470](#)
  - main memory, [466–467](#)
  - multi-threading and resource control, [467](#)
  - payment transactions, [429–430](#), [429f](#)
- “OLTP: Through the Looking Glass” paper (Harizopoulos), [152](#)
- One-shot applications, [472](#)
- One-shot transactions, [247](#), [474](#), [476](#)
- One size does not fit all
  - BigDAWG polystore system, [282–284](#)
  - in end of architectural era seminal work, [483–486](#)
  - overview, [4–5](#)
  - research contributions, [187–188](#)
  - special-purpose database systems, [103](#)
- One size fits all: An idea whose time has come and gone seminal work
  - abstract, [441](#)
  - conclusion, [460](#)
  - correct primitives, [451–452](#)
  - data warehouses, [443–445](#), [444–445f](#), [455–456](#)
  - DBMS processing and application logic integration, [452–454](#), [453f](#)
  - factoring, [458–459](#), [459f](#)
  - financial-feed processing, [446–447](#), [448f](#)
  - high availability, [454–455](#)
  - inbound versus outbound processing, [449–450f](#), [449–451](#)
  - introduction, [441–442](#)
  - performance, [448–455](#)
  - scientific databases, [457](#)
  - sensor-based applications, [445–446](#)
  - sensor networks, [456](#)
  - stream processing, [445–447](#), [448f](#)
  - synchronization, [455](#)
  - text search, [457](#)
  - XML databases, [457–458](#)
- One-variable detachment in Ingres, [592–593](#)
- One-Variable Query Processor (OVQP) in Ingres, [597–599](#)
- O’Neil, Pat, C-Store seminal work, [491–518](#)
- Ong, James, [206](#)
- Open source
  - BSD and Ingres, [166–167](#)
  - BSD license, [165](#)
  - Ingres impact, [167–168](#)
  - open source Ingres, [309–310](#)
  - post-Ingres, [168–169](#)
  - Postgres, [218–220](#)
  - PostgreSQL, [318–319](#)
  - research impact, [169–171](#)
- OPENR function in Ingres, [583](#)
- “Operating System Support for Data Management” (Stonebraker and Kumar), [47](#)
- Operators
  - C-Store queries, [509–510](#)
  - Postgres, [525–526](#), [528](#)
  - scientific data management, [262–263](#)
- Optimizations in Shore, [423](#)
- OQL language, [208](#)
- Oracle Corporation
  - competition with, [102–103](#)

- Oracle Corporation (*continued*)
  - performance claims, 149
  - Postgres attack by, 209
  - Tamr, 363–364
- Orca optimizer, 221
- OS/2 Database Manager, 381–383
- OS/2 system, 381–383
- Ousterhout, John, 216
- Ouzzani, Mourad
  - Data Civilizer, 189
  - Data Civilizer article, 291–300
- “Over My Dead Body” issues in StreamBase, 328–329
- Overheads in OLTP, 414
- Pagan, Alexander
  - Data Tamer project, 269–270, 273
  - Tamr company, 120, 274, 365f
- Palmer, Andy, 104
  - 2014 Turing Award Ceremony, 130f
  - “Cue Ball”, 144
  - Data Tamer project, 269
  - Festschrift, 143
  - startup company article, 129–138
  - Tamr CEO, 105, 123, 274
  - Tamr company, 365f
  - Vertica Systems, 142–143, 242–243
- Paper deluge, 159–160
- Paper requirements, 159–161
- ParAccel, 222
- Paradigm4, 89
- Paradise project, 113
- Parallel databases ideas source, 151
- Parallel Data Warehouse project, Microsoft, 47
- Parallel DBMS in Postgres, 316
- Parallel Sysplex, 400
- Parallelism in Tamr, 363
- PARAMD function in Ingres, 584
- PARAMI function in Ingres, 584
- Parsers in Ingres, 586
- Partitions in C-Store, 498
- Partnerships
  - leadership approach, 131–132
  - startup companies, 132–135
- Partridge, John
  - Aurora/Borealis/StreamBase reunion, 332f
  - Aurora language, 227n
  - connection points, 229n
  - RFID tagging, 225n
  - StreamBase customers, 327–328
  - StreamBase founding, 232
  - StreamBase issues, 328–329
- Past data access as DBMS limitation, 226
- Patents, 126
- Path expressions in Postgres, 528
- Patterson, Dave, 216
- Pavlo, Andrew
  - H-Store project, 187
  - H-Store project article, 245–251
  - VoltDB executor in H-Store, 75
- PayPal, 250
- Pearson Correlation Coefficient (PCC), 292
- People for startup companies, 138
- Performance
  - BigDAWG polystore system, 287, 288f
  - bottleneck studies, 436
  - C-Store, 151, 511–515
  - Data Unification, 358
  - H-Store, 479–483, 479f
  - Ingres, 602
  - locking methods, 90–91
  - OLTP, 424–432
  - one size fits all, 448–455
  - Postgres, 30–32, 549–550, 554–555, 555f
- Permute operators in C-Store, 509
- Perna, Janet, 382, 384, 401
- Persistence
  - Postgres, 522
  - VoltDB, 346–347
- Persistent CLOS, 522
- Persistent redo logs, 470
- Personal life of Stonebraker, 50
- Peterlee IS/1 System, 404
- Ph.D. paper requirements, 159
- Pipes, 572
- Pirahesh, Hamid, 382, 384

- Pitch decks in startup company guidelines, [123–125](#)
- Pivotal company, [221–222](#)
- Plans for C-Store queries, [509–510](#)
- Poliakov, Alex, SciDB codeline article, [349–355](#)
- Polybase system, [284](#)
- Polystores. *See* BigDAWG polystore system
- Portable code base for IBM relational databases, [381–384](#)
- Post-Ingres years
  - open source, [168–169](#)
  - overview, [45–46](#)
- Postgres codelines, [311](#)
  - conclusion, [319](#)
  - PostgreSQL, [317–319](#)
  - prototype, [311–313](#)
- Postgres design, [15–16](#)
  - base typess, [523](#)
  - bicycle trip metaphor, [32–35](#)
  - conclusion, [35–36](#)
  - Illustra buyout, [32](#)
  - inheritance, [524](#), [529](#)
  - Ingres help for, [27–28](#)
  - Internet takeoff, [30](#)
  - Land Sharks, [29](#)
  - marketing challenge, [28–29](#)
  - performance benchmark, [30–32](#)
  - speedbumps, [24–26](#)
  - start, [15–22](#), [22f](#)
- Postgres implementation seminal work
  - abstract, [519](#)
  - conclusion, [555–557](#)
  - data model, [523–528](#)
  - data model and query language overview, [521–523](#)
  - data model critique, [532–537](#)
  - design problems, [549–550](#)
  - dynamic loading and process structure, [551–552](#)
  - fast path feature, [530–531](#)
  - implementation introduction, [550–554](#)
  - introduction, [520–521](#)
  - object-orientation, [531–532](#)
- POSTQUEL query language, [528–530](#)
- programming language, [552–554](#)
- rules systems, [538–547](#)
- status and performance, [554–555](#), [555f](#)
- storage systems, [547–550](#)
- Postgres perspective, [205](#)
  - active databases and rule systems, [212–213](#)
  - commercial adaptations, [220–223](#)
  - context, [205–206](#)
  - deep storage technologies, [215–216](#)
  - language support, [217–218](#)
  - lessons, [223–224](#)
  - log-centric storage and recovery, [213–215](#)
  - overview, [206–207](#)
  - software impact, [218–223](#)
  - XPRS architecture, [216–217](#)
- Postgres project, [102](#)
  - abstract data types, [88](#), [149](#)
  - description, [186](#)
  - ideas source, [149–150](#)
  - Illustra purchase, [111–112](#)
  - impact, [3](#)
  - POSTQUEL, [400](#)
  - productionization, [315](#)
  - satisfaction with, [78–79](#)
  - and SQL, [315–316](#)
  - start of, [203–204](#)
- PostgreSQL
  - creation, [170](#), [400](#)
  - impact, [3](#)
  - open source, [218–220](#), [317–319](#)
  - software architecture, [210](#)
- POSTQUEL query language
  - features, [528–530](#)
  - functions, [526–528](#), [555–556](#), [555f](#)
- Potamianos, Spyros, [212](#)
- Pragmatism in startup companies, [135–137](#)
- Pricing models, [134](#)
- Primary-copy replication control, [90](#)
- Primary indexes, [493](#)
- Primitives in one size fits all, [451–452](#)
- Princeton University, [43](#)



- Probabilistic reasoning in scientific data management, 266
- Problems
  - ignoring, 163
  - solving, 276–277
- Process 2 in Ingres, 585–591
- Process 3 in Ingres, 591–599
- Process 4 in Ingres, 599–602
- Process structure
  - Ingres, 571–575, 572f, 574f
  - Postgres, 551–552
- Project operators in C-Store, 509
- Project Oxygen, 225
- Project Sequoia
  - 2000, 112–113
  - Postgres, 215
- Projections in C-Store, 496–498, 510
- PROTECTION catalog in Ingres, 580
- Prototypes
  - ADT-Ingres, 202, 203f
  - Data Tamer project, 270–273
  - H-Store project, 247–250
  - Mariposa, 150
  - Morpheus, 150
  - MUFFIN, 200
  - noise in, 389
  - Postgres, 311–313
  - Shore, 152
  - startup companies, 120–121
  - Tamr project, 121
- PRS2 system, 212
- Punctuated Streams Team, 325
- Purify tool for Postgres productionization, 315
- Putzolu, Franco, 101
- PVLDB 2016 paper, 297–298
- Qatar Computing Research Institute (QCRI)
  - creation, 105
  - Data Civilizer project, 189
  - Data Tamer, 152
  - Tamr project, 120
- Quel language, 102
  - comments, 570–571
  - complex objects, 203
  - description, 195–196, 197f
  - overview, 563–565
  - and standardization, 398–399
  - utility commands, 565–568
- Query classes in H-Store, 480–482
- Query decomposition in Ingres, 591–597
- Query execution
  - C-Store, 509–510
  - H-Store, 474–475
- Query modeling and optimization in
  - BigDAWG, 285–286
- Query modification in Ingres, 90, 585–588
- Query optimization in C-Store, 510
- Query rewrite implementation in rules
  - systems, 538–541
- Quiet (friend), 27
- R-Tree index structure
  - Ingres, 201–202
  - and Postgres, 210
- Radical simplicity for Postgres transactional storage, 214
- RAID storage architectures, 216
- Raising money for startup companies, 127
- RAP project, 109
- RARES project, 109
- Rasin, Alex, 334
  - on Aurora project, 324f
- RCA company, 97–98
- Ré, Chris, 6
- Read-optimized systems, 492
- Real-time requirements as DBMS limitation, 227
- Real-world impact in rules of thumb, 236–238
- Record deduplication in Data Tamer project, 272–273
- Recorded Future company, 296–297
- Recovery
  - C-Store, 506–508
  - database logs for, 21–22
  - H-Store, 476–478
  - Ingres, 600–602

- Postgres, [213–216](#)
- Red Brick Systems, [66](#)
- Redo logs
  - H-Store, [482](#)
  - OLTP design, [470](#)
- “Reduction of Large Scale Markov Models for Random Chains” dissertation (Stonebraker), [43](#)
- Referential integrity, [538–539](#)
- Reformatting tuples in Ingres, [593](#)
- Relational-CODASYL debate, [404–406](#)
- Relational database industry birth, [94](#)
  - Ingres competition, [100–103](#)
  - Ingres team, [99–100](#)
  - Ingres timing, [98–99](#)
  - maturity stage, [103–105](#)
  - overview, [97–98](#)
- Relational database management systems (RDBMS)
  - industry birth timeline, [173–179](#), [174–176f](#), [178–179f](#)
  - Ingres later years, [194–198](#), [197f](#)
- Relational databases, brief history of, [2–6](#)
- Relational era, [3](#)
- Relational models in one size does not fit all world, [483–485](#)
- Relations
  - Ingres, [576–580](#)
  - QUEL, [566](#)
- Remote direct memory access (RDMA), [67](#)
- Rendezvous system, [404](#)
- REPLACE command in Ingres, [584](#), [600–601](#)
- Replication management in OLTP, [434–435](#)
- Research, open source impact on, [169–171](#)
- Research Applied to the National Needs (RANN) program, [147–148](#)
- Research contributions
  - 2010s and beyond, [188–189](#)
  - Berkeley years, [185–186](#)
  - MIT, [186](#)
  - one size doesn’t fit all era, [187–188](#)
  - technical rules of engagement, [183–185](#)
- Research story about Aurora project, [387–391](#)
- Resident set size (RSS) in VoltDB, [342](#)
- Resource control in OLTP design, [467](#)
- RETRIEVE commands in Ingres, [573](#)
- Reviews, unsatisfactory, [160–161](#)
- RFID (radio frequency identification)
  - tagging, [225n](#)
- Ries, Dan, [200](#)
- Rivers, Jonathan, [350](#)
- Robinson, John “JR”
  - Shared Nothing band, [138f](#)
  - Tamr, [365f](#)
  - Vertica Systems, [339](#)
- Rock fetches, [124](#)
- Rogers, Jennie, BigDAWG polystore system
  - article, [279–289](#)
- Rollbacks in C-Store, [506](#)
- Roots in OLTP tables, [471](#)
- Route 66 TV show, [81](#)
- Row store architecture, [492](#)
- Rowe, Larry
  - commercial Ingres codeline, [305](#)
  - Ingres founding, [303](#)
  - Postgres, [36](#), [88](#), [206](#), [217](#)
  - Postgres implementation seminal work, [519–559](#)
  - RTI founding, [101](#), [303](#)
- RS column store in C-Store, [500–501](#)
- RTI (Relational Technology, Inc.), [101](#)
  - commercial version, [111](#)
  - founding, [166–167](#), [398](#)
  - Ingres basis of, [198](#)
- Rubenstein, Brad, [205](#)
- Ruby-on-Rails system, [486](#)
- Rules systems in Postgres, [212–213](#)
  - complexity, [541–543](#)
  - implementation efficiency, [544](#)
  - introduction, [538–541](#)
  - knowledge management, [520](#)
  - push for, [316](#)
  - second system, [545–547](#)
  - views, [543–544](#)
- S-Store project, [234](#), [331](#)
- Sales in startup company guidelines, [128](#)

- Sales problem in one size fits all, 442
- Salz, Jon, 232, 328–330
- Sarawagi, Sunita, 215
- ScaLAPACK analytics, 349
- Scaling in Tamr, 362–365
- Schek, Hans, 92
- Schema mapping in Data Tamer project, 270–272, 271f
- Schieffer, Berni, 384, 385f
- Schlamb, Kelly, 385f
- Schreiber, Mark
  - Data Civilizer, 64, 294
  - Tamr, 361
- Schultz, Hayden, 331
- Schuster, Stu, 101
- SciDB codeline, 349
  - connectivity, 349–351
  - features focus, 351–352
  - genomic data, 354–355
  - hard numbers, 352–353
  - languages, 353–354
  - security, 354
- SciDB project
  - contributions for, 89
  - description, 47
  - one size doesn't fit all era, 188
- Scientific data management, 253–254
  - beginning tasks, 260–263, 261f
  - current users, 267–268
  - first users, 263–267
  - logistics, 259–260
  - mountain representation, 254–256
  - planning, 256–259
- Scientific databases in one size fits all, 457
- Scope in BigDAWG polystore system, 284
- SDTM (Study Data Tabulation Model), 364
- Search in one size fits all, 457
- Search User Interfaces (Hearst), 394
- Second System Effect, 223
- Secondary indexes, 493
- Secrecy in startup company guidelines, 127
- Security in SciDB, 354
- Segments in C-Store, 498
- Select operators in C-Store, 509
- Self-funded companies, 74
- Self-order values in C-Store, 501
- Selinger, Pat, 382, 384, 407
- Sensor-based applications in one size fits all, 445–446, 456
- SEQUEL language, 102
- Service of Stonebraker, 49
- Shankland, Jim, on Miro team, 314f
- Shared-nothing architecture in SciDB
  - codeline, 351
- Shared Nothing band, 80, 138f
- Sharma, Kristi Sen, SciDB codeline article, 349–355
- She, Zuohao (Jack), 373
- Shims in BigDAWG polystore system, 284
- Shore (Scalable Heterogeneous Object Repository), 418–419
  - architecture, 419–422, 420f
  - prototype, 152
  - removing components, 422–424
- Shore Storage Manager (SSM), 418
- Short One (friend), 27
- Sibley, Ed, 404, 406
- SIGFIDET conference, 403–408
- Singer, Adam, on Aurora project, 324f
- Single-partition transactions in H-Store
  - project, 247
- Single-sited transactions, 472, 474, 476
- Single threading in OLTP, 413, 416–417
- Skeen, Dale, 195, 200
- Skok, David, 151–152
- Sleepycat startup, 171
- Sloan Digital Sky Survey (SDSS), 113, 255–256
- Slotted pages in Shore, 419
- SMALLTALK language, 553
- Smooth (friend), 28
- Snapshot isolation, 495, 503–505
- Snodgrass, Rick, 198
- Software impact in Postgres, 218–223
- Solicitation in startup company guidelines, 123–125
- Sort keys and operators in C-Store, 497–498, 509

- Source code for Ingres project, [109–110](#)
- SOURCE directory in Ingres, [576](#)
- Space budget in C-Store, [500](#)
- Spanner, [70–71](#)
- Spark, [402](#)
- Spending money guidelines for startup companies, [125–126](#)
- Sprite distributed OS, [216](#)
- SQL language, [102](#)
  - introduction, [404](#)
  - MapReduce, [114](#)
  - one size does not fit all world, [485–486](#)
  - and Postgres, [315–316](#)
  - vs. Quel, [196](#)
- SQuAl system, [323](#)
- Stable memory in Postgres, [548](#)
- Stanford Linear Accelerator (SLAC) facility, [257](#)
- Star schema in data warehousing, [443, 444f](#)
- Starburst project, [213](#)
- Startup companies, founded, [48–49](#)
- Startup companies, guidelines
  - business acumen on team, [122–123](#)
  - company control, [127](#)
  - first customers, [126](#)
  - ideas, [119–120](#)
  - intellectual property, [126](#)
  - introduction, [119](#)
  - lighthouse customers, [122](#)
  - mistakes, [128](#)
  - pitch deck and VC solicitation, [123–125](#)
  - raising money, [127](#)
  - sales, [128](#)
  - secrecy, [127](#)
  - spending money, [125–126](#)
  - summary, [128](#)
  - teams and prototypes, [120–121](#)
  - venture capitalists, [127](#)
- Startup companies, running
  - introduction, [129–130](#)
  - overview, [130–132](#)
  - partnerships, [132–135](#)
  - people in, [138](#)
  - pragmatism, [135–137](#)
- State storage in one size fits all, [458](#)
- Status in Postgres, [554–555, 555f](#)
- Sterile transaction classes, [473, 476](#)
- Stonebraker, Beth, [50, 141, 144](#)
- Stonebraker, Leslie, [50, 141, 144](#)
- Stonebraker, Michael
  - collected works, [607–633](#)
  - failures article, [155–164, 156f](#)
  - ideas article, [147–153](#)
  - Postgres design, construction, and commercialization story, [15–37](#)
  - startup company guidelines. *See* Startup companies, guidelines
  - Winslett interview, [59–83](#)
- Stonebraker, Michael, biography overview
  - academic career and birth of Ingres, [43–45](#)
  - academic positions, [43](#)
  - advocacy, [50, 91–92](#)
  - awards and honors, [49](#)
  - career flowchart, [54–55](#)
  - companies founded, [48–49](#)
  - early years and education, [42–43](#)
  - industry, MIT, and new millennium, [46–47](#)
  - legacy, [47–48](#)
  - personal life, [50](#)
  - post-Ingres years, [45–46](#)
  - sabbatical at MIT, [113–114](#)
  - student genealogy chart, [52–53, 56f](#)
  - synopsis, [41–42](#)
- Stonebraker, Michael, seminal works
  - C-Store, [491–518](#)
  - end of architectural era, [463–489](#)
  - Ingres implementation, [561–605](#)
  - OLTP databases, [411–439](#)
  - one size fits all, [441–462](#)
  - Postgres implementation, [519–559](#)
- Stonebraker's good ideas
  - abstract data types, [148](#)
  - Data Tamer, [152](#)
  - data warehouses, [151](#)
  - distributed databases, [150](#)
  - H-Store/VoltDB, [151–152](#)

- Stonebraker's good ideas (*continued*)
  - how to exploit, 153
  - Ingres, 147–148
  - parallel databases, 151
  - Postgres, 149–150
  - startup company guidelines, 119–120
- Stonebraker, Sandra, 50, 141, 144
- Storage allocators in C-Store, 502–503
- Storage keys in C-Store, 498
- Storage management and structures
  - C-Store, 502–503
  - commercial Ingres codeline, 306–308
  - Ingres, 580–582, 582f
  - Postgres, 213–216, 547–550
  - QUEL, 566–567
- Stored procedures, 91
- STRATEGY program in OVQP, 597–599
- Stream processing era
  - Aurora and Borealis origins, 225–227
  - Aurora and Borealis systems, 227–231
  - concurrent efforts, 231–232
  - current systems, 233–234
  - StreamBase Systems, 232–233
- Stream processing in one size fits all, 445–447, 448f
- STREAM project, 231
- Stream-SQL, enthusiasm for, 484
- STREAM Team, 325
- StreamBase codelines, 321–322, 322f
  - April Fool's Day joke, 330–331
  - conclusion, 331–332
  - customers, 327–328
  - development, 326–327
  - issues, 328–330
- StreamBase Systems
  - Architecture Committee, 329
  - aggregation systems, 451
  - from Aurora, 89
  - founding, 46, 232–233
  - Grassy Brook renamed to, 142
  - textual language, 229
- Strongly two-phase applications, 473
- Student genealogy chart, 52–53, 56f
- Student perspective, 393–396, 395f
- Subject matter experts (SMEs) in Tamr, 361–362
- Survey of Income and Program Participation (SIPP) data, 197–198
- Sybase, 400
- Synchronization in one size fits all, 455
- Sysplex Coupling Facility, 380
- System catalogs in Ingres, 578–580
- System i, 380
- System-level data management problems
  - and approaches, 91–92
- System R system, 88, 100
  - architectural features, 465
  - code base, 380
  - development, 44
  - vs. Ingres, 196
- Systems, leadership and advocacy, 87–90
- Szalay, Alex, 112–113
- Szolovits, Peter, 370
- T-trees, 436
- Table fragments in H-Store, 475
- Tall Shark (friend), 27
- Tamr codeline, 357
  - algorithmic complexity, 359–361
  - conclusion, 365–366, 366f
  - Data Unification, 358–359
  - user emphasis, 361–362
  - variety, 362–365
- Tamr project and company
  - creation, 105
  - from Data Tamer, 90
  - founding, 273–275
  - idea for, 120, 152
  - prototype, 120
- Tandem Computers, 101
- Tang, Nan
  - Data Civilizer, 189
  - Data Civilizer article, 291–300
- Tango, Jo
  - at Highland event, 132
  - venture capitalist perspective article, 139–144
- Tarashansky, Igor, 352

- Tatbul, Nesime
  - Aurora/Borealis/StreamBase codelines article, [321](#)–[332](#)
  - Aurora/Borealis/StreamBase reunion, [332f](#)
- Taylor, Cimarron, on Miro team, [314f](#)
- Teams for startup companies, [120](#)–[121](#)
- Technical rules of engagement, [183](#)–[185](#)
- Technology Licensing Offices (TLOs), [126](#)
- Telegraph Team, [325](#)
- TelegraphCQ project, [231](#)
- Telnav company, [90](#)
- Temporal Functional Dependencies in Data Civilizer, [297](#)
- Tenure paper requirements, [159](#)
- Teradata, [222](#)
- Term sheets in startup company guidelines, [124](#)–[125](#)
- Terminal monitor in Ingres, [572](#)
- Test-of-time award, [202](#)
- Text search in one size fits all, [457](#)
- Thomson Reuters (TR) company, [274](#)
- Thread support in Shore, [420](#)–[421](#)
- Three dimensional problems, [520](#)
- Throughput in OLTP, [428](#)
- Tibbetts, Richard
  - Aurora/Borealis/StreamBase reunion, [332f](#)
  - StreamBase development, [326](#)–[327](#)
  - StreamBase issues, [328](#)–[330](#)
  - StreamBase Systems, [232](#)
- TIBCO Software, Inc., [233](#), [321](#)
- Time travel feature in Postgres, [316](#), [529](#), [548](#), [550](#)
- TimeSeries DataBlade in Postgres, [317](#)
- Timestamp authorities (TAs) in C-Store, [504](#)
- TimesTen system, [436](#)
- TMP directory in Ingres, [576](#)
- TPC (Transaction Processing Performance Council) benchmark
  - Data Unification, [358](#)
  - H-Store, [479](#)–[483](#), [479f](#)
  - OLTP, [424](#)–[425](#), [425f](#)
  - TPC-B, [382](#), [436](#)
- Training workloads in C-Store, [500](#)
- Trajman, Omer, [335](#)
- Tran, Nga, C-Store seminal work, [491](#)–[518](#)
- Transaction-less databases, [413](#)
- Transactions
  - C-Store, [503](#)–[508](#), [505f](#)
  - concurrency control. *See* Concurrency control
  - features, [470](#)–[471](#)
  - H-Store, [247](#), [476](#)–[478](#)
  - OLTP, [417](#)–[418](#)
  - rollbacks, [506](#)
  - schema characteristics, [471](#)–[473](#)
- Transitive closure in Postgres, [529](#)
- Trees schemas, [471](#)–[472](#)
- Triggers
  - DBMS limitation, [226](#)
  - one size fits all, [450](#)
  - Postgres, [213](#)
  - rules systems, [212](#), [539](#)–[540](#)
- Triple Rock (friend), [27](#)
- Trust with venture capitalists, [140](#)
- Tsichritzis, Dennis, [403](#)–[404](#)
- Tuple movers
  - C-Store, [494](#)–[495](#), [508](#)
  - Vertica Systems, [337](#)
- Tuple storage in VoltDB, [343](#)
- Tuples in Ingres
  - AMI, [583](#)
  - substitution, [592](#)
  - TIDs, [581](#)
  - variables, [563](#)
- Turing Award in 2014
  - citation, [130](#)
  - overview, [114](#)–[115](#)
  - perspectives, [93](#)–[95](#)
- Two-phase applications, [473](#), [476](#)
- Types in Postgres, [523](#), [532](#)–[537](#)
- Ubell, Michael
  - Illustra, [221](#)
  - Miro team, [314f](#)
  - Postgres productionization, [315](#)

- Undo logs
  - H-Store, 482
  - OLTP design, 471
- UNIN process structure, 571–572
- Union types in Postgres, 532–534
- Unix platforms
  - Ingres, 101, 571–573
  - systems based on, 47
- Updates
  - C-Store, 503–508, 505f
  - Ingres, 600–602
- Uptone (friend), 28, 30
- Urban Dynamics, 43
- Urban systems, 147–148
- User-Defined Aggregate (UDA) functions in
  - Postgres, 209
- User-defined extensions (UDXs)
  - Ingres prototype, 202
  - SciDB codeline, 350
- User-defined functions (UDFs) in Postgres, 209, 211
- User-defined types (UDTs) in commercial
  - Ingres codeline, 308–309
- User emphasis in Tamr, 361–362
- User experience (UX) design and
  - implementation in Tamr, 361
- User feedback for Ingres, 602–603
- Utility commands in Ingres, 563–568, 599–602
  
- VanderPlas, Jake, 353
- Varaiya, Pravin, 17, 148
- Variety in Tamr, 361–365
- Venture capitalists
  - perspective, 139–144
  - in startup company guidelines, 127
- Verisk Health, 121, 152
- Vernica, Rares, 350, 353–354
- Vertica Systems
  - from C-Store, 89
  - creation, 104
  - founding, 133–135, 242–244
  - HP purchase of, 67
  - impact of, 113–114
  - patent infringement suit, 126
  - satisfaction with, 78–79
  - Tamr, 364
  - venture capitalist perspective, 142–143
- Vertica Systems codeline, 333
  - architectural decisions, 336–339
  - building, 333–334
  - conclusion, 340
  - customers, 334–335, 339–340
  - features discussion, 335–336
- Video data management, 265
- Vietnam war, 81
- Views
  - Ingres, 44, 586
  - rules systems, 543–544
- Vincent, Tim, 383–384, 385f
- VLDB demo paper for H-Store prototype, 249
- VLSI CAD design era, 201–202
- Voice-of-Experience (friend), 28, 30
- VoltDB
  - creation, 104
  - from H-Store, 89
  - H-Store executor, 75
  - H-Store split, 251
  - PayPal interest in, 250
- VoltDB codeline, 341–342
  - compaction, 342–344
  - disk persistence, 346–347
  - latency, 344–348
- Volume in Tamr, 361–362
  
- Weak consistency in OLTP, 435
- Wei Hong Optimizer, 217
- Weisberg, Ariel, 249
- Whales, 158
- Whitney, Kevin, 404
- Whittaker, Andrew, 370
- Whyte, Nick, 99
- Widom, Jennifer, 228–229, 325
- Winer, Mike, 385f
- WinFS project, 215n
- Winslete, Marianne, interview with
  - Stonebraker, 59–83

- Wisconsin, 1996 fall, 108–111
- Wisconsin Benchmark, 436
- Wong, Eugene, 98
  - Ingres, 88
  - Ingres founding, 41, 148, 303
  - Ingres implementation seminal work, 561–605
  - RTI founding, 398
  - Stonebraker guided by, 3, 43
- Worker sites in H-Store, 477
- Workflow-based diagrammatic languages, 227n
- Workload in OLTP, 425–427, 426f
- Write-ahead logging in Postgres, 214
- Write-optimized systems, 492
- WS column store in C-Store, 502, 507–508
  
- Xiao, Min
  - OMDB, 338
  - Vertica Systems, 339
- Xing, Ying, on Aurora project, 324f
- XML databases in one size fits all, 457–458
- XPRS architecture in Postgres, 216–217
- XQuery language, 208
- XRM-An Extended (N-ary) Relational Memory, 404
  
- Yan, Robin, on Aurora project, 324f
- Youssefi, Karel
  - Ingres team, 99
  - Tandem Computers, 101
- Yu, Andrew
  - Postgres parser, 218
  - PostgreSQL, 170
  - SQLization project, 317–318
- Yu, Katherine, 376
  
- Zaniolo, Carlo, 217
- Zdonik, Stan, 103
  - Aurora/Borealis/StreamBase reunion, 332f
  - Aurora project, 322–323, 324f
  - Borealis project, 186
  - expert sourcing, 273
  - H-Store project, 245–246
  - Shared Nothing band, 138f
  - stream processing era article, 225–234
  - StreamBase Systems, 232
  - Tamr project, 120
  - Vertica Systems, 334, 339
- Zero-billion-dollar ideas, 185
- Zhang, Donghui, 353–354
- Zilles, Stephen, 88
- Zook, Bill, 99





## Biographies

### Editor

#### Michael L. Brodie



**Michael L. Brodie** has over 45 years of experience in research and industrial practice in databases, distributed systems, integration, artificial intelligence, and multidisciplinary problem-solving. Dr. Brodie is a research scientist at the Computer Science and Artificial Intelligence Laboratory, Massachusetts Institute of Technology; advises startups; serves on advisory boards of national and international research organizations; and is an adjunct professor at the National University of Ireland, Galway and at the University of Technology, Sydney. As Chief Scientist of IT at Verizon for over 20 years, he was responsible for advanced technologies, architectures, and methodologies for IT strategies and for guiding industrial-scale deployments of emerging technologies. He has served on several National Academy of Science committees. Current interests include Big Data, Data Science, and Information Systems evolution. Dr. Brodie holds a Ph.D. in databases from the University of Toronto and a Doctor of Science (*honoris causa*) from the National University of Ireland. Visit [www.Michaelbrodie.com](http://www.Michaelbrodie.com) for further information.

## Authors

### Daniel J. Abadi



**Daniel J. Abadi** is the Darnell-Kanal Professor of Computer Science at the University of Maryland, College Park. He performs research on database system architecture and implementation, especially at the intersection of scalable and distributed systems. He is best known for the development of the storage and query execution engines of the C-Store (column-oriented database) prototype, which was commercialized by Vertica and eventually acquired by Hewlett-Packard, and for his HadoopDB research on fault-tolerant scalable analytical database systems, which was commercialized by Hadapt and acquired by Teradata in 2014. Abadi has been a recipient of a Churchill Scholarship, a NSF CAREER Award, a Sloan Research Fellowship, a VLDB Best Paper Award, a VLDB 10-year Best Paper Award, the 2008 SIGMOD Jim Gray Doctoral Dissertation Award, the 2013–2014 Yale Provost's Teaching Prize, and the 2013 VLDB Early Career Researcher Award. He received his Ph.D. in 2008 from MIT. He blogs at DBMS Musings (<http://dbmsmusings.blogspot.com>) and Tweets at @daniel\_abadi.

### Magdalena Balazinska



**Magdalena Balazinska** is a professor in the Paul G. Allen School of Computer Science and Engineering at the University of Washington and is the director of the University's eScience Institute. She's also director of the IGERT PhD Program in Big Data and Data Science and the associated Advanced Data Science PhD Option. Her research interests are in database management systems with a current focus on data management for data science, big data systems, and cloud computing. Magdalena holds a Ph.D. from the Massachusetts Institute of Technology (2006). She is a Microsoft Research New Faculty Fellow (2007) and received the inaugural VLDB Women in Database Research Award (2016), an ACM SIGMOD Test-of-Time Award (2017), an NSF CAREER Award (2009), a 10-year most influential paper award (2010), a Google Research Award (2011),

an HP Labs Research Innovation Award (2009 and 2010), a Rogel Faculty Support Award (2006), a Microsoft Research Graduate Fellowship (2003–2005), and multiple best-paper awards.

### **Nikolaus Bates-Haus**



**Nikolaus Bates-Haus** is Technical Lead at Tamr Inc., an enterprise-scale data unification company, where he assembled the original engineering team and led the development of the first generation of the product. Prior to joining Tamr, Nik was Lead Architect and Director of Engineering at Endeca (acquired by Oracle in 2011), where he led development of the MDEX analytical database engine, a schema-on-read column store designed for large-scale parallel query evaluation. Previously, Nik worked in data integration, machine learning, parallel computation, and real-time processing at Torrent Systems, Thinking Machines, and Philips Research North America. Nik holds an M.S. in Computer Science from Columbia University and a B.A. in Mathematics/Computer Science from Wesleyan University. Tamr is Nik's seventh startup.

### **Philip A. Bernstein**



**Philip A. Bernstein** is a Distinguished Scientist at Microsoft Research, where he has worked for over 20 years. He is also an Affiliate Professor of Computer Science at the University of Washington. Over the last 20 years, he has been a product architect at Microsoft and Digital Equipment Corp., a professor at Harvard University and Wang Institute of Graduate Studies, and a VP Software at Sequoia Systems. He has published over 150 papers and 2 books on the theory and implementation of database systems, especially on transaction processing and data integration, and has contributed to a variety of database products. He is an ACM Fellow, an AAAS Fellow, a winner of ACM SIGMOD's Codd Innovations Award, a member of the Washington State Academy of Sciences, and a member of the U.S. National Academy of Engineering. He received a B.S. from Cornell and M.Sc. and Ph.D. degrees from the University of Toronto.

### Janice L. Brown




**Janice L. Brown** is president and founder of Janice Brown & Associates, Inc., a communications consulting firm. She uses strategic communications to help entrepreneurs and visionary thinkers launch technology companies, products, and ventures, as well as sell their products and ideas. She has been involved in three ventures (so far) with 2014 Turing Award-winner Michael Stonebraker: Vertica Systems, Tamr, and the Intel Science and Technology Center for Big Data. Her background includes positions at several public relations and advertising agencies, and product PR positions at two large technology companies. Her work for the Open Software Foundation won the PRSA's Silver Anvil Award, the "Oscar" of the PR industry. Brown has a B.A. from Simmons College. Visit [www.janicebrown.com](http://www.janicebrown.com).

### Paul Brown



**Paul Brown** first met Mike Stonebraker in early 1992 at Brewed Awakening coffee shop on Euclid Avenue in Berkeley, CA. Mike and John Forrest were interviewing Paul to take over the job Mike Olson had just left. Paul had a latte. Mike had tea. Since then, Paul has worked for two of Mike's startups: Illustra Information Technologies and SciDB / Paradigm4. He was co-author with Mike of a book and a number of research papers. Paul has worked for a series of DBMS companies all starting with the letter "I": Ingres, Illustra, Informix, and IBM. Alliterative ennui setting in, Paul joined Paradigm4 as SciDB's Chief Architect. He has since moved on to work for Teradata. Paul likes dogs, DBMSs, and (void \*). He hopes he might have just picked up sufficient gravitas in this industry to pull off the beard.

### Paul Butterworth



**Paul Butterworth** served as Chief Systems Architect at Ingres from 1980–1990. He is currently co-founder and Chief Technology Officer (CTO) at VANTIQ, Inc. His past roles include Executive Vice President, Engineering at You Technology Inc., and co-founder and CTO of Emotive Communications, where he conceived and designed the Emotive Cloud Platform for enterprise mobile computing. Before that, Paul was an architect at Oracle and a founder & CTO at AmberPoint, where he directed the technical strategy for the AmberPoint SOA governance products. Prior to AmberPoint, Paul was a Distinguished Engineer and Chief Technologist for the Developer Tools Group at Sun Microsystems and a founder, Chief Architect, and Senior Vice President of Forte Software. Paul holds undergraduate and graduate degrees in Computer Science from UC Irvine.

### Michael J. Carey



**Michael J. Carey** received his B.S. and M.S. from Carnegie-Mellon University and his Ph.D. from the University of California, Berkeley, in 1979, 1981, and 1983, respectively. He is currently a Bren Professor of Information and Computer Sciences at the University of California, Irvine (UCI) and a consulting architect at Couchbase, Inc. Before joining UCI in 2008, Mike worked at BEA Systems for seven years and led the development of BEA's AquaLogic Data Services Platform product for virtual data integration. He also spent a dozen years teaching at the University of Wisconsin-Madison, five years at the IBM Almaden Research Center working on object-relational databases, and a year and a half at Propel Software, an e-commerce platform startup, during the infamous 2000–2001 Internet bubble. He is an ACM Fellow, an IEEE Fellow, a member of the National Academy of Engineering, and a recipient of the ACM SIGMOD E.F. Codd Innovations Award. His current interests center on data-intensive computing and scalable data management (a.k.a. Big Data).

**Fred Carter**

**Fred Carter**, a software architect in a variety of software areas, worked at Ingres Corporation in several senior positions, including Principal Scientist/Chief Architect. He is currently a principal architect at VANTIQ, Inc. Prior to VANTIQ, Fred was the runtime architect for AmberPoint, which was subsequently purchased by Oracle. At Oracle, he continued in that role, moving the AmberPoint system to a cloud-based, application performance monitoring service. Past roles included architect for EAI products at Forte (continuing at Sun Microsystems) and technical leadership positions at Oracle, where he designed distributed object services for interactive TV, online services, and content management, and chaired the Technical Committee for the Object Definition Alliance to foster standardization in the area of network-based multimedia systems. Fred has an undergraduate degree in Computer Science from Northwestern University and received his M.S. in Computer Science from UC Berkeley.

**Raul Castro Fernandez**

**Raul Castro Fernandez** is a postdoc at MIT, working with Samuel Madden and Michael Stonebraker on data discovery—how to help people find relevant data among databases, data lakes, and the cloud. Raul built Aurum, a data discovery system, to identify relevant data sets among structured data. Among other research lines, he is looking at how to incorporate unstructured data sources, such as PDFs and emails. More generally, he is interested in data-related problems, from efficient data processing to machine learning engineering. Before MIT, Raul completed his Ph.D. at Imperial College London, where he focused on designing new abstractions and building systems for large-scale data processing.



### Ugur Çetintemel



**Ugur Çetintemel** is a professor in the department of Computer Science at Brown University. His research is on the design and engineering of high-performance, user-friendly data management and processing systems that allow users to analyze large data sets interactively. Ugur chaired SIGMOD '09 and served on the editorial boards of *VLDB Journal*, *Distributed and Parallel Databases*, and *SIGMOD Record*. He is the recipient of a National Science Foundation Career Award and an IEEE 10-year test of time award in Data Engineering, among others. Ugur was a co-founder and a senior architect of StreamBase, a company that specializes in high-performance data processing. He was also a Brown Manning Assistant Professor and has been serving as the Chair of the Computer Science Department at Brown since July 2014.

### Xuedong Chen

**Xuedong Chen** is currently an Amazon.com Web Services software developer in Andover, Massachusetts. From 2002–2007 he was a Ph.D. candidate at UMass Boston, advised by Pat and Betty O'Neil. He, along with Pat O'Neil and others, were co-authors with Mike Stonebraker.

### Mitch Cherniack



**Mitch Cherniack** is an Associate Professor at Brandeis University. He is a previous winner of an NSF Career Award and co-founder of Vertica Systems and StreamBase Systems. His research in Database Systems has focused on query optimization, streaming data systems, and column-based database architectures. Mitch received his Ph.D. from Brown University in 1999, an M.S. from Concordia University in 1992, and a B.Ed. from McGill University in 1984.



### David J. DeWitt



**David J. DeWitt** joined the Computer Sciences Department at the University of Wisconsin in September 1976 after receiving his Ph.D. from the University of Michigan. He served as department chair from July 1999 to July 2004. He held the title of John P. Morgridge Professor of Computer Sciences when he retired from the University of Wisconsin in 2008. In 2008, he joined Microsoft as a Technical Fellow to establish and manage the Jim Gray Systems Lab in Madison. In 2016, he moved to Boston to join the MIT Computer Science and AI Laboratory as an Adjunct Professor. Professor DeWitt is a member of the National Academy of Engineering (1998), a fellow of the American Academy of Arts and Sciences (2007), and an ACM Fellow (1995). He received the 1995 Ted Codd SIGMOD Innovations Award. His pioneering contributions to the field of scalable database systems for “big data” were recognized by ACM with the 2009 Software Systems Award.

### Aaron J. Elmore



**Aaron J. Elmore** is an assistant professor in the Department of Computer Science and the College of the University of Chicago. Aaron was previously a postdoctoral associate at MIT working with Mike Stonebraker and Sam Madden. Aaron’s thesis on *Elasticity Primitives for Database-as-a-Service* was completed at the University of California, Santa Barbara under the supervision of Divy Agrawal and Amr El Abbadi. Prior to receiving a Ph.D., Aaron spent several years in industry and completed an M.S. at the University of Chicago.

### Miguel Ferreira

**Miguel Ferreira** is an alumnus of MIT. He was coauthor of the paper, “Integrating Compression and Execution in Column-Oriented Database Systems,” while working with Samuel Madden and Daniel Abadi, and “C-store: A Column-Oriented DBMS,” with Mike Stonebraker, Daniel Abadi, and others.

### Vijay Gadepally




**Vijay Gadepally** is a senior member of the technical staff at the Massachusetts Institute of Technology (MIT) Lincoln Laboratory and works closely with the Computer Science and Artificial Intelligence Laboratory (CSAIL). Vijay holds an M.Sc. and Ph.D. in Electrical and Computer Engineering from The Ohio State University and a B.Tech in Electrical Engineering from the Indian Institute of Technology, Kanpur. In 2011, Vijay received an Outstanding Graduate Student Award at The Ohio State University. In 2016, Vijay received the MIT Lincoln Laboratory's Early Career Technical Achievement Award and in 2017 was named to AFCEA's inaugural 40 under 40 list. Vijay's research interests are in high-performance computing, machine learning, graph algorithms, and high-performance databases.

### Nabil Hachem



**Nabil Hachem** is currently Vice President, Head of Data Architecture, Technology, and Standards at MassMutual. He was formerly Global Head of Data Engineering at Novartis Institute for Biomedical Research, Inc. He also held senior data engineering posts at Vertica Systems, Inc., Infinity Pharmaceuticals, Upromise Inc., Fidelity Investments Corp., and Ask Jeeves Inc. Nabil began his career as an electrical engineer and operations department manager for a data telecommunications firm in Lebanon. In addition to his commercial career, Nabil taught computer science at Worcester Polytechnic Institute. He co-authored dozens of papers on scientific databases, file structures, and join algorithms, among others. Nabil received a degree in Electrical Engineering from the American University of Beirut and earned his Ph.D. in Computer Engineering from Syracuse University.

**Don Haderle**

**Don Haderle** joined IBM in 1968 as a software developer and retired in 2005 as the software executive operating as Chief Technology Officer (CTO) for Information Management. He consulted with venture capitalists and advised startups. He currently sits on technical advisory boards for a number of companies and consults independently. Considered the father of commercial high-performance, industrial-strength relational database systems, he was the technical leader and chief architect of DB2 from 1977–1998. He

led DB2's overall architecture and development, making key personal contributions to and holding fundamental patents in all key elements, including: logging primitives, memory management, transaction fail-save and recovery techniques, query processing, data integrity, sorting, and indexing. As CTO, Haderle collaborated with researchers to incubate new product directions for the information management industry. Don was appointed an IBM Fellow in 1989 and Vice President of Advanced Technology in 1991; named an ACM Fellow in 2000; and elected to the National Academy of Engineering in 2008. He is a graduate of UC Berkeley (B.A., Economics, 1967).

**James Hamilton**

**James Hamilton** is Vice President and Distinguished Engineer on the Amazon Web Services team, where he focuses on infrastructure efficiency, reliability, and scaling. He has spent more than 20 years working on high-scale services, database management systems, and compilers. Prior to joining AWS, James was architect on the Microsoft Data Center Futures team and the Windows Live Platform Services team. He was General Manager of the Microsoft Exchange Hosted Services team and has led many of the SQL Server en-

gineering teams through numerous releases. Before joining Microsoft, James was Lead Architect on the IBM DB2 UDB team. He holds a B.Sc. in Computer Science from the University of Victoria and a Master's in Math, Computer Science from the University of Waterloo.

### Stavros Harizopoulos



**Stavros Harizopoulos** is currently a Software Engineer at Facebook, where he leads initiatives on Realtime Analytics. Before that, he was a Principal Engineer at AWS Redshift, a petabyte-scale columnar Data Warehouse in the cloud, where he was leading efforts on performance and scalability. In 2011, he co-founded Amiato, a fully managed real-time ETL cloud service, which was later acquired by Amazon. In the past, Stavros has held research-scientist positions at HP Labs and MIT CSAIL, working on characterizing the energy efficiency of database servers, as well as dissecting the performance characteristics of modern in-memory and column-store databases. He is a Carnegie Mellon Ph.D. and a Y Combinator alumnus.

### Marti Hearst



**Marti Hearst** is a professor in the School of Information and the EECS Department at UC Berkeley. She was formerly a member of the research staff at Xerox PARC and received her Ph.D. from the CS Division at UC Berkeley. Her primary research interests are user interfaces for search engines, information visualization, natural language processing, and improving education. Her book *Search User Interfaces* was the first of its kind in academics. Prof. Hearst was named a Fellow of the ACM in 2013 and a member of the CHI Academy in 2017, and is president of the Association for Computational Linguistics. She has received four student-initiated Excellence in Teaching Awards.

**Jerry Held**

**Jerry Held** has been a successful Silicon Valley entrepreneur, executive, and investor for over 40 years. He has managed all growth stages of companies, from conception to multi-billion-dollar global enterprise. He is currently chairman of Tamr and Madaket Health and serves on the boards of NetApp, Informatica, and Copia Global. His past board service includes roles as executive chairman of Vertica Systems and MemSQL and lead independent director of Business Objects. Previously, Dr. Held was “CEO-in-residence” at venture capital firm Kleiner Perkins Caufield & Byers. He was senior vice president of Oracle Corporation’s server product division and a member of the executive team that grew Tandem Computers from pre-revenue to multi-billion-dollar company. Among many other roles, he led pioneering work in fault-tolerant, shared-nothing, and scale-out relational database systems. He received his Ph.D. in Computer Science from the University of California, Berkeley, where he led the initial development of the Ingres relational database management system.

**Pat Helland**

**Pat Helland** has been building databases, transaction systems, distributed systems, messaging systems, multiprocessor hardware, and scalable cloud systems since 1978. At Tandem Computers, he was Chief Architect of the transaction engine for NonStop SQL. At Microsoft, he architected Microsoft Transaction Server, Distributed Transaction Coordinator, SQL Service Broker, and evolved the Cosmos big data infrastructure to include optimizing database features as well as petabyte-scale transactionally correct event processing. While at Amazon, Pat contributed to the design of the Dynamo eventually consistent store and also the Product Catalog. Pat attended the University of California, Irvine from 1973–1976 and was in the inaugural UC Irvine Information and Computer Science Hall of Fame. Pat chairs the Dean’s Leadership Council of the Donald Bren School of Information and Computer Sciences (ICS), UC Irvine.



### Joseph M. Hellerstein



**Joseph M. Hellerstein** is the Jim Gray Professor of Computer Science at the University of California, Berkeley, whose work focuses on data-centric systems and the way they drive computing. He is an ACM Fellow, an Alfred P. Sloan Research Fellow, and the recipient of three ACM-SIGMOD “Test of Time” awards for his research. In 2010, *Fortune Magazine* included him in their list of 50 smartest people in technology, and MIT’s *Technology Review* magazine included his work on their TR10 list of the 10 technologies “most likely to change our world.” Hellerstein is the co-founder and Chief Strategy Officer of Trifacta, a software vendor providing intelligent interactive solutions to the messy problem of wrangling data. He serves on the technical advisory boards of a number of computing and Internet companies including Dell EMC, SurveyMonkey, Captricity, and Datometry, and previously served as the Director of Intel Research, Berkeley.

### Wei Hong



**Wei Hong** is an engineering director in Google’s Data Infrastructure and Analysis (DIA) group, responsible for the streaming data processing area including building and maintaining the infrastructure for some of Google’s most revenue-critical data pipelines in Ads and Commerce. Prior to joining Google, he co-founded and led three startup companies: Illustra and Cohera with Mike Stonebraker in database systems and Arch Rock in Internet of Things. He also held senior engineering leadership positions at Informix, PeopleSoft, Cisco, and Nest. He was a senior researcher at Intel Research Berkeley working on sensor networks and streaming database systems and won an ACM SIGMOD Test of Time Award. He is a co-inventor of 80 patents. He received his Ph.D. from UC Berkeley and has ME, BE, and BS from Tsinghua University.

### John Hugg



**John Hugg** has had a deep love for problems relating to data. He's worked at three database product startups and worked on database problems within larger organizations as well. Although John dabbled in statistics in graduate school, Dr. Stonebraker lured him back to databases using the nascent VoltDB project. Working with the very special VoltDB team was an unmatched opportunity to learn and be challenged. John received an M.S in 2007 and a B.S. in 2005 from Tufts University.

### Ihab Ilyas



**Ihab Ilyas** is a professor in the Cheriton School of Computer Science at the University of Waterloo, where his main research focuses on the areas of big data and database systems, with special interest in data quality and integration, managing uncertain data, rank-aware query processing, and information extraction. Ihab is also a co-founder of Tamr, a startup focusing on large-scale data integration and cleaning. He is a recipient of the Ontario Early Researcher Award (2009), a Cheriton Faculty Fellowship (2013), an NSERC Discovery Accelerator Award (2014), and a Google Faculty Award (2014), and he is an ACM Distinguished Scientist. Ihab is an elected member of the VLDB Endowment board of trustees, elected SIGMOD vice chair, and an associate editor of *ACM Transactions on Database Systems* (TODS). He holds a Ph.D. in Computer Science from Purdue University and a B.Sc. and an M.Sc. from Alexandria University.

### Jason Kinchen



an avid cyclist and a Red Cross disaster action team volunteer.

**Jason Kinchen**, Paradigm4's V.P. of Engineering, is a software professional with over 30 years' experience in delivering highly complex products to life science, automotive, aerospace, and other engineering markets. He is an expert in leading technical teams in all facets of a project life cycle from feasibility analysis to requirements to functional design to delivery and enhancement, and experienced in developing quality-driven processes improving the software development life cycle and driving strategic planning. Jason is an

### Moshe Tov Kreps

**Moshe Tov Kreps** (formerly known as Peter Kreps) is a former researcher at the University of California at Berkeley and the Lawrence Berkeley National Laboratory. He was coauthor, with Mike Stonebraker, Eugene Wong, and Gerald Held, of the seminal paper, "The Design and Implementation of INGRES," published in the ACM Transactions on Database Systems in September 1976.

### Edmond Lau



directly with CTO's, directors, managers, and other emerging leaders to unlock what's possible for them. Edmond has been featured in the *New York Times*, *Forbes*, *Time*, *Slate*, *Inc.*, *Fortune*, and *Wired*. He blogs at [coleadership.com](http://coleadership.com), has a website ([www.theeffectiveengineer.com](http://www.theeffectiveengineer.com)), and tweets at @edmondlau.

**Edmond Lau** is the co-founder of Co Leadership, where his mission is to transform engineers into leaders. He runs leadership experiences, multi-week programs, and online courses to bridge people from where they are to the lives and careers they dream of. He's the author of *The Effective Engineer*, the now the de facto onboarding guide for many engineering teams. He's spent his career leading engineering teams across Silicon Valley at Quip, Quora, Google, and Ooyala. As a leadership coach, Edmond also works



## Shilpa Lawande



**Shilpa Lawande** is CEO and co-founder of postscript .us, an AI startup on a mission to free doctors from clinical paperwork. Previously, she was VP/GM HPE Big Data Platform, including its flagship Vertica Analytics Platform. Shilpa was a founding engineer at Vertica and led its Engineering and Customer Success teams from startup through the company's acquisition by HP. Shilpa has several patents and books on data warehousing to her name, and was named to the 2012 Mass High Tech Women to Watch list and Rev

Boston 20 in 2015. Shilpa serves as an advisor at Tamr, and as mentor/volunteer at two educational initiatives, Year Up (Boston) and CPathshala (India). Shilpa has a M.S. in Computer Science from the University of Wisconsin-Madison and a B.S. in Computer Science and Engineering from the Indian Institute of Technology, Mumbai.

## Amerson Lin



**Amerson Lin** received his B.S. and M.Eng both in Computer Science at MIT, the latter in 2005. He returned to Singapore to serve in the military and government before returning to the world of software. He was a consultant at Pivotal and then a business development lead at Palantir in both Singapore and the U.S. Amerson currently runs his own Insurtech startup—Gigacover—which delivers digital insurance to Southeast Asia.

### Samuel Madden



**Samuel Madden** is a professor of Electrical Engineering and Computer Science in MIT's Computer Science and Artificial Intelligence Laboratory. His research interests include databases, distributed computing, and networking. He is known for his work on sensor networks, column-oriented database, high-performance transaction processing, and cloud databases. Madden received his Ph.D. in 2003 from the University of California at Berkeley, where he worked on the TinyDB system for data collection from sensor networks. Madden was named one of Technology Review's Top 35 Under 35 (2005), and is the recipient of several awards, including an NSF CAREER Award (2004), a Sloan Foundation Fellowship (2007), VLDB best paper awards (2004, 2007), and a MobiCom 2006 best paper award. He also received "test of time" awards in SIGMOD 2013 and 2017 (for his work on Acquisitional Query Processing in SIGMOD 2003 and on Fault Tolerance in the Borealis system in SIGMOD 2007), and a ten-year best paper award in VLDB 2015 (for his work on the C-Store system).

### Tim Mattson



**Tim Mattson** is a parallel programmer. He earned his Ph.D. in Chemistry from the University of California, Santa Cruz for his work in molecular scattering theory. Since 1993, Tim has been with Intel Corporation, where he has worked on High Performance Computing: both software (OpenMP, OpenCL, RCCE, and OCR) and hardware/software co-design (ASCI Red, 80-core TFLOP chip, and the 48 core SCC). Tim's academic collaborations include work on the fundamental design patterns of parallel programming, the BigDAWG polystore system, the TileDB array storage manager, and building blocks for graphs "in the language of linear algebra" (the GraphBLAS). Currently, he leads a team of researchers at Intel working on technologies that help application programmers write highly optimized code that runs on future parallel systems. Outside of computing, Tim fills his time with coastal sea kayaking. He is an ACA-certified kayaking coach (level 5, advanced open ocean) and instructor trainer (level three, basic coastal).

### Felix Naumann



**Felix Naumann** studied Mathematics, Economics, and Computer Science at the University of Technology in Berlin. He completed his Ph.D. thesis on “Quality-driven Query Answering” in 2000. In 2001 and 2002, he worked at the IBM Almaden Research Center on topics of data integration. From 2003–2006, he was assistant professor for information integration at the Humboldt-University of Berlin. Since then, he has held the chair for information systems at the Hasso Plattner Institute at the University of Potsdam in Germany. He is Editor-in-Chief of *Information Systems*, and his research interests are in data profiling, data cleansing, and text mining.

### Mike Olson



**Mike Olson** co-founded Cloudera in 2008 and served as its CEO until 2013 when he took on his current role of chief strategy officer (CSO). As CSO, Mike is responsible for Cloudera’s product strategy, open-source leadership, engineering alignment, and direct engagement with customers. Prior to Cloudera, Mike was CEO of Sleepycat Software, makers of Berkeley DB, the open-source embedded database engine. Mike spent two years at Oracle Corporation as Vice President for Embedded Technologies after Oracle’s acquisition of Sleepycat in 2006. Prior to joining Sleepycat, Mike held technical and business positions at database vendors Britton Lee, Illustra Information Technologies, and Informix Software. Mike has a B.S. and an M.S. in Computer Science from the University of California, Berkeley. Mike tweets at @mikeolson.

### Elizabeth O’Neil

**Elizabeth O’Neil** (Betty) is a Professor of Computer Science at the University of Massachusetts, Boston. Her focus is research, teaching, and software development in database engines: performance analysis, transactions, XML support, Unicode support, buffering methods. In addition to her work for UMass Boston, she was, among other pursuits, a long-term (1977–1996) part-time Senior Scientist for Bolt, Beranek, and Newman, Inc., and during two sabbaticals was a full-time consultant for Microsoft Corporation. She is the owner of two patents owned by Microsoft.

### Patrick O'Neil

**Patrick O'Neil** is Professor Emeritus at the University of Massachusetts, Boston. His research has focused on database system cost-performance, transaction isolation, data warehousing, variations of bitmap indexing, and multi-dimensional databases/OLAP. In addition to his research, teaching, and service activities, he is the coauthor—with his wife Elizabeth (Betty)—of a database management textbook, and has been active in developing database performance benchmarks and corporate database consulting. He holds several patents.

### Mourad Ouzzani



**Mourad Ouzzani** is a principal scientist with the Qatar Computing Research Institute, HBKU. Before joining QCRI, he was a research associate professor at Purdue University. His current research interests include data integration, data cleaning, and building large-scale systems to enable science and engineering. He is the lead PI of Rayyan, a system for supporting the creation of systematic reviews, which had more than 11,000 users as of March 2017. He has extensively published in top-tier venues including SIGMOD, PVLDB, ICDE, and TKDE. He received Purdue University Seed for Success Awards in 2009 and 2012. He received his Ph.D. from Virginia Tech and his M.S. and B.S. from USTHB, Algeria.

### Andy Palmer



**Andy Palmer** is co-founder and CEO of Tamr, Inc., the enterprise-scale data unification company that he founded with fellow serial entrepreneur and 2014 Turing Award winner Michael Stonebraker, Ph.D., and others. Previously, Palmer was co-founder and founding CEO of Vertica Systems (also with Mike Stonebraker), a pioneering analytics database company (acquired by HP). He founded Koa Labs, a seed fund supporting the Boston/Cambridge entrepreneurial ecosystem, is a founder-partner at The Founder Collective, and holds a research affiliate position at MIT CSAIL. During his career as an entrepreneur, Palmer has served as Founder, founding investor, BoD member, or advisor to more than 60 startup companies in technology, healthcare, and the

life sciences. He also served as Global Head of Software and Data Engineering at Novartis Institutes for BioMedical Research (NIBR) and as a member of the start-up team and Chief Information and Administrative Officer at Infinity Pharmaceuticals (NASDAQ: INFI). Previously, he held positions at innovative technology companies Bowstreet, pcOrder.com, and Trilogy. He holds a BA from Bowdoin (1988) and an MBA from the Tuck School of Business at Dartmouth (1994).

### Andy Pavlo



**Andy Pavlo** is an assistant professor of Databaseology in the Computer Science Department at Carnegie Mellon University. He also used to raise clams. Andy received a Ph.D. in 2013 and an M.Sc. in 2009, both from Brown University, and an M.Sc. in 2006 and a B.Sc., both from Rochester Institute of Technology.

### Alex Poliakov



**Alex Poliakov** has over a decade of experience developing distributed database internals. At Paradigm4, he helps set the vision for the SciDB product and leads a team of Customer Solutions experts who help researchers in scientific and commercial applications make optimal use of SciDB to create new insights, products, and services for their companies. Alex previously worked at Netezza, after graduating from MIT's Course 6. Alex is into flying drones and producing drone videos.

**Alexander Rasin**

**Alexander Rasin** is an Associate Professor in the College of Computing and Digital Media (CDM) at DePaul University. He received his Ph.D. and M.Sc. in Computer Science from Brown University, Providence, RI. He is a co-Director of Data Systems and Optimization Lab at CDM and his primary research interest is in database forensics and cybersecurity applications of forensic analysis. Dr. Rasin's other research projects focus on building and tuning performance of domain-specific data management systems—currently in the areas of computer-aided diagnosis and software analytics. Several of his current research projects are supported by NSF.

**Jennie Rogers**

**Jennie Rogers** is the Lisa Wissner-Slivka and Benjamin Slivka Junior Professor in Computer Science and an Assistant Professor at Northwestern University. Before that she was a postdoctoral associate in the Database Group at MIT CSAIL where she worked with Mike Stonebraker and Sam Madden. She received her Ph.D. from Brown University under the guidance of Ugur Çetintemel. Her research interests include the management of science data, federated databases, cloud computing, and database performance modeling. Her

Erdős number is 3.




**Lawrence A. Rowe**

**Lawrence A. Rowe** is an Emeritus Professor of Electrical Engineering and Computer Science at U.C. Berkeley. His research interests are software systems and applications. His group developed the Berkeley Lecture Webcasting System that produced 30 course lecture webcasts each week viewed by over 500,000 people per month. His publications received three “best paper” and two “test of time” awards. He is an investor/advisor in The Batchery a Berkeley-based seed-stage incubator. Rowe is an ACM Fellow, a co-recipient of the 2002 U.C. Technology Leadership Council Award for IT Innovation, the recipient of the 2007 U.C. Irvine Donald Bren School of ICS Distinguished Alumni Award, the 2009 recipient of the ACM SIGMM Technical Achievement Award, and a co-recipient of the Inaugural ACM SIGMOD Systems Award for the development of modern object-relational DBMS. Larry and his wife Jean produce and sell award-winning premium wines using Napa Valley grapes under the Greyscale Wines brand.

**Kriti Sen Sharma**

**Kriti Sen Sharma** is a Customer Solutions Architect at Paradigm4. He works on projects spanning multiple domains (genomics, imaging, wearables, finance, etc.). Using his skills in collaborative problem-solving, algorithm development, and programming, he builds end-to-end applications that address customers’ big-data needs and enable them to gain business insights rapidly. Kriti is an avid blogger and also loves biking and hiking. Kriti received a Ph.D. in 2013 and an M.Sc. in 2009, both from Virginia Polytechnic Institute and State University, and an B.Tech. from Indian Institute of Technology, Kharagpur, in 2005.

### Nan Tang



**Nan Tang** is a senior scientist at Qatar Computing Research Institute, HBKU, Qatar Foundation, Qatar. He received his Ph.D. from the Chinese University of Hong Kong in 2007. He worked as a research staff member at CWI, the Netherlands, from 2008–2010. He was a research fellow at University of Edinburgh from 2010–2012. His current research interests include data curation, data visualization, and intelligent and immersive data analytics.

### Jo Tango



**Jo Tango** founded Kepha Partners. He has invested in the e-commerce, search engine, Internet ad network, wireless, supply chain software, storage, database, security, on-line payments, and data center virtualization spaces. He has been a founding investor in many Stonebraker companies: Goby (acquired by NAVTEQ), Paradigm4, StreamBase Systems (acquired by TIBCO), Vertica Systems (acquired by Hewlett-Packard), and VoltDB. Jo previously was at Highland Capital Partners for nearly nine years, where he was a General Partner. He also spent five years with Bain & Company, where he was based in Singapore, Hong Kong, and Boston, and focused on technology and startup projects. Jo attended Yale University (B.A., *summa cum laude* and Phi Beta Kappa) and Harvard Business School (M.B.A., Baker Scholar). He writes a personal blog at [jtangoVC.com](http://jtangoVC.com).



### Nesime Tatbul



**Nesime Tatbul** is a senior research scientist at the Intel Science and Technology Center at MIT CSAIL. Before joining Intel Labs, she was a faculty member at the Computer Science Department of ETH Zurich. She received her B.S. and M.S. in Computer Engineering from the Middle East Technical University (METU) and her M.S. and Ph.D. in Computer Science from Brown University. Her primary research area is database systems. She is the recipient of an IBM Faculty Award in 2008, a Best System Demonstration Award at SIGMOD 2005, and the Best Poster and the Grand Challenge awards at DEBS 2011. She has served on the organization and program committees for various conferences including SIGMOD (as an industrial program co-chair in 2014 and a group leader in 2011), VLDB, and ICDE (as a PC track chair for Streams, Sensor Networks, and Complex Event Processing in 2013).

### Nga Tran

**Nga Tran** is currently the Director of Engineering in the server development team at Vertica, where she has worked for the last 14 years. Previously, she was a Ph.D. candidate at Brandeis University, where she participated in research that contributed to Mike Stonebraker's research.

### Marianne Winslett



**Marianne Winslett** has been a professor in the Department of Computer Science at the University of Illinois since 1987, and served as the Director of Illinois's research center in Singapore, the Advanced Digital Sciences Center, from 2009–2013. Her research interests lie in information management and security, from the infrastructure level on up to the application level. She is an ACM Fellow and the recipient of a Presidential Young Investigator Award from the U.S. National Science Foundation. She is the former Vice-Chair of ACM SIGMOD and the former co-Editor-in-Chief of *ACM Transactions on the Web*, and has served on the editorial boards of *ACM Transactions on Database Systems*, *IEEE*

*Transactions on Knowledge and Data Engineering*, *ACM Transactions on Information and System Security*, *The Very Large Data Bases Journal*, and *ACM Transactions on the Web*. She has received two best paper awards for research on managing regulatory compliance data (VLDB, SSS), one best paper award for research on analyzing browser extensions to detect security vulnerabilities (USENIX Security), and one for keyword search (ICDE). Her Ph.D. is from Stanford University.

### Eugene Wong



**Eugene Wong** is Professor Emeritus at the University of California, Berkeley. His distinguished career includes contributions to academia, business, and public service. As Department Chair of EECS, he led the department through its greatest period of growth and into one of the highest ranked departments in its field. In 2004, the Wireless Foundation was established in Cory Hall upon completion of the Eugene and Joan C. Wong Center for Communications Research. He authored or co-authored over 100 scholarly articles and published 4 books, mentored students, and supervised over 20 dissertations. In 1980, he co-founded (with Michael Stonebraker and Lawrence A. Rowe) the INGRES Corporation. He was the Associate Director of the Office of Science and Technology Policy, under George H. Bush; from 1994–1996, he was Vice President for Research and Development for Hong Kong University of Science and Technology. He received the ACM Software System Award in 1988 for his work on INGRES, and was awarded the 2005 IEEE Founders Medal, with the apt citation: “For leadership in national and international engineering research and technology policy, for pioneering contributions in relational databases.”

**Stan Zdonik**

**Stan Zdonik** is a tenured professor of Computer Science at Brown University and a noted researcher in database management systems. Much of his work involves applying data management techniques to novel database architectures, to enable new applications. He is co-developer of the Aurora and Borealis stream processing engines, C-Store column store DBMS, and H-Store NewSQL DBMS, and has contributed to other systems including SciDB and the BigDAWG polystore system. He co-founded (with Michael Stonebraker)

two startup companies: StreamBase Systems and Vertica Systems. Earlier, while at Bolt Beranek and Newman Inc., Dr. Zdonik worked on the Prophet System, a data management tool for pharmacologists. He has more than 150 peer-reviewed papers in the database field and was named an ACM Fellow in 2006. Dr. Zdonik has a B.S in Computer Science and one in Industrial Management, an M.S. in Computer Science, and the degree of Electrical Engineer, all from MIT, where he went on to receive his Ph.D. in database management under Prof. Michael Hammer.